

Comparing Categories - Refinements

Maithreyi Gopalan
Week 5

Data viz in Wild

Saratessa

Sophia

Will and Maiko on deck

Reviewing Lab PS-2

Agenda

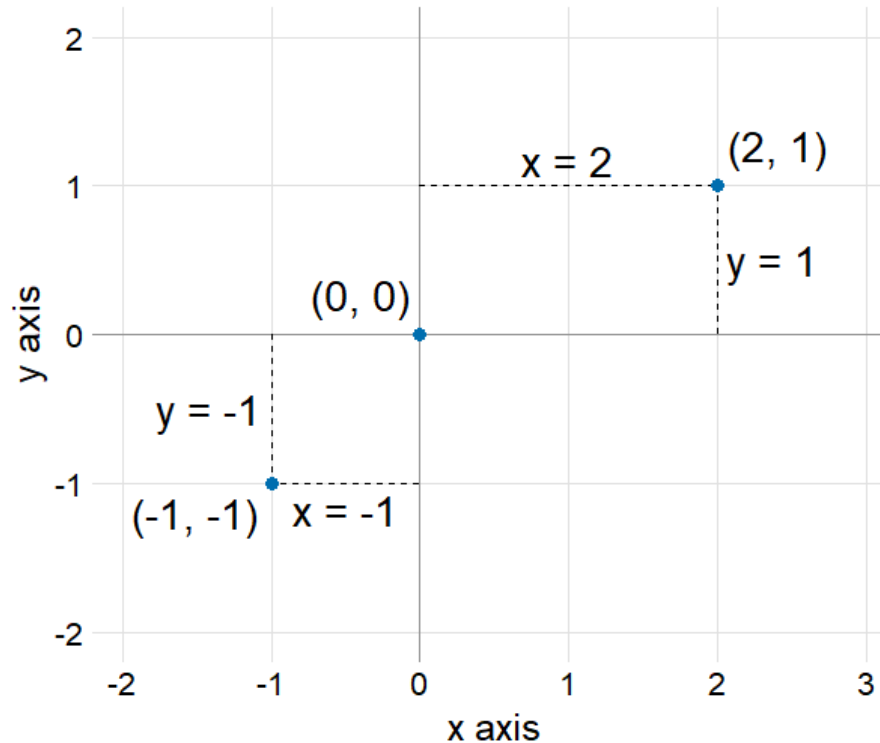
- Visualization for comparing categories - more considerations
- Aspect ratios, scales, and labels
- Annotations (most of the day)
- Saving plots (pretty quick)
- Compound figures
- Themes refresher (if time permits)
- Introduce Lab-PS3 quickly
- Lab 5
- Q & A with Dr. Daniel AndersonBecause we have a quest

Learning Objectives

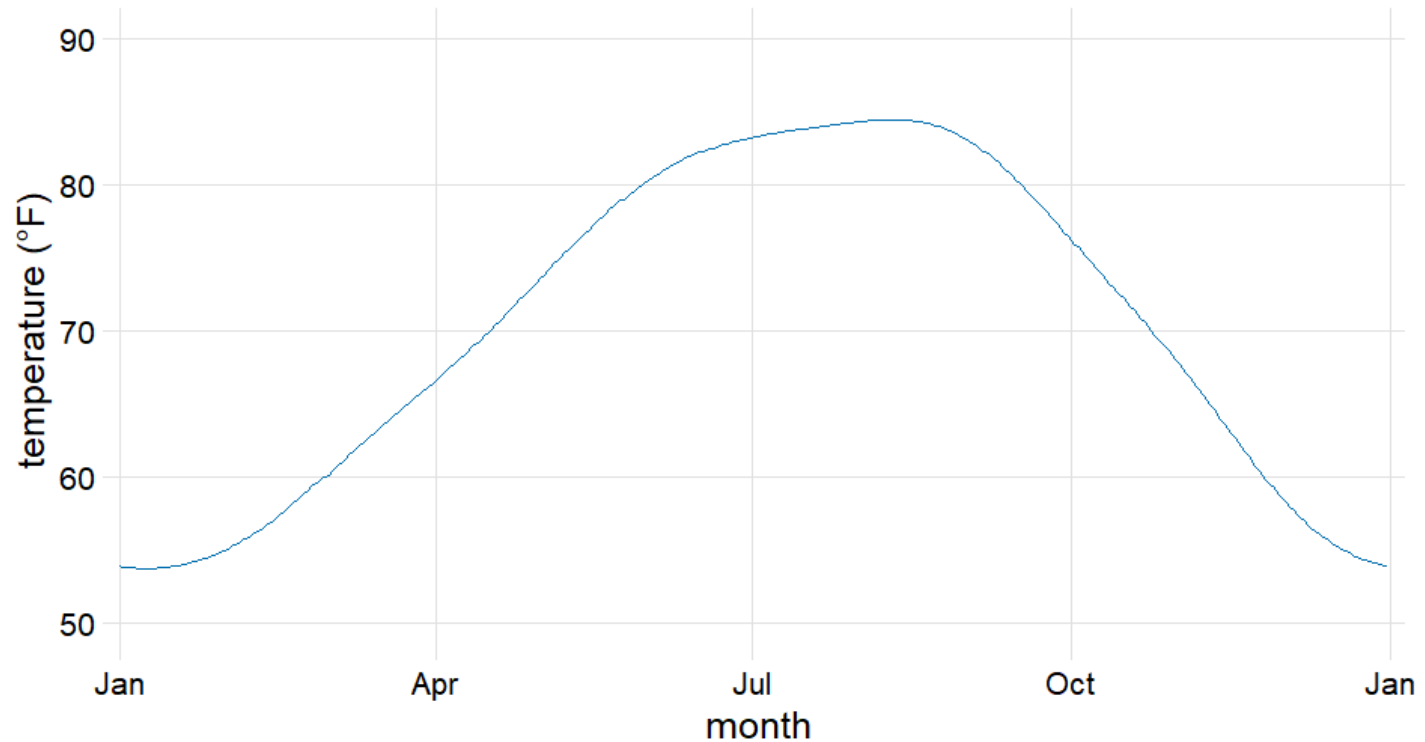
- Additional Considerations when visualizing/comparing categories including over time
- Understand how to make a wide variety of tweaks to ggplot.
- Understand common modifications to plots to make them more clear and reduce cognitive load
 - And ways to implement them

Axes

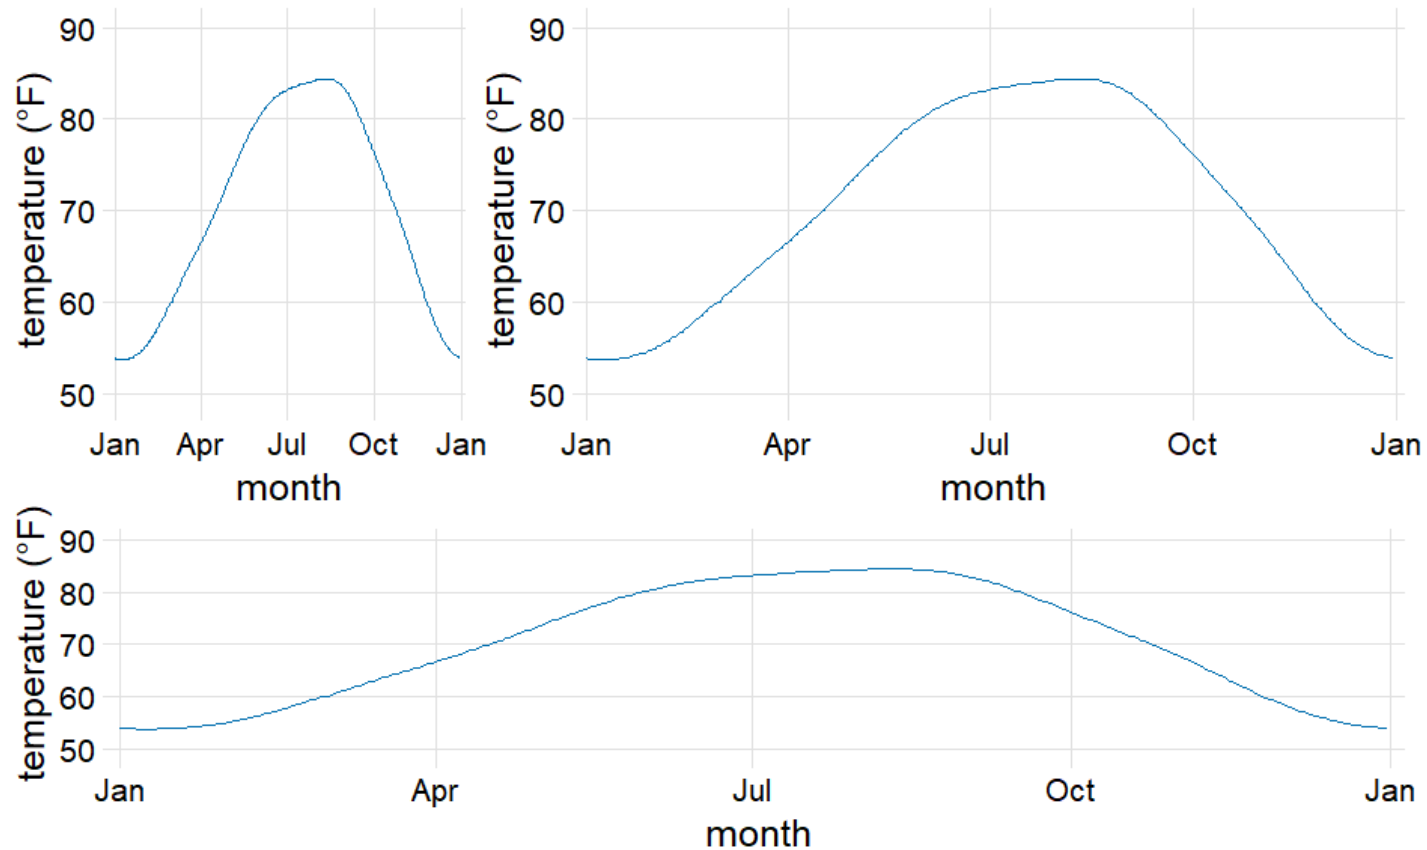
- Cartesian coordinates - what we generally use

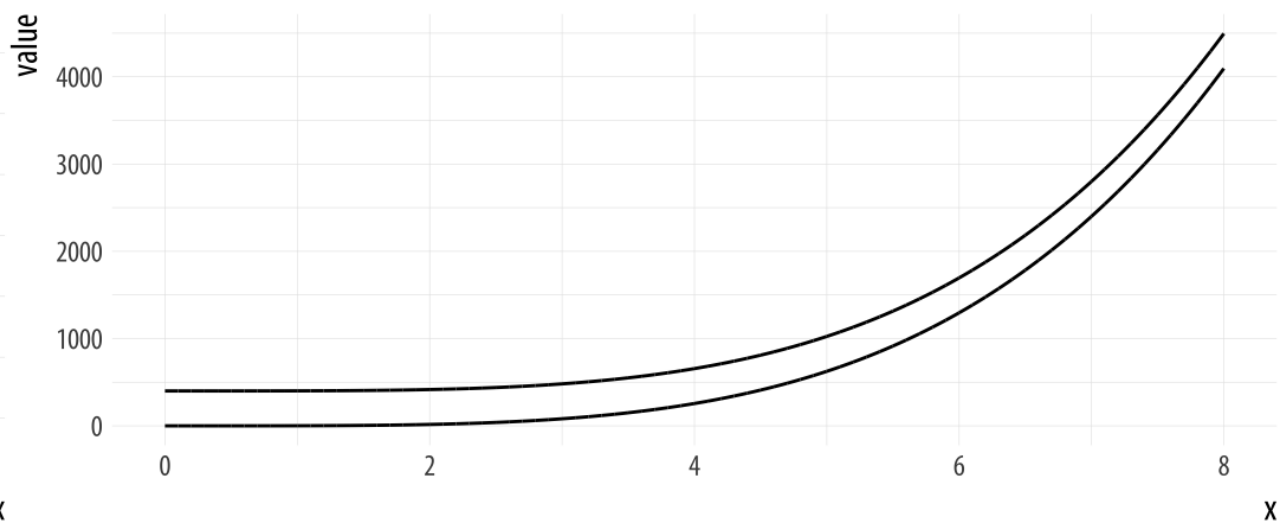
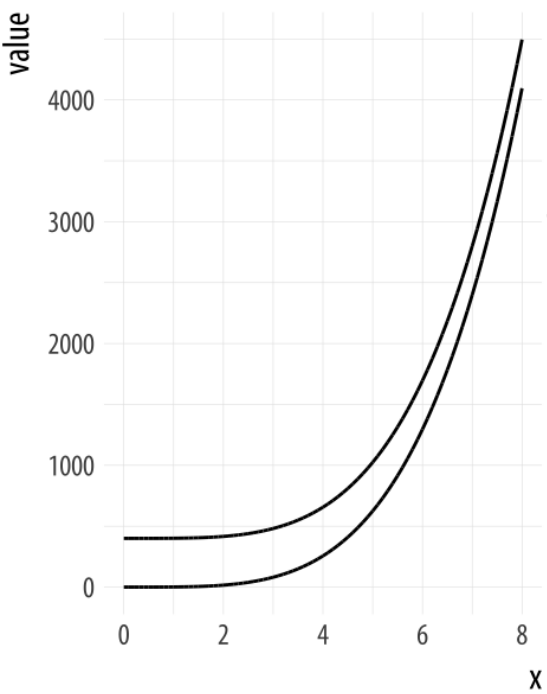


Different units



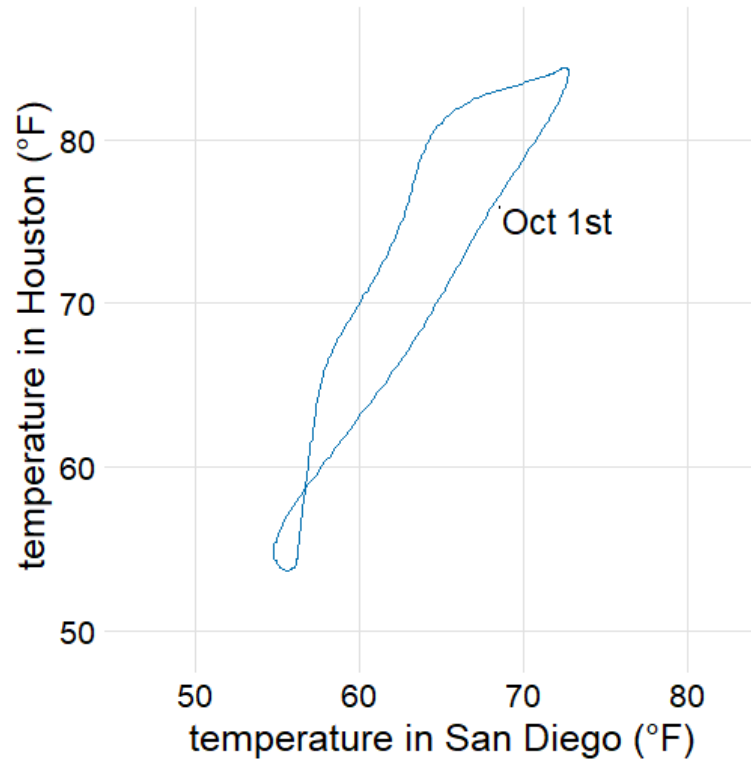
Aspect ratio





Same scales

Use `coord_fixed()`



Note - I think this is a weird plot, but the point remains

Changing aspect ratio

- Explore how your plot will look in its final size
- No hard/fast rules (if on different scales)
- Not even really rules of thumb
- Keep visual perception in mind
- Try your best to be truthful - show the trend/relation, but don't exaggerate/hide it

Handy function

(from an apparently deleted tweet from [@tjmahr](#))

here's my favorite helper [#rstats](#) function. preview
ggsave() output

```
ggpreview <- function (... , device = "png") {  
  fname <- tempfile(fileext = paste0(".", device))  
  ggplot2::ggsave(filename = fname, device =  
    device, ...)  
  system2("open", fname)  
  invisible(NULL)  
}
```

— tj mahr 🍕 🍍 (@tjmahr)

Gist

(side note: gists are a good way to share things)

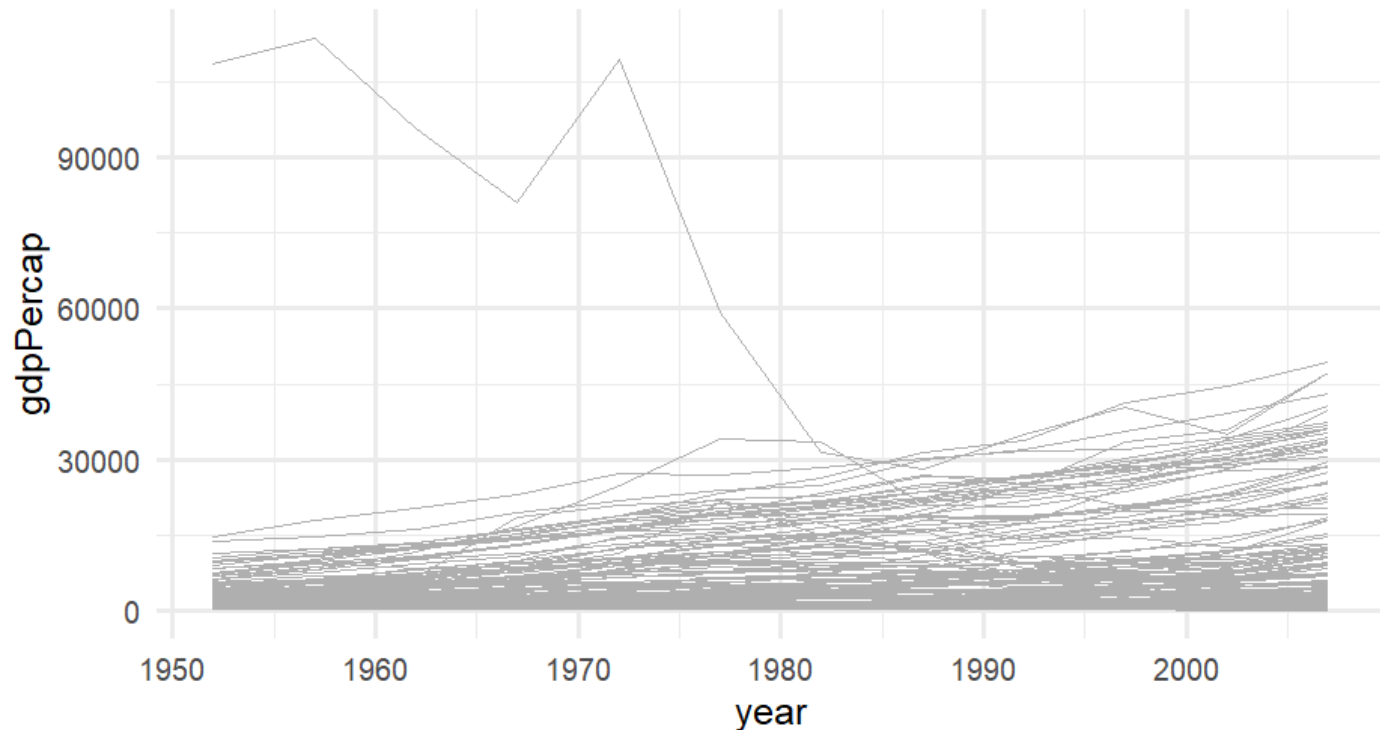
- See the full code/example [here](#)
- Let's take 3 minutes to play around:
 - Create a plot (could even be the example in the gist)
 - Try different aspect ratios by changing the width/length

03:00

Scale transformation

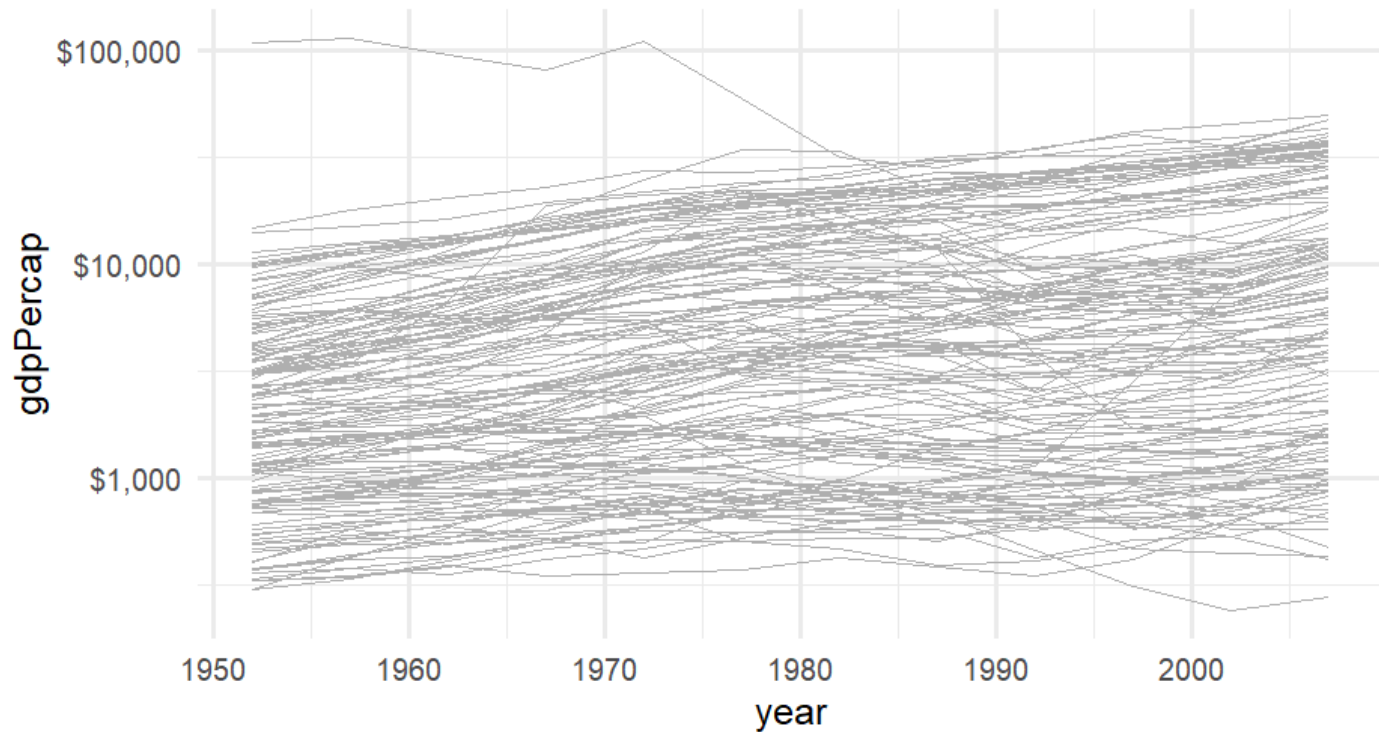
Raw scale

```
library(gapminder)
ggplot(gapminder, aes(year, gdpPercap)) +
  geom_line(aes(group = country),
            color = "gray70")
```

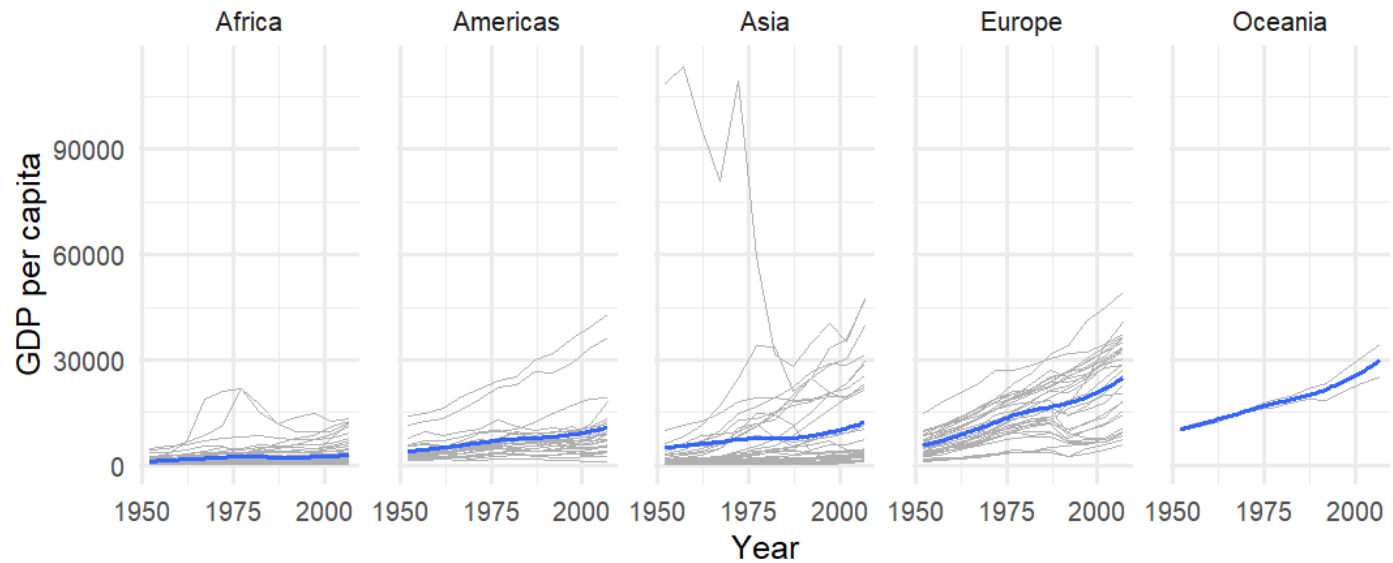


Log10 scale

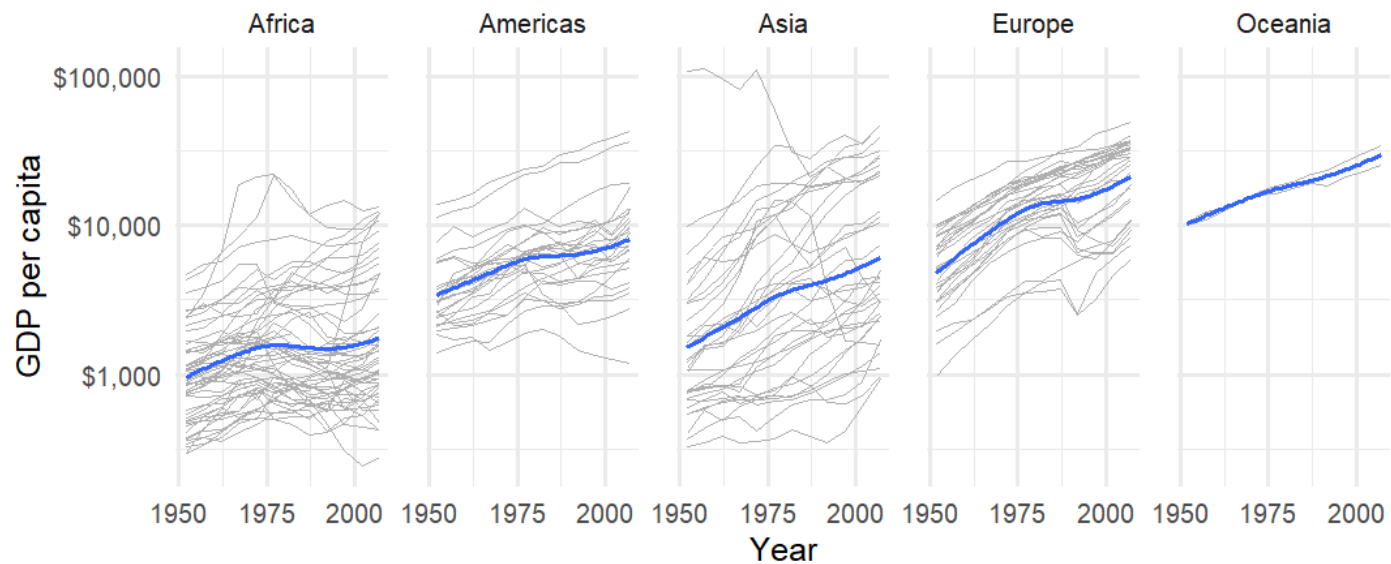
```
ggplot(gapminder, aes(year, gdpPercap)) +  
  geom_line(aes(group = country),  
            color = "gray70") +  
  scale_y_log10(labels = scales::dollar)
```



GDP per capita on Five Continents



GDP per capita on Five Continents



Scales

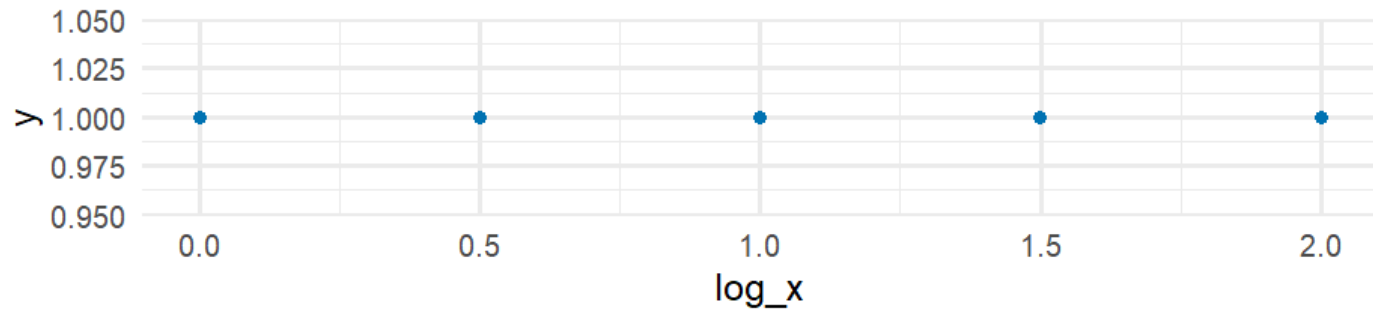
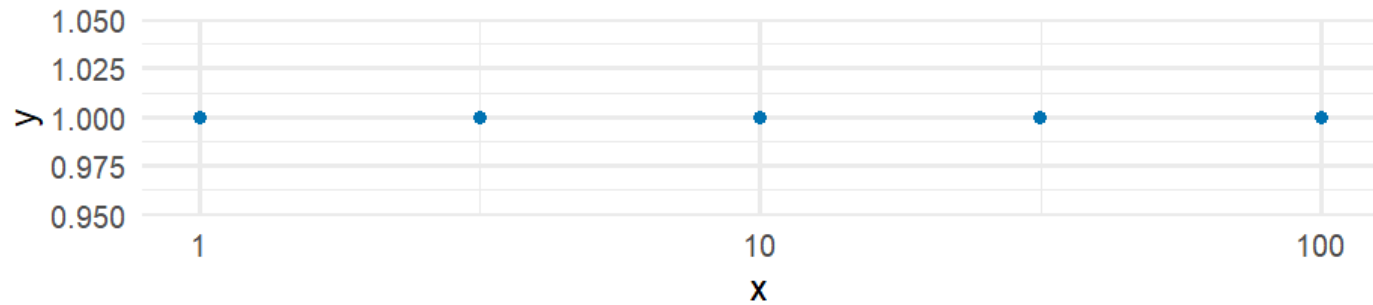
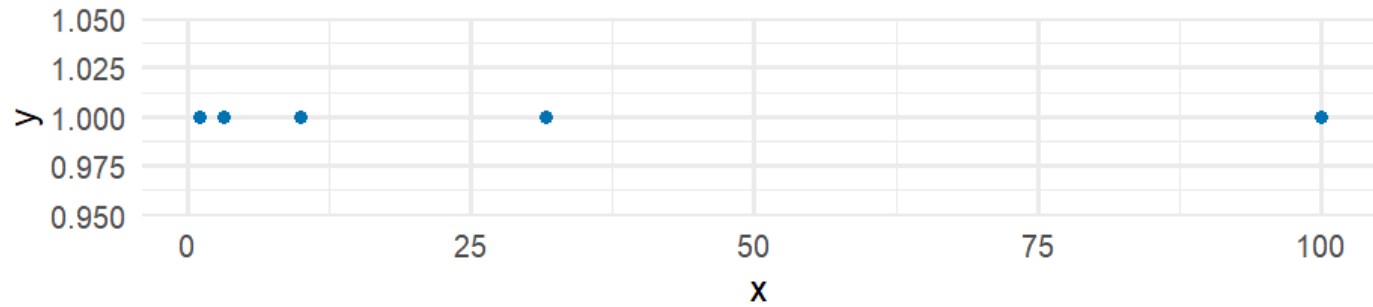
```
d <- tibble(x = c(1, 3.16, 10, 31.6, 100),
             log_x = log10(x))

ggplot(d, aes(x, 1)) +
  geom_point(color = "#0072B2")

ggplot(d, aes(x, 1)) +
  geom_point(color = "#0072B2") +
  scale_x_log10()

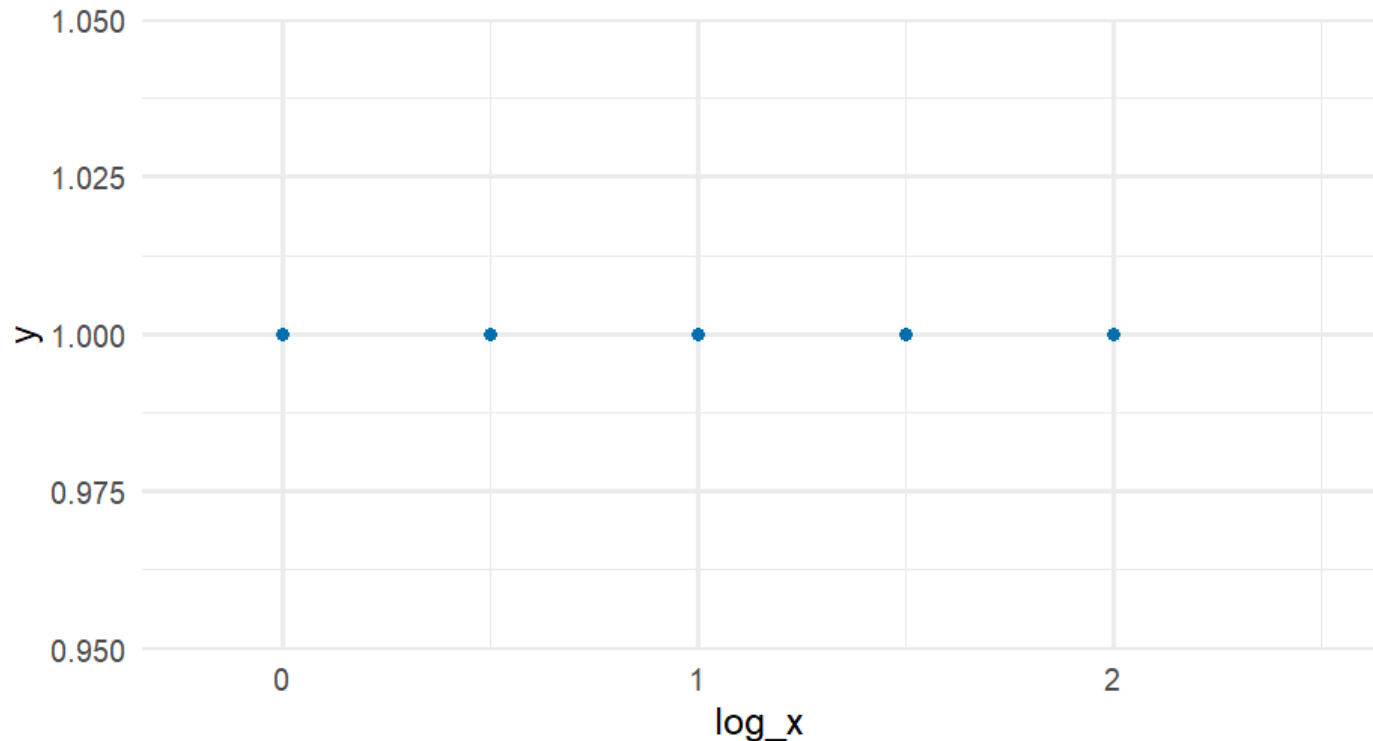
ggplot(d, aes(log_x, 1)) +
  geom_point(color = "#0072B2")
```

Scales



Don't transform twice

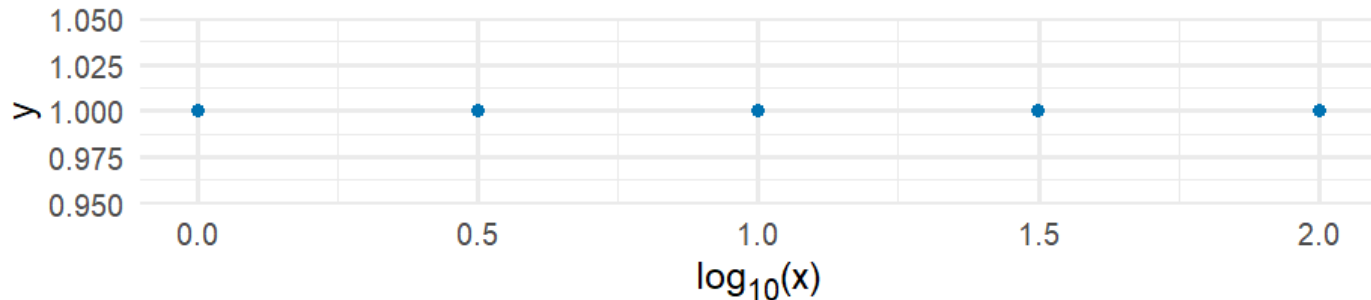
```
ggplot(d, aes(log_x, 1)) +  
  geom_point(color = "#0072B2") +  
  scale_x_log10() +  
  xlim(-0.2, 2.5)
```



Careful with labeling

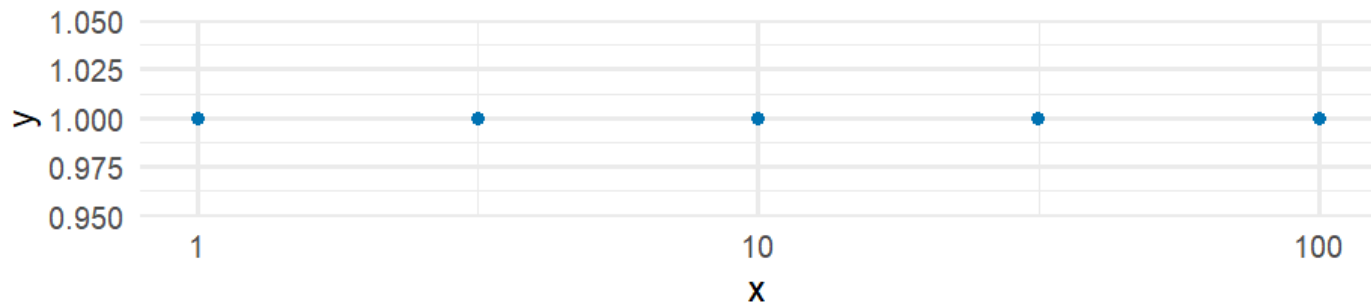
- Has the scale or the data been log transformed?
- Specify the base

```
library(ggtext)
ggplot(d, aes(log_x, 1)) +
  geom_point(color = "#0072B2") +
  labs(x = "log<sub>10</sub>(x)") +
  theme(axis.title.x = element_markdown())
```



Labels should denote the data, not the scale of the axis

```
ggplot(d, aes(x, 1)) +  
  geom_point(color = "#0072B2") +  
  scale_x_log10()
```



Labeling the above with $\log_{10}(x)$ would be ambiguous and confusing

Interpretation

Log scales show relative numbers, not raw

Difference between 100 and 101 (1% change) will be much smaller than difference between 1 and 2 (100% change)

More resources to learn more

- [Blog post by Lisa Charlotte Muth](#)
- The "Logarithmic or Linear scales" section [here](#)

Labels and captions

Disclaimer

- APA style requires the labels be made in specific ways
- Much of the following discussion still applies
- Our book (Wilke) uses a similar style throughout

Title

What is the point of your figure?

What are you trying to communicate

- Figures should have only one title
- Use integrated title/subtitles for sharing with a broad audience
 - Blog posts
 - Social media
 - Reports to stakeholders
- Make sure your figure has a title
 - Should not start with "This figure displays/shows..."

Caption

Consider stating the data source

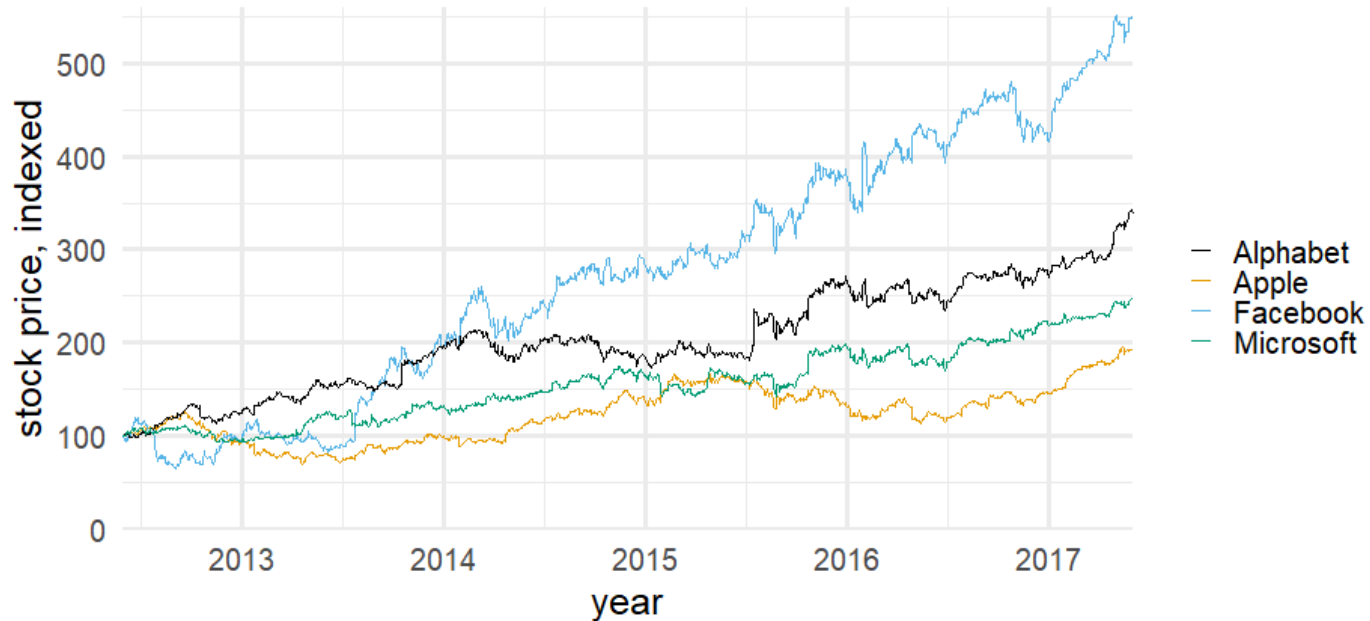
Other details relevant to the figure but not important enough for a subtitle

Axis labels

- The title for the axis
- Critical for communication
- **Never** use variable names (very common and very poor practice)
- State the measure and the unit (if quantitative)
 - e.g., "Brain Mass (grams)", "Support for Measure (millions of people)", "Dollars spent"
 - Categorical variable likely will not need to the measurement unit

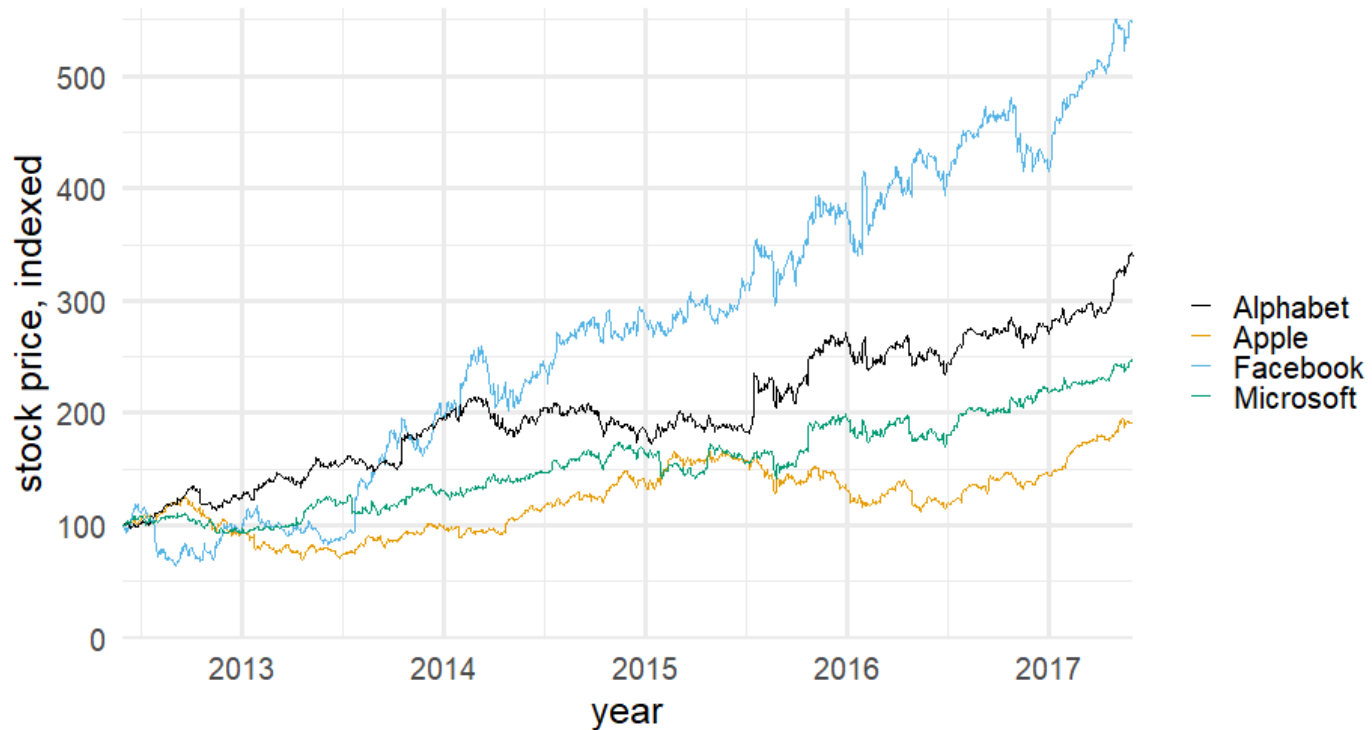
Omission

- Consider omitting obvious or redundant labels
 - Use `labs(x = NULL)` or `labs(x = "")`
 - If already using `scale_x/y_*` just supply the `name` argument

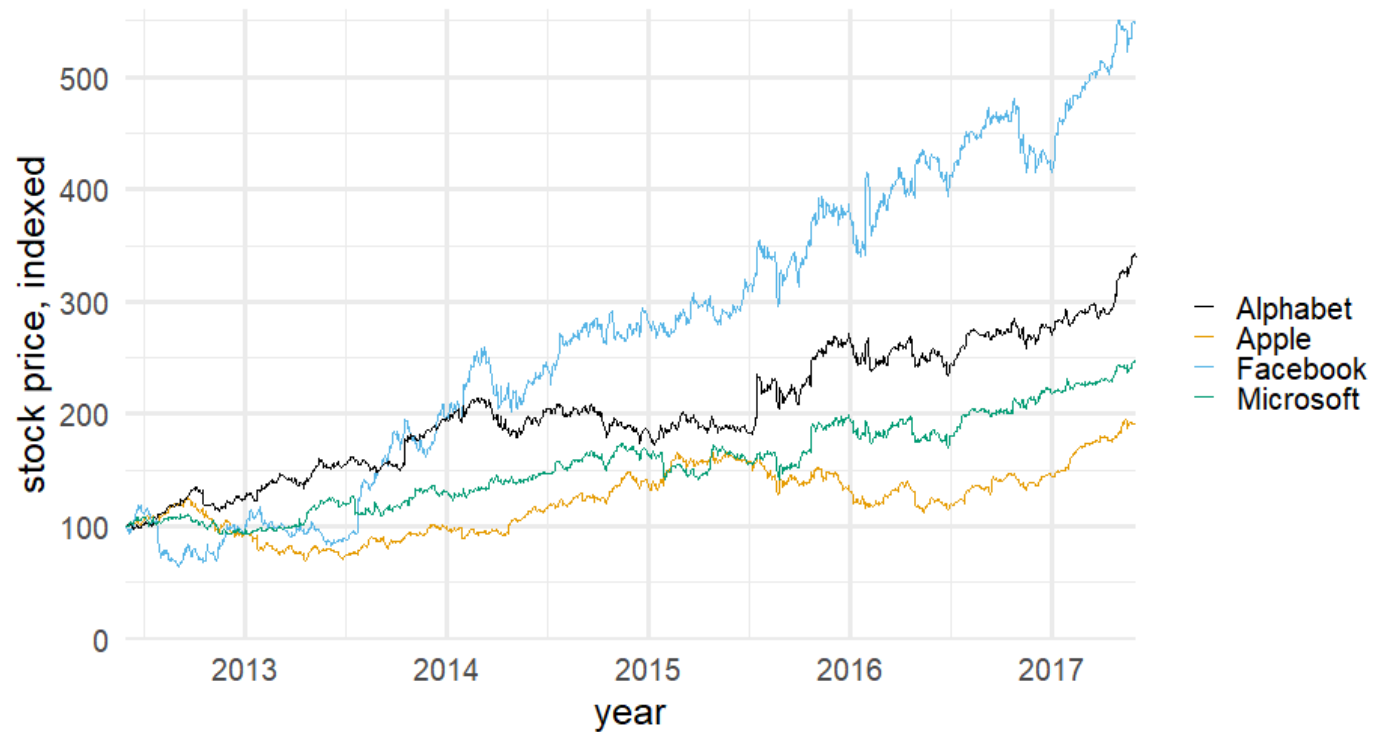


Omission

- Do not omit axis titles that are not obvious



Don't overdo it



Practice

Let's use the `ggplot2::diamonds` dataset.

- Plot the relationship between `carat` and `price`
- Play with scale transformations
- Give it some good labels
- Make any other modifications you'd like that you think makes it prettier and/or easier to interpret.

05:00

Annotations

The big topic for the day

Among the most effective

- If possible, try to remove legends, and just include annotations
- Warning - this is often fairly difficult in ggplot (requires a lot of fiddling)
- Consider saving and making final annotations outside of R
 - Bad for reproducibility, but good for interpretability
- There are some packages that can help

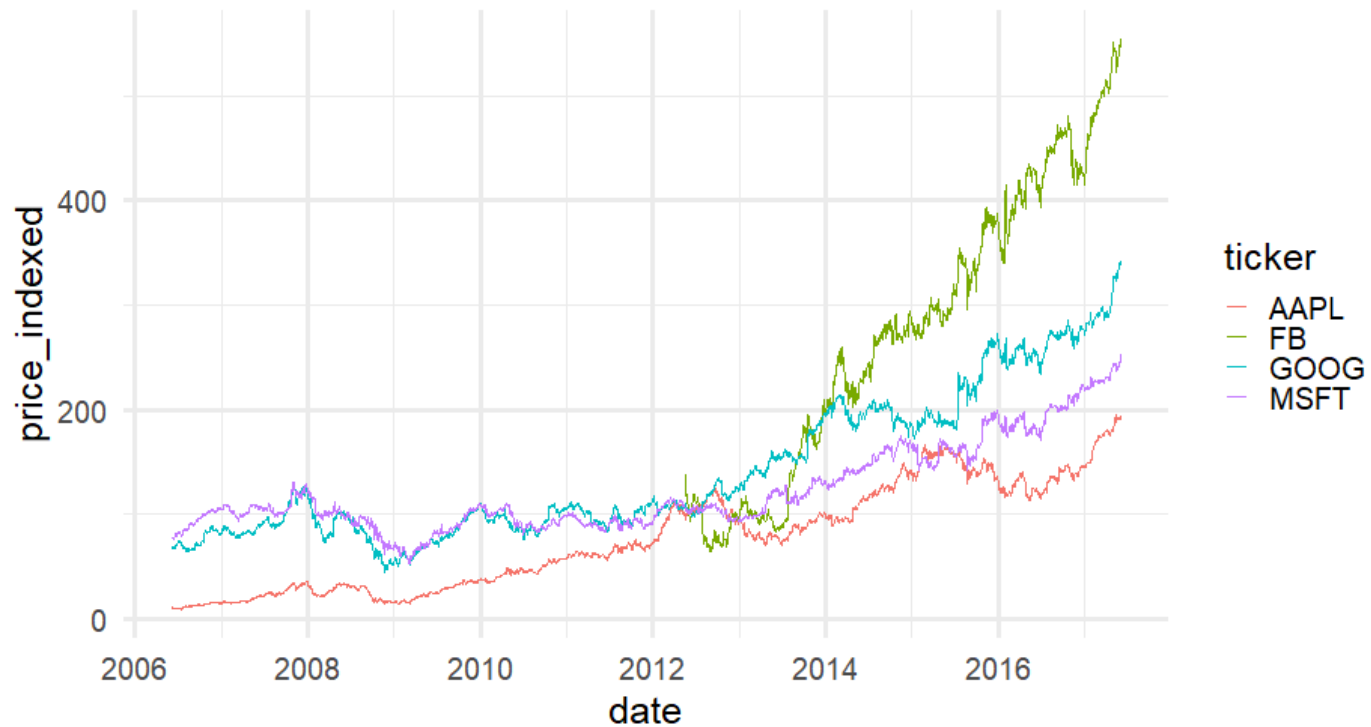
Building up a plot

```
remotes::install_github("clauswilke/dviz.supp")  
head(tech_stocks)
```

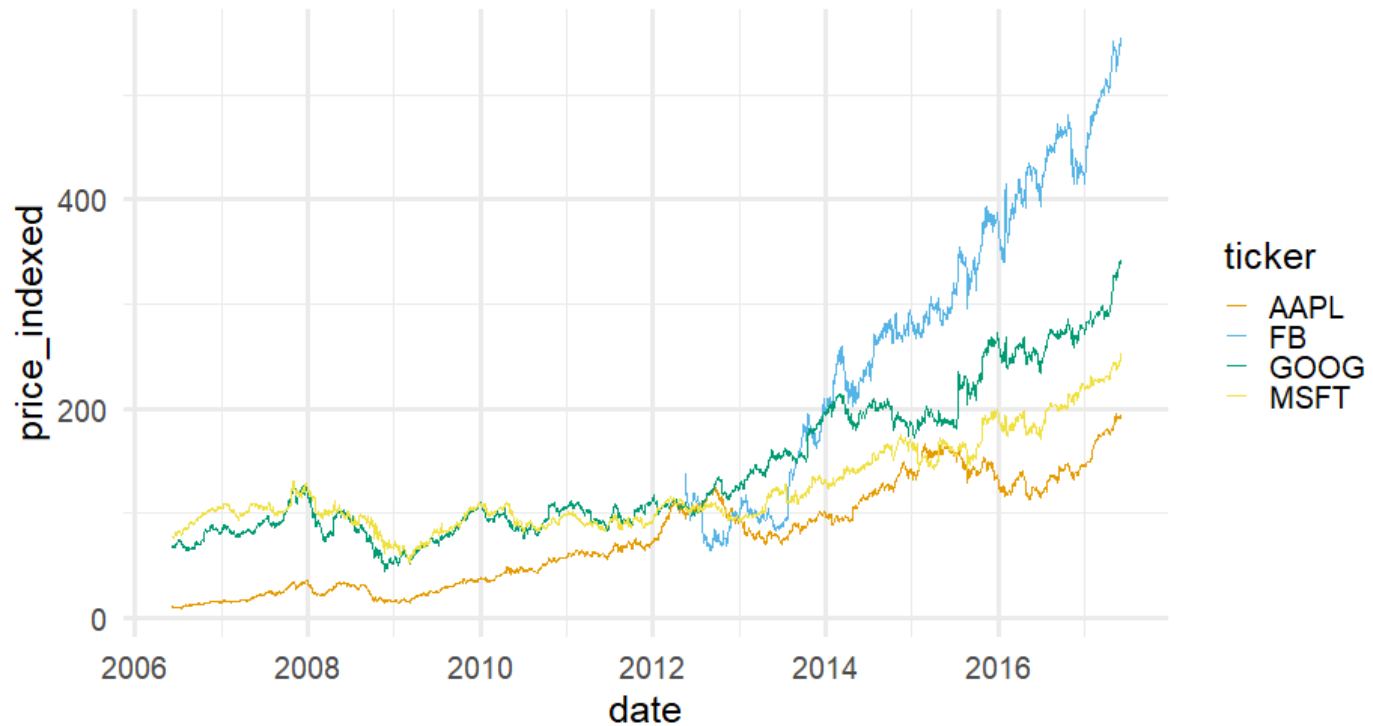
```
## # A tibble: 6 × 6
```

##	company	ticker	date	price	index_price	price_indexed
##	<chr>	<chr>	<date>	<dbl>	<dbl>	<dbl>
## 1	Alphabet	GOOG	2017-06-02	976.	285.	342.
## 2	Alphabet	GOOG	2017-06-01	967.	285.	339.
## 3	Alphabet	GOOG	2017-05-31	965.	285.	338.
## 4	Alphabet	GOOG	2017-05-30	976.	285.	342.
## 5	Alphabet	GOOG	2017-05-26	971.	285.	341.
## 6	Alphabet	GOOG	2017-05-25	970.	285.	340.

```
ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +  
  geom_line()
```

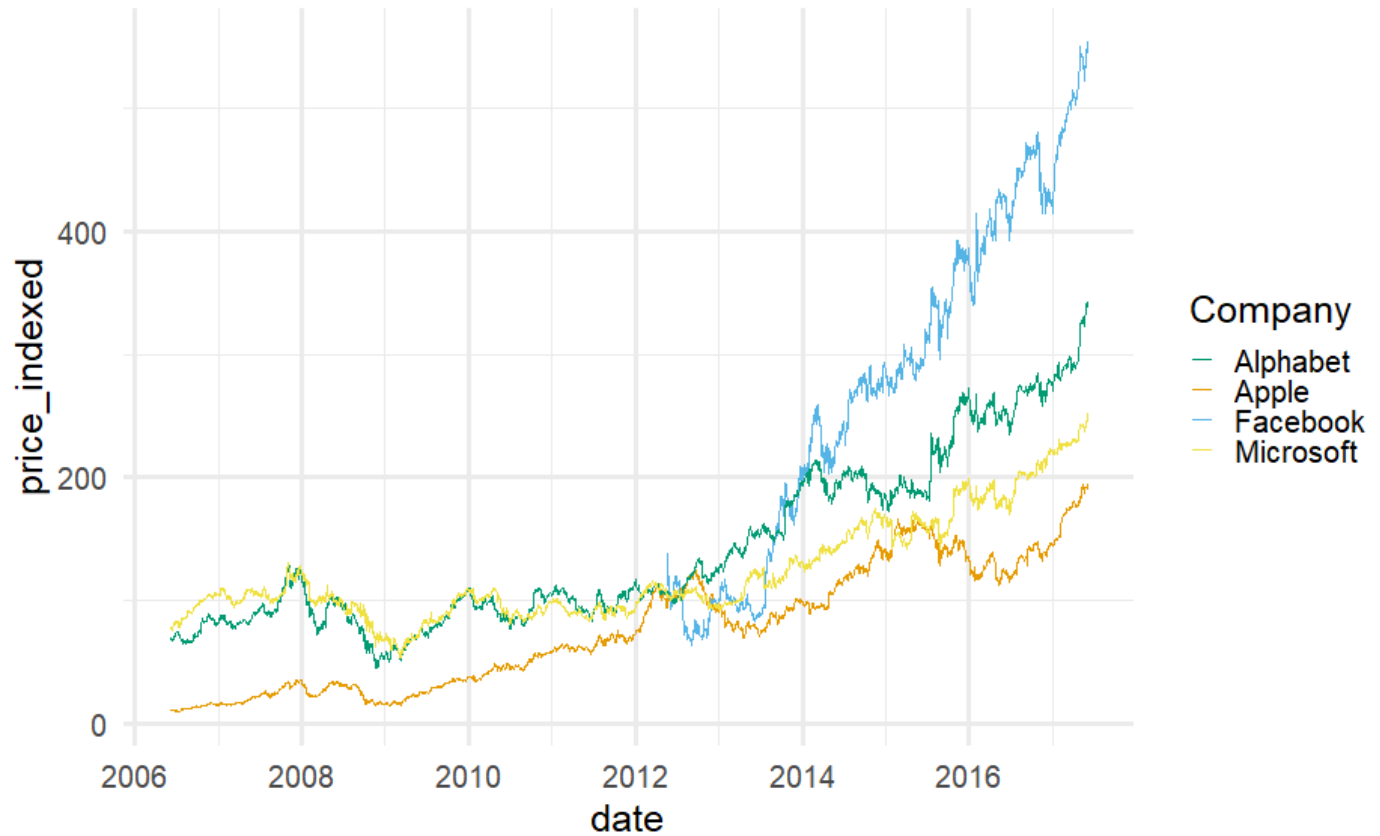


```
ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +  
  geom_line() +  
  scale_color_okabeIto()
```



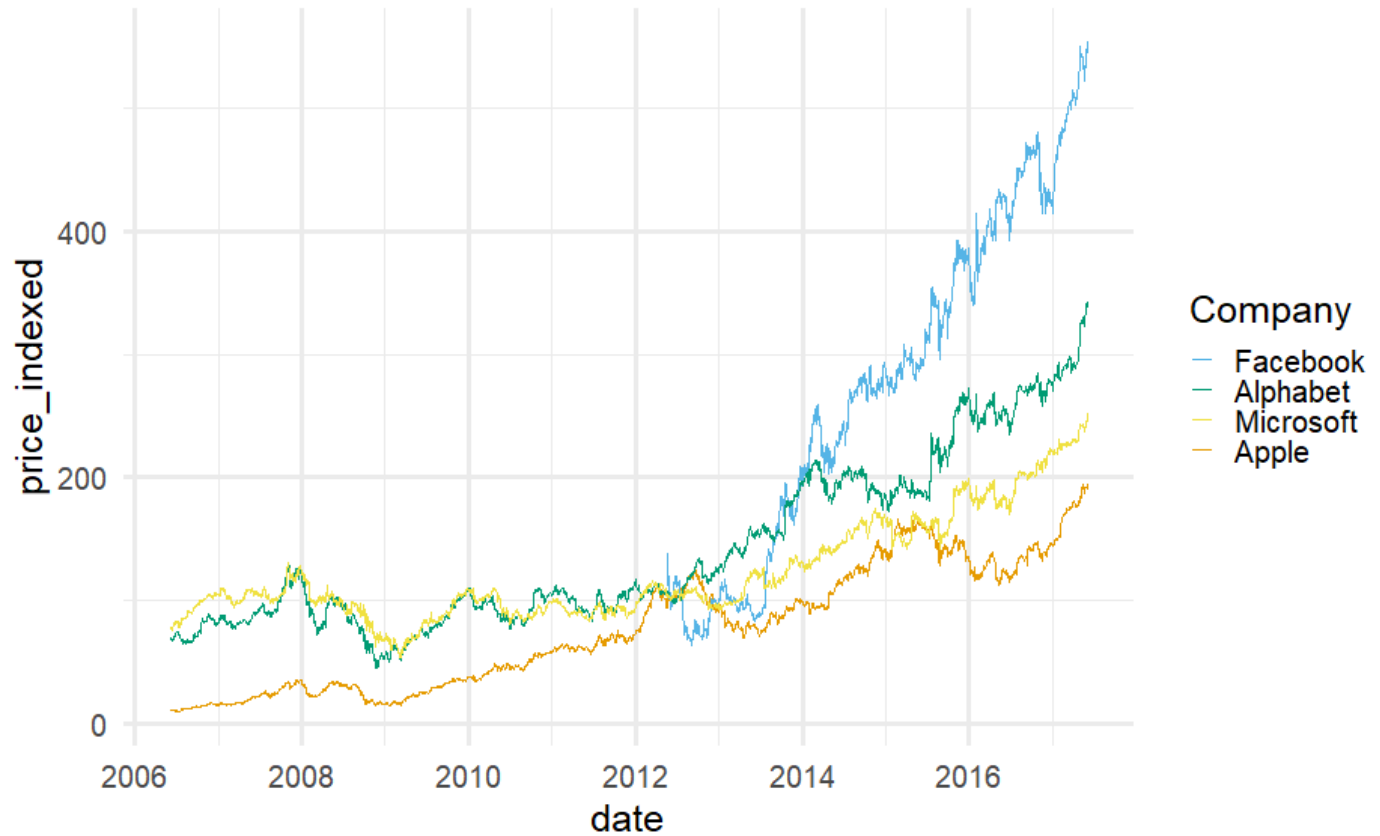
```
ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +  
  geom_line() +  
  scale_color_OkabeIto(  
    name = "Company",  
    breaks = c("GOOG", "AAPL", "FB", "MSFT"),  
    labels = c("Alphabet", "Apple", "Facebook", "Microsoft")  
  )
```

Bad

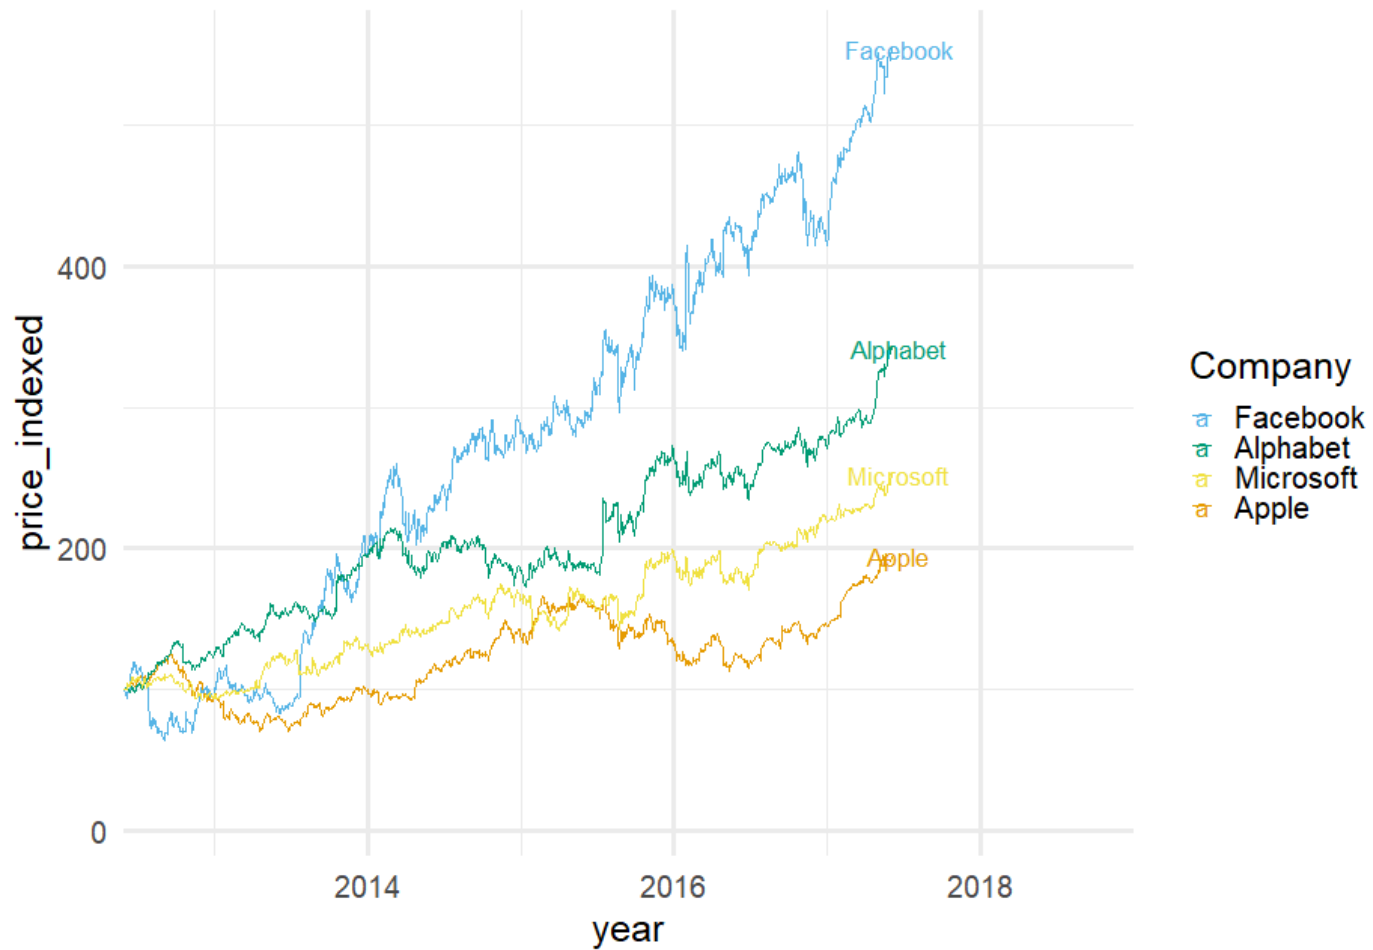



```
ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +  
  geom_line() +  
  scale_color_OkabeIto(  
    name = "Company",  
    breaks = c("FB", "GOOG", "MSFT", "AAPL"),  
    labels = c("Facebook", "Alphabet", "Microsoft", "Apple")  
  )
```

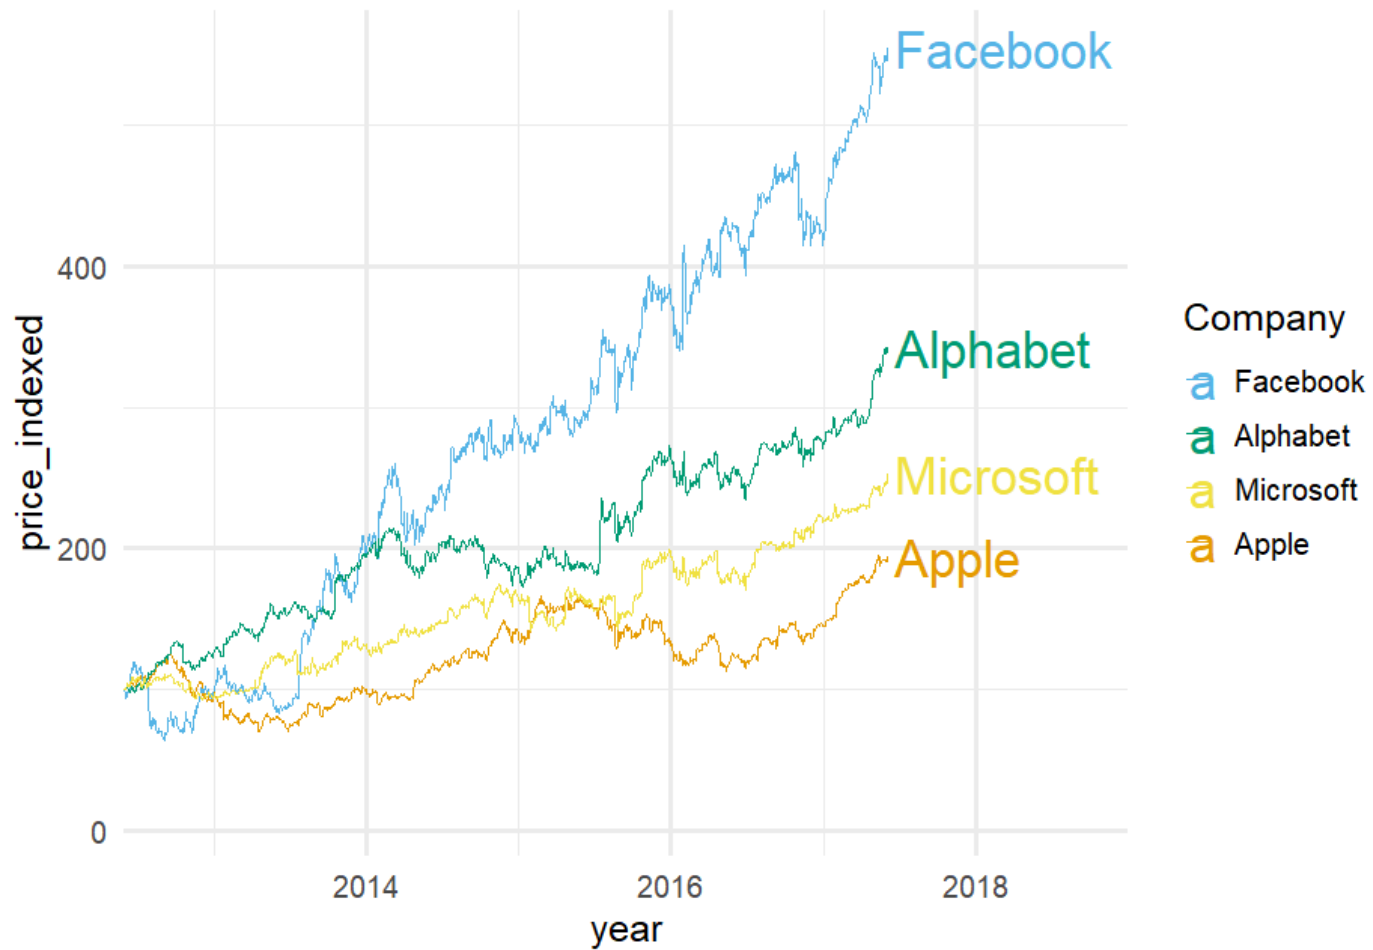
Good



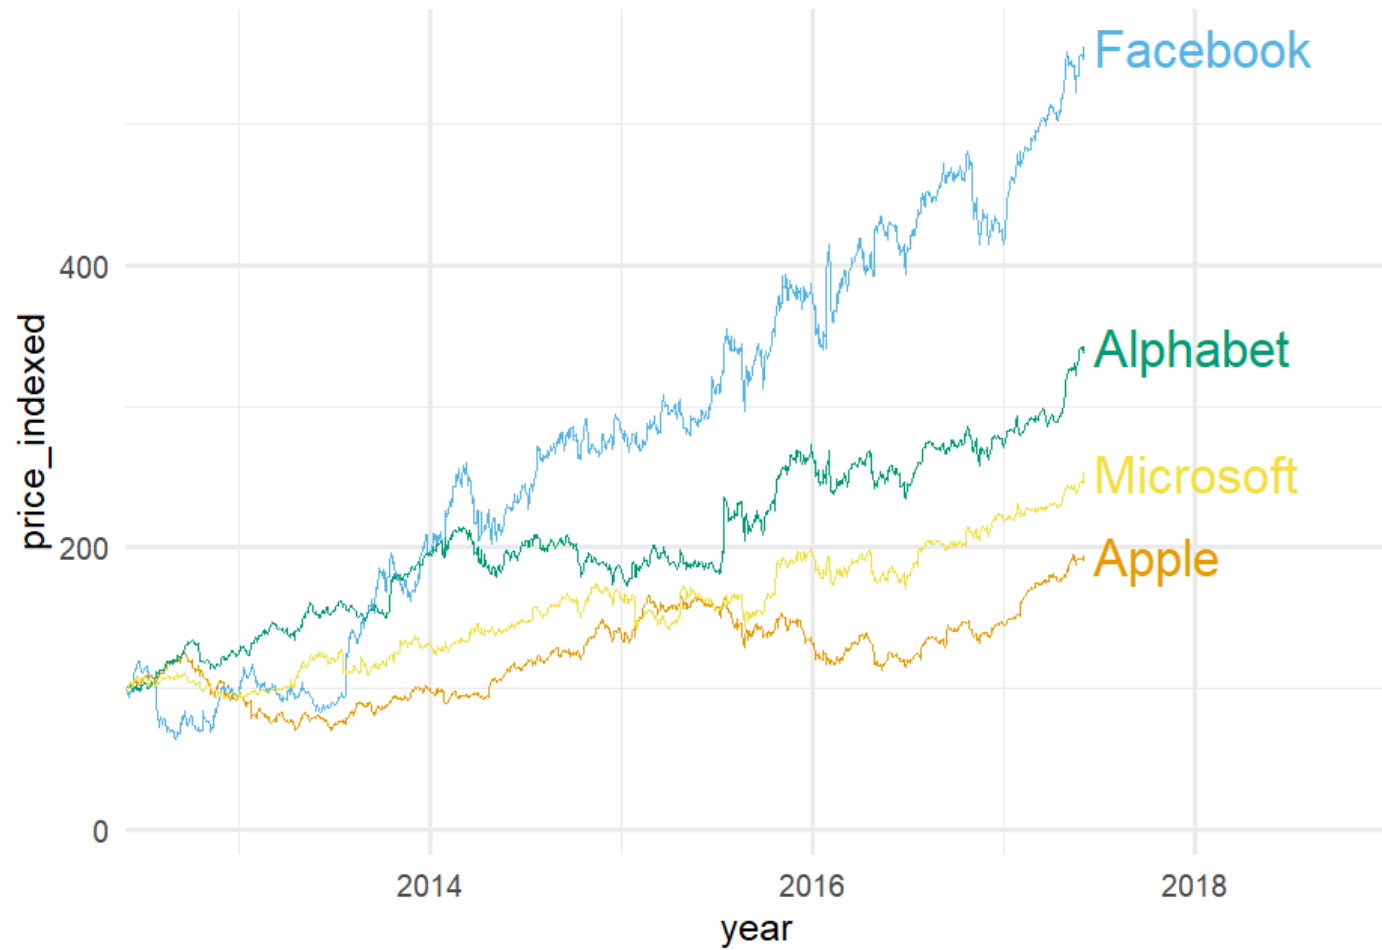
```
ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +  
  geom_line() +  
  scale_color_OkabeIto(  
    name = "Company",  
    breaks = c("FB", "GOOG", "MSFT", "AAPL"),  
    labels = c("Facebook", "Alphabet", "Microsoft", "Apple")  
  ) +  
  scale_x_date(  
    name = "year",  
    limits = c(ymd("2012-06-01"), ymd("2018-12-31")),  
    expand = c(0,0)  
  ) +  
  geom_text(  
    data = filter(tech_stocks, date == "2017-06-02"),  
    aes(y = price_indexed, label = company),  
    nudge_x = 20  
  )
```



```
ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +  
  geom_line() +  
  scale_color_OkabeIto(  
    name = "Company",  
    breaks = c("FB", "GOOG", "MSFT", "AAPL"),  
    labels = c("Facebook", "Alphabet", "Microsoft", "Apple")  
  ) +  
  scale_x_date(  
    name = "year",  
    limits = c(ymd("2012-06-01"), ymd("2018-12-31")),  
    expand = c(0,0)  
  ) +  
  geom_text(  
    data = filter(tech_stocks, date == "2017-06-02"),  
    aes(y = price_indexed, label = company),  
    nudge_x = 20,  
    hjust = 0,  
    size = 12  
  )
```



```
ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +  
  geom_line() +  
  scale_color_OkabeIto(  
    name = "Company",  
    breaks = c("FB", "GOOG", "MSFT", "AAPL"),  
    labels = c("Facebook", "Alphabet", "Microsoft", "Apple")  
  ) +  
  scale_x_date(  
    name = "year",  
    limits = c(ymd("2012-06-01"), ymd("2018-12-31")),  
    expand = c(0,0)  
  ) +  
  geom_text(  
    data = filter(tech_stocks, date == "2017-06-02"),  
    aes(y = price_indexed, label = company),  
    nudge_x = 20,  
    hjust = 0,  
    size = 12  
  ) +  
  guides(color = "none")
```

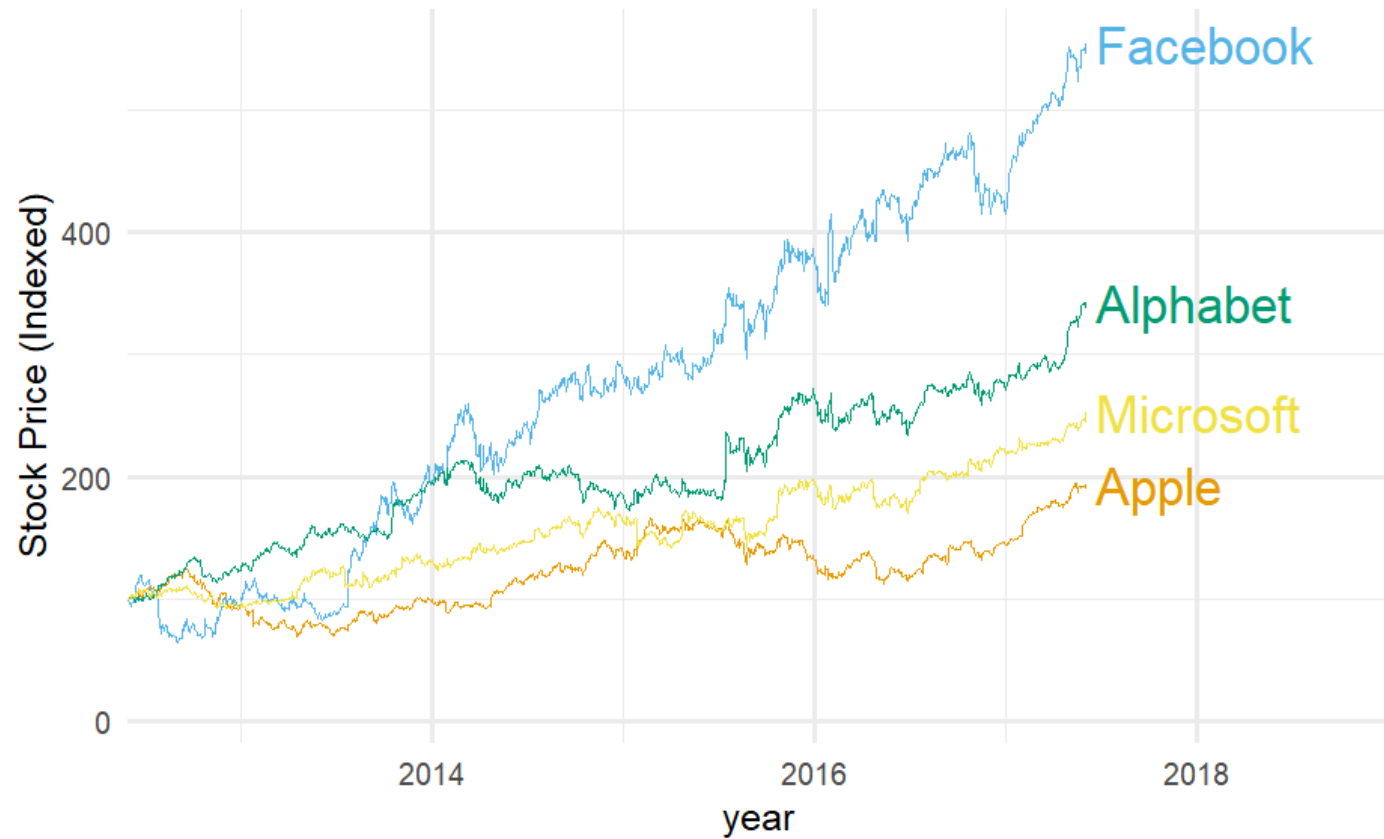



```

ggplot(tech_stocks, aes(date, price_indexed, color = ticker)) +
  geom_line() +
  scale_color_OkabeIto(
    name = "Company",
    breaks = c("FB", "GOOG", "MSFT", "AAPL"),
    labels = c("Facebook", "Alphabet", "Microsoft", "Apple")
  ) +
  scale_x_date(
    name = "year",
    limits = c(ymd("2012-06-01"), ymd("2018-12-31")),
    expand = c(0,0)
  ) +
  geom_text(
    data = filter(tech_stocks, date == "2017-06-02"),
    aes(y = price_indexed, label = company),
    nudge_x = 20,
    hjust = 0,
    size = 12
  ) +
  guides(color = "none") +
  labs(
    title = "Tech growth over time",
    caption = "Data from Wilke (2019): Fundamentals of Data Visualiza
    y = "Stock Price (Indexed)"
  )

```

Tech growth over time



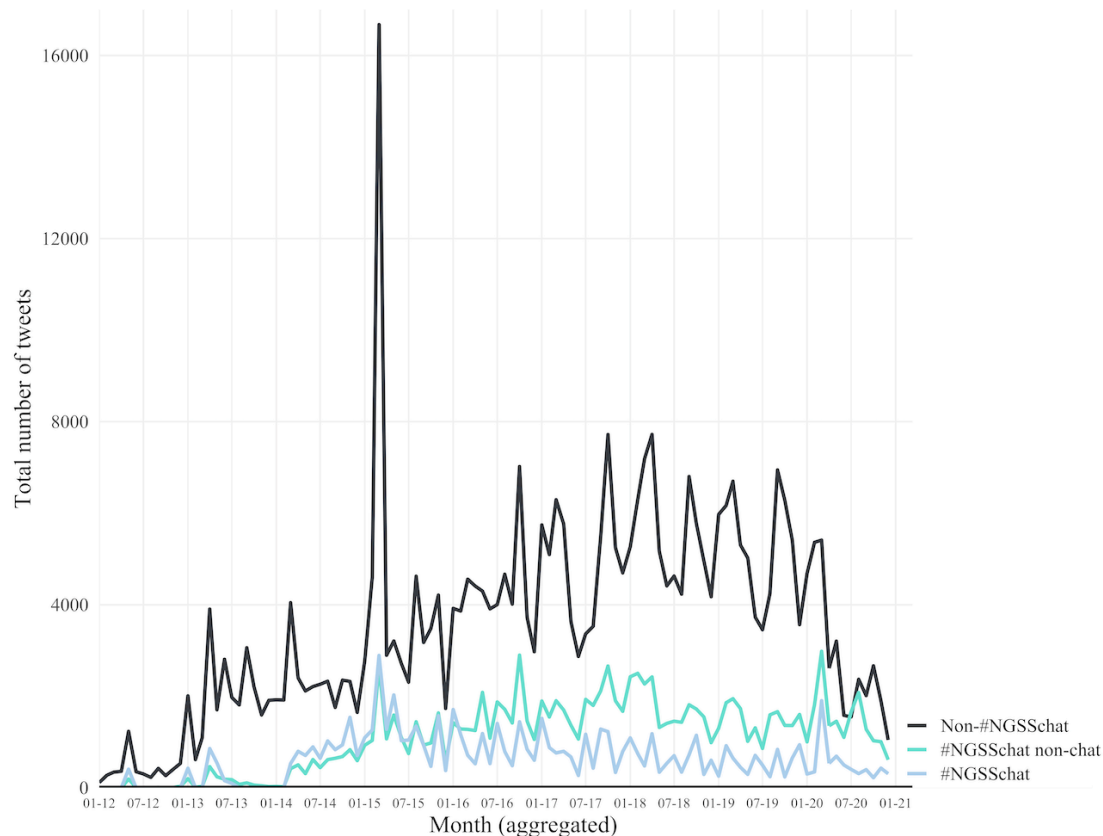
Data from Wilke (2019): Fundamentals of Data Visualization

A few more notes

- You might want to try `geom_label()` instead of `geom_text()`, or perhaps layering them with the first providing the white space for the second (as we saw in Lab 2)
- Could consider not making the font color vary with the lines (the labels are close enough)
- Depending on you you use the legend, it can work almost as well.

Example

From an actual publication, where I used the legend instead of direct annotations



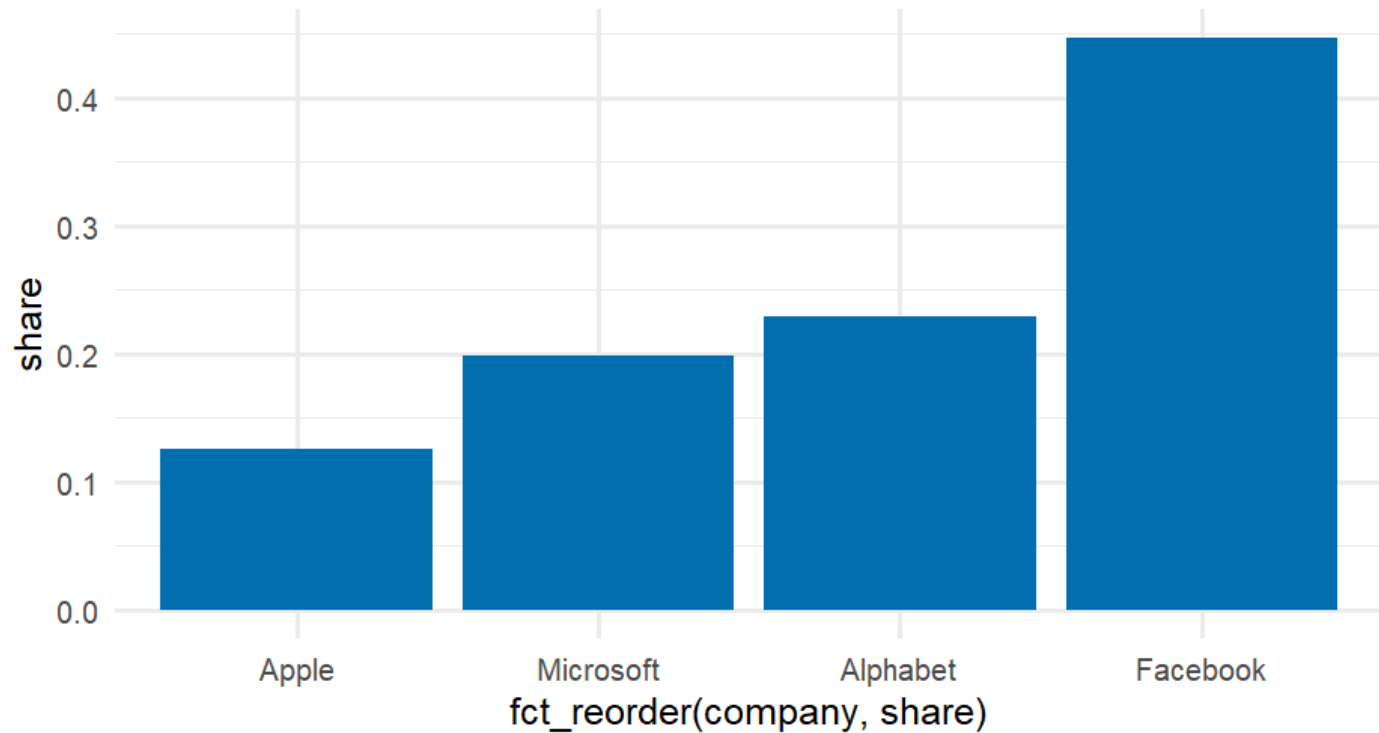
Labeling bars

```
avs <- tech_stocks %>%  
  group_by(company) %>%  
  summarize(stock_av = mean(price_indexed)) %>%  
  ungroup() %>%  
  mutate(share = stock_av / sum(stock_av))  
avs
```

```
## # A tibble: 4 × 3  
##   company    stock_av share  
##   <chr>      <dbl> <dbl>  
## 1 Alphabet    141.  0.229  
## 2 Apple       77.1  0.125  
## 3 Facebook    275.  0.447  
## 4 Microsoft   122.  0.199
```

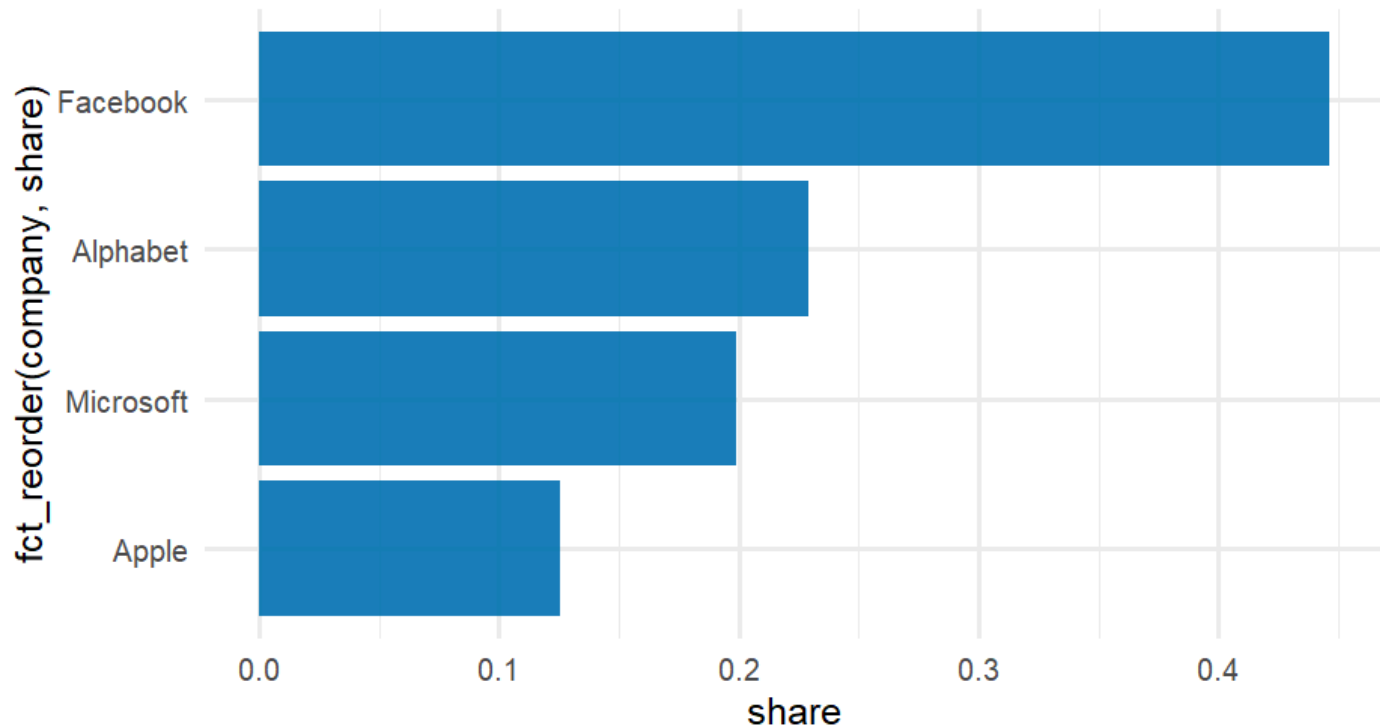
Bar plot

```
ggplot(avs, aes(fct_reorder(company, share), share)) +  
  geom_col(fill = "#0072B2")
```

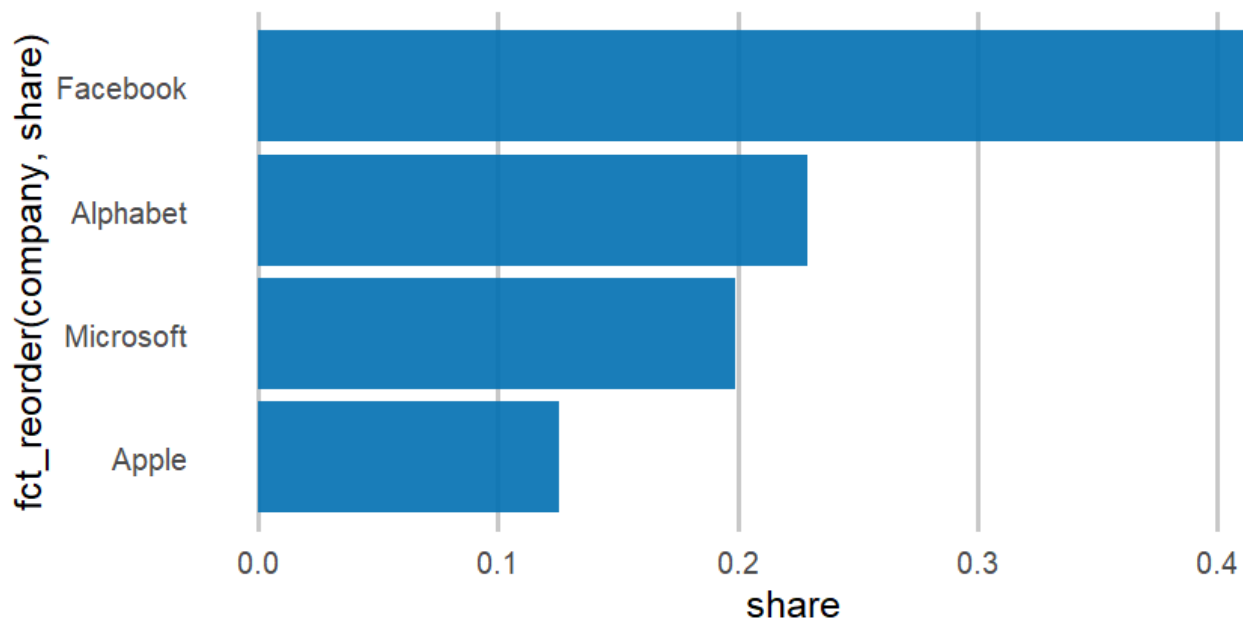


Horizontal

```
ggplot(avs, aes(share, fct_reorder(company, share))) +  
  geom_col(fill = "#0072B2", alpha = 0.9)
```



```
ggplot(avs, aes(share, fct_reorder(company, share))) +  
  geom_col(fill = "#0072B2", alpha = 0.9) +  
  theme(  
    panel.grid.major.y = element_blank(),  
    panel.grid.minor.x = element_blank(),  
    panel.grid.major.x = element_line(color = "gray80")  
  )
```



Quick aside

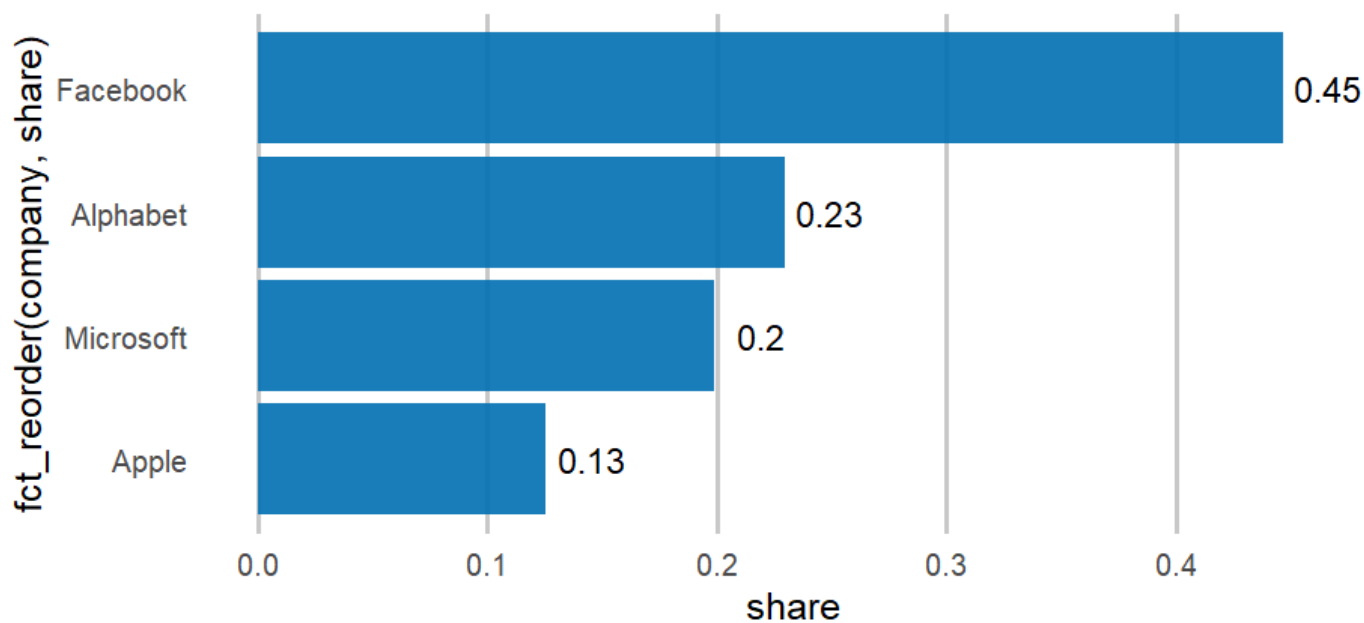
Let's actually make a bar plot theme

```
bp_theme <- function(...) {  
  theme_minimal(...) +  
    theme(  
      panel.grid.major.y = element_blank(),  
      panel.grid.minor.x = element_blank(),  
      panel.grid.major.x = element_line(color = "gray80"),  
      plot.title.position = "plot"  
    )  
}
```

```

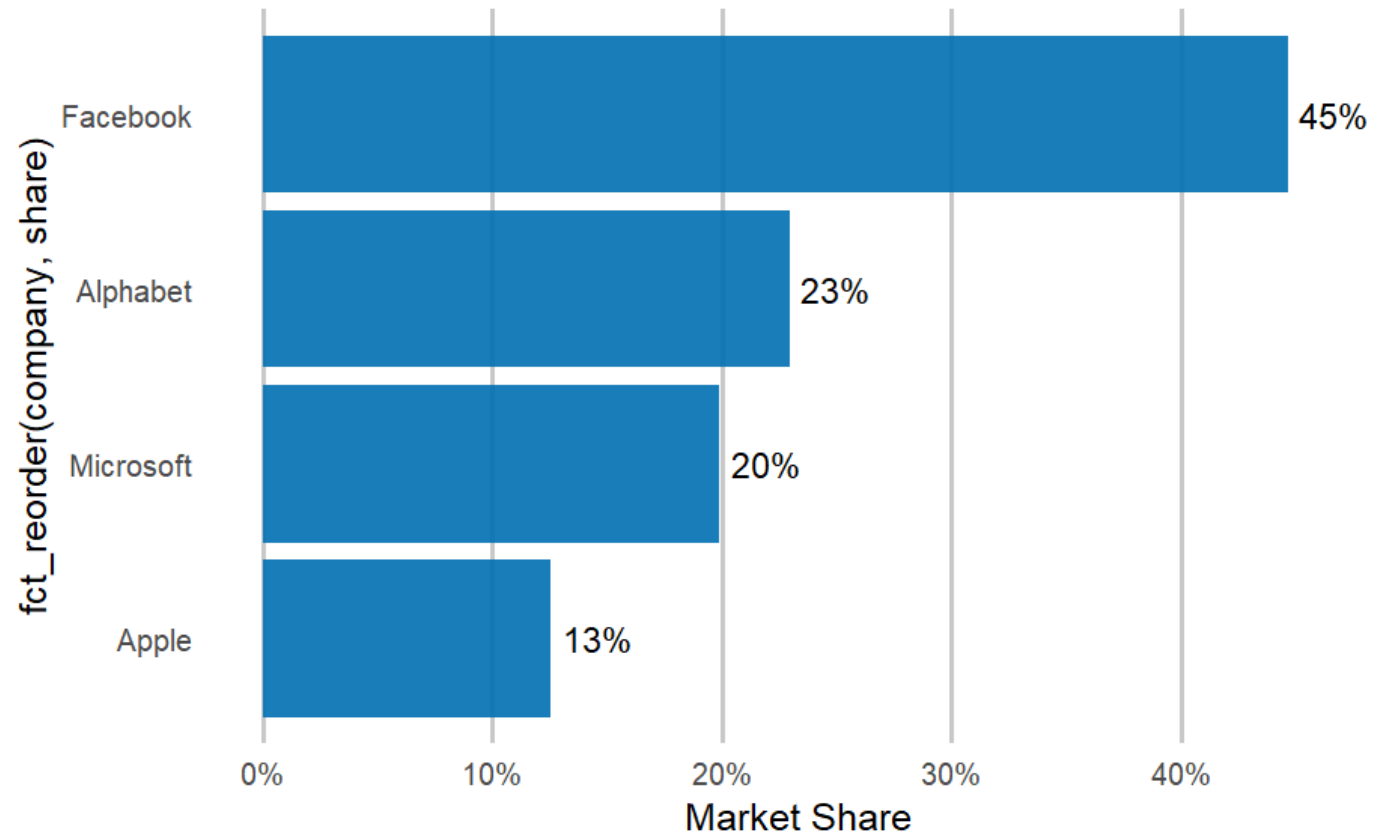
ggplot(avs, aes(share, fct_reorder(company, share))) +
  geom_col(fill = "#0072B2", alpha = 0.9) +
  geom_text(
    aes(share, company, label = round(share, 2)),
    nudge_x = 0.02,
    size = 8
  ) +
  bp_theme(base_size = 25)

```



```
ggplot(avs, aes(share, fct_reorder(company, share))) +  
  geom_col(fill = "#0072B2", alpha = 0.9) +  
  geom_text(  
    aes(share, company, label = paste0(round(share*100), "%")),  
    nudge_x = 0.02,  
    size = 8  
  ) +  
  scale_x_continuous(  
    name = "Market Share",  
    labels = scales::percent  
  ) +  
  labs(  
    x = NULL,  
    title = "Tech company market control",  
    caption = "Data from Clause Wilke Book: Fundamentals of Data  
  ) +  
  bp_theme(base_size = 25)
```

Tech company market control

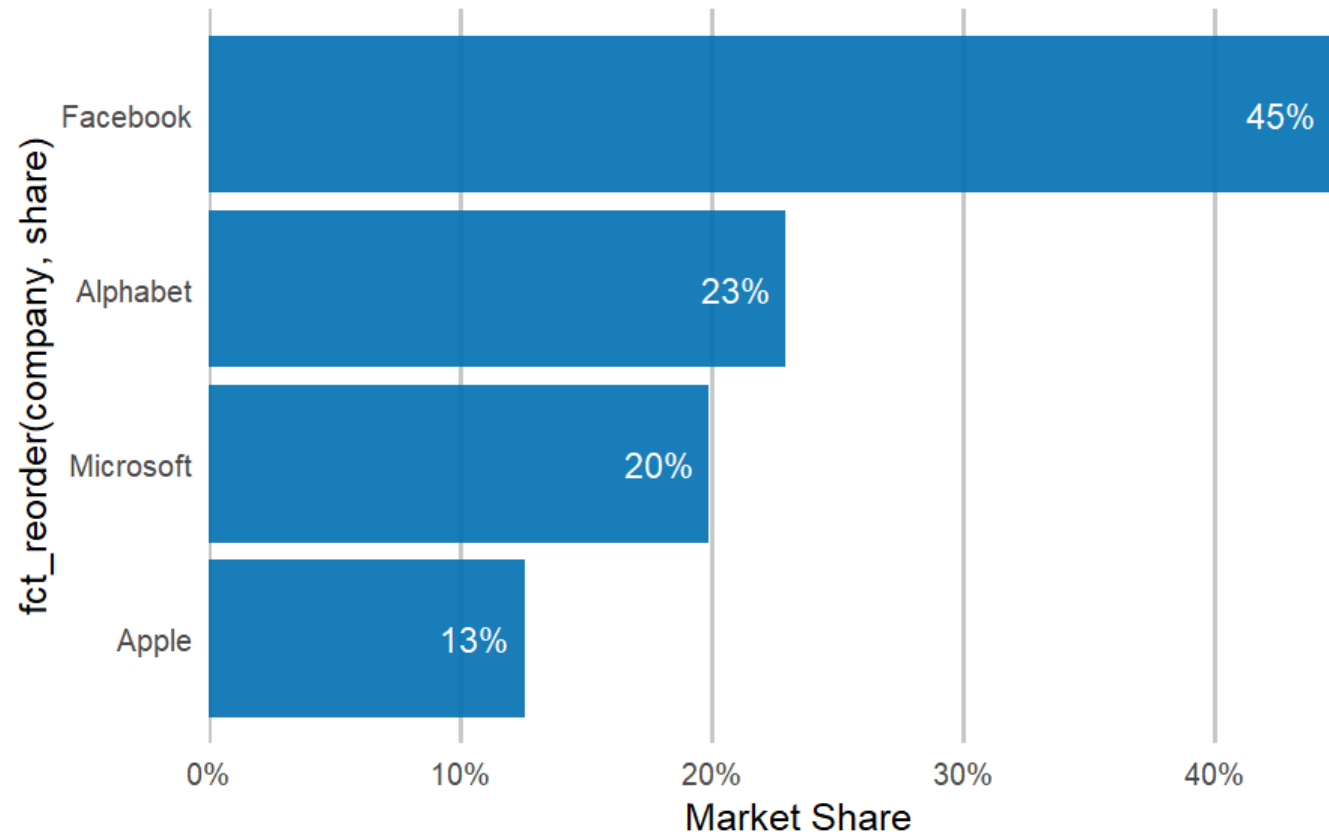


Data from Clause Wilke Book: Fundamentals of Data Visualizations

Another alternative

```
ggplot(avs, aes(share, fct_reorder(company, share))) +  
  geom_col(fill = "#0072B2", alpha = 0.9) +  
  geom_text(  
    aes(share, company, label = paste0(round(share*100), "%")),  
    nudge_x = -0.02,  
    size = 8,  
    color = "white"  
  ) +  
  scale_x_continuous(  
    "Market Share",  
    labels = scales::percent,  
    expand = c(0, 0, 0.05, 0)  
  ) +  
  labs(  
    x = NULL,  
    title = "Tech company market control",  
    caption = "Data from Clause Wilke Book: Fundamentals of Data  
  ) +  
  bp_theme(base_size = 25)
```

Tech company market control



Data from Clause Wilke Book: Fundamentals of Data Visualizations

Last example

This is a bit artificial in this case, but...

It is very common to have small bars. You may want most labels inside, but some outside

First, create variables specifying what you want.

Here I'm using 0.2 as my cutoff for whether the label is inside the bar or outside

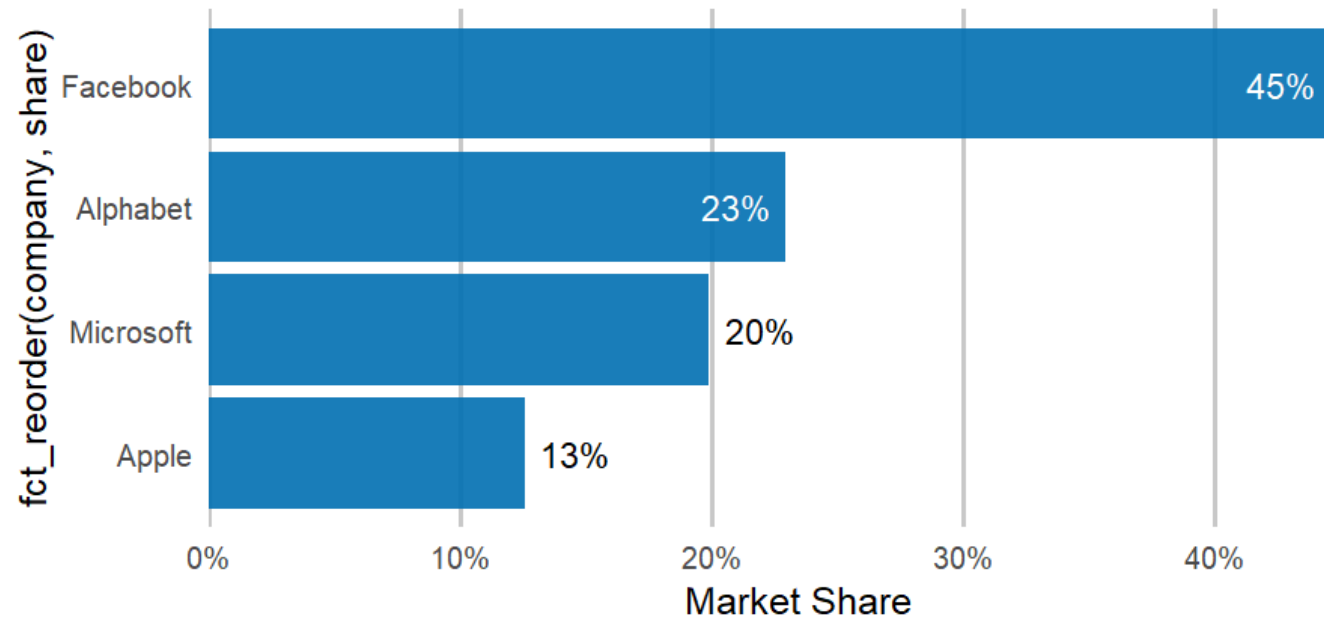
```
avs <- avs %>%  
  mutate(  
    nudge_amount = ifelse(share < 0.2, 0.02, -0.02),  
    label_color = ifelse(share < 0.2, "black", "white")  
  )  
avs
```

```
## # A tibble: 4 × 5  
##   company    stock_av share nudge_amount label_color  
##   <chr>      <dbl> <dbl>      <dbl> <chr>  
## 1 Alphabet    141.  0.229      -0.02 white  
## 2 Apple        77.1  0.125       0.02 black  
## 3 Facebook    275.  0.447      -0.02 white  
## 4 Microsoft   122.  0.199       0.02 black
```


nudge_* doesn't work inside aes 🙋

```
ggplot(avs, aes(share, fct_reorder(company, share))) +  
  geom_col(fill = "#0072B2", alpha = 0.9) +  
  geom_text(  
    aes(  
      share,  
      company,  
      label = paste0(round(share*100), "%"),  
      color = label_color  
    ),  
    nudge_x = avs$nudge_amount,  
    size = 8,  
  ) +  
  scale_x_continuous(  
    "Market Share",  
    labels = scales::percent,  
    expand = c(0, 0, 0.05, 0)  
  ) +  
  scale_color_identity() +  
  labs(  
    x = NULL,  
    title = "Tech company market control",  
    caption = "Data from Clause Wilke Book: Fundamentals of Data  
  ) +  
  bp_theme(base_size = 25)
```

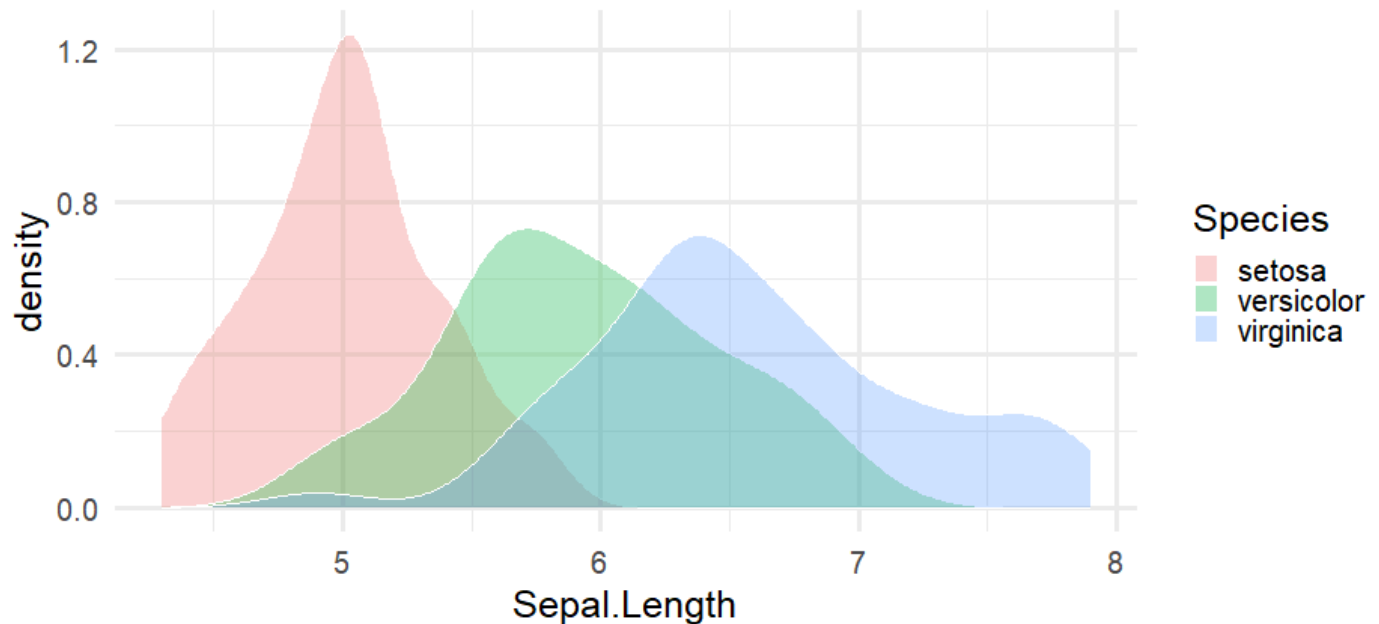
Tech company market control



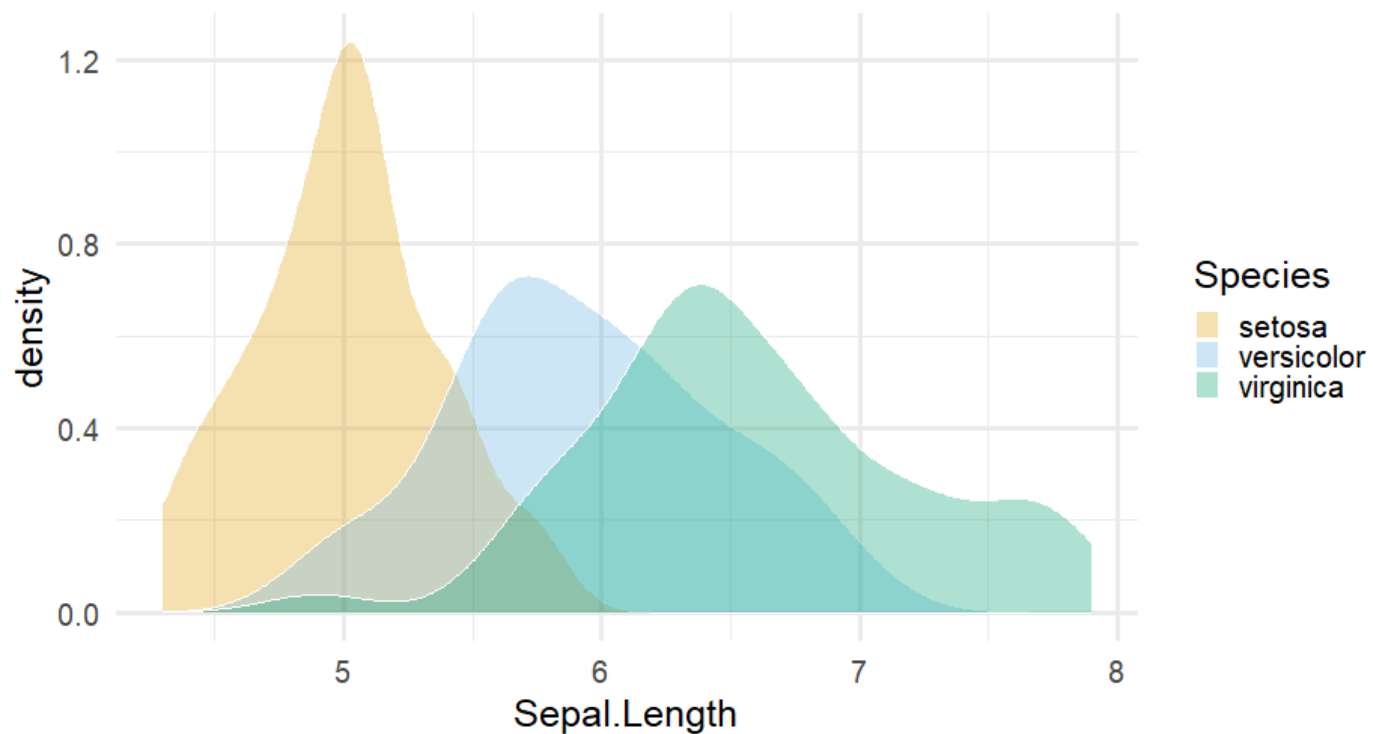
Data from Clause Wilke Book: Fundamentals of Data Visualizations

Distributions

```
ggplot(iris, aes(Sepal.Length, fill = Species)) +  
  geom_density(alpha = 0.3,  
               color = "white")
```



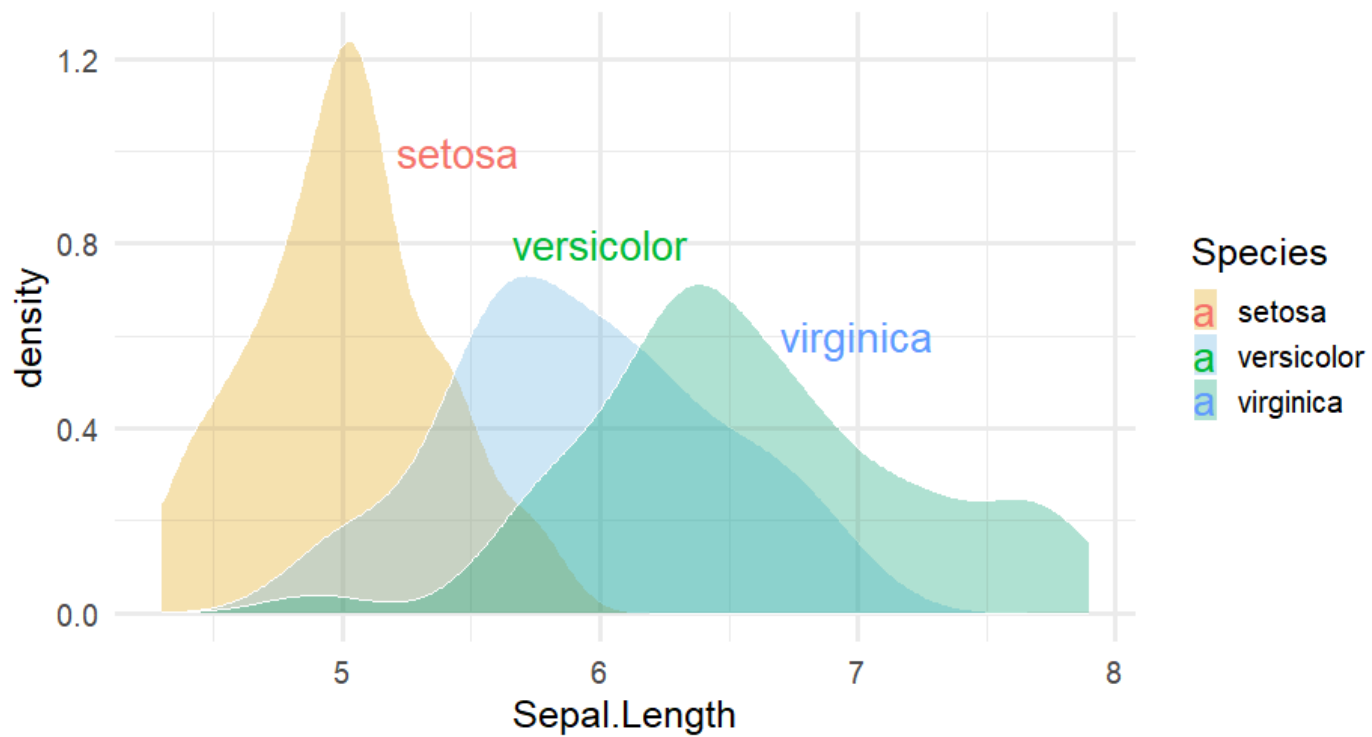
```
ggplot(iris, aes(Sepal.Length, fill = Species)) +  
  geom_density(alpha = 0.3,  
               color = "white") +  
  scale_fill_OkabeIto()
```



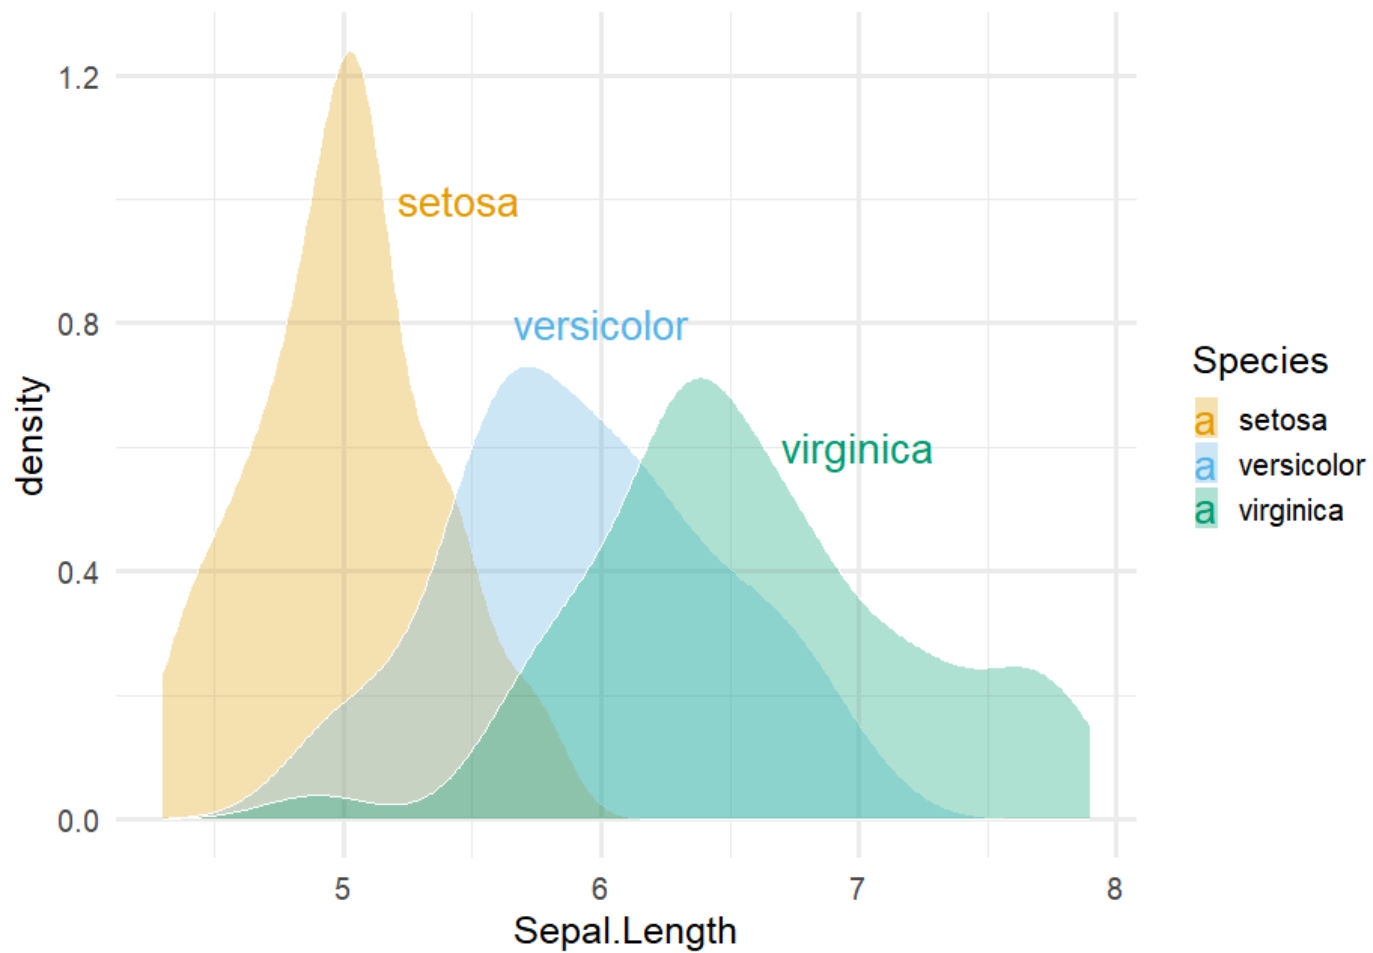
Labeling

One method

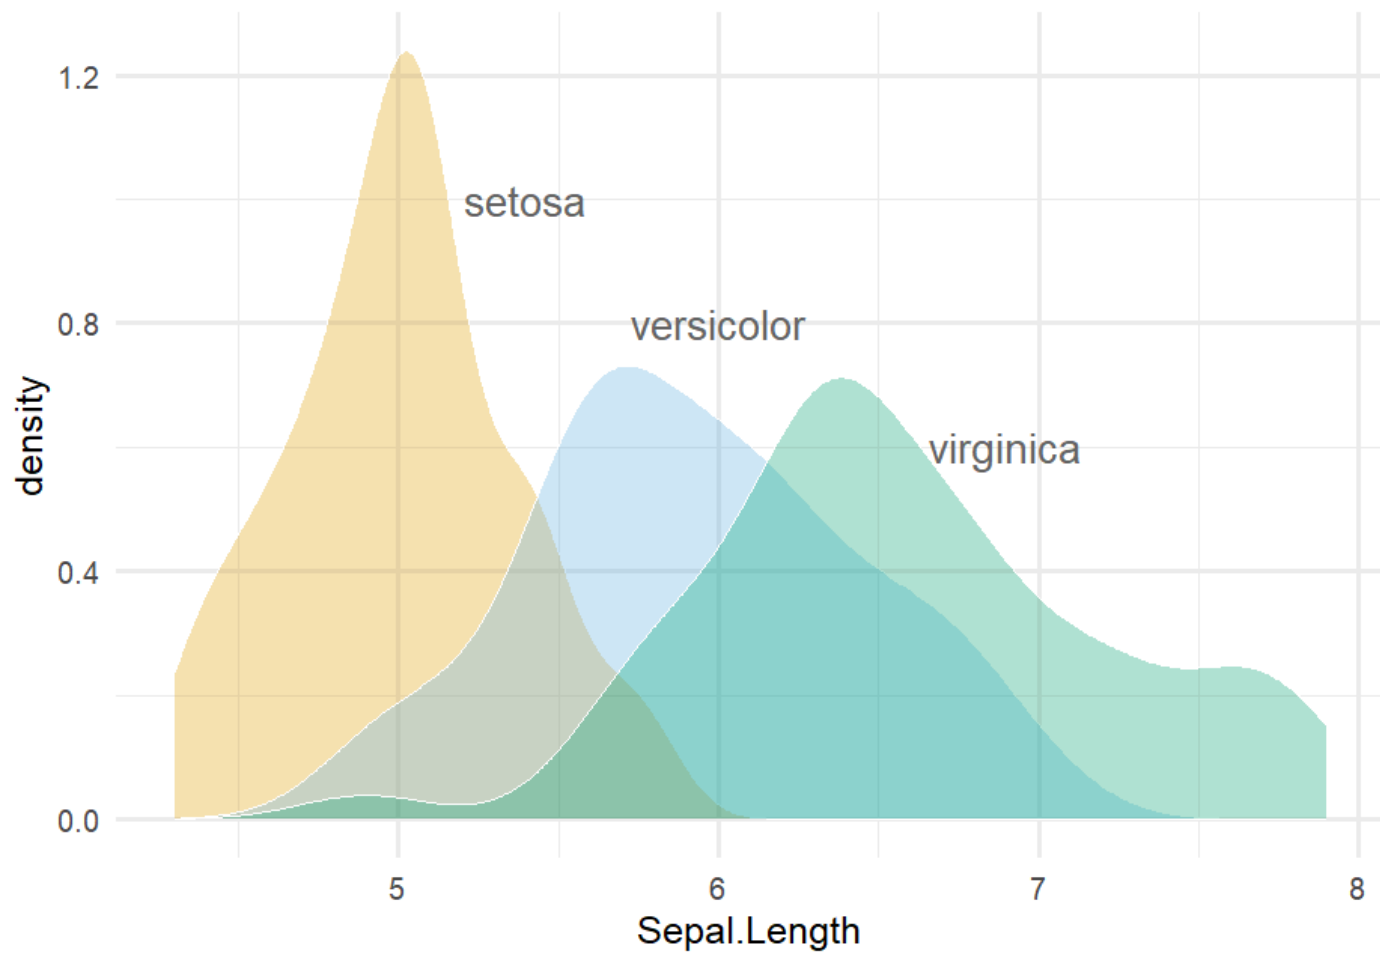
```
label_locs <- tibble(Sepal.Length = c(5.45, 6, 7),  
                     density = c(1, 0.8, 0.6),  
                     Species = c("setosa", "versicolor", "virginica"))  
  
ggplot(iris, aes(Sepal.Length, fill = Species)) +  
  geom_density(alpha = 0.3,  
              color = "white") +  
  scale_fill_OkabeIto() +  
  geom_text(aes(label = Species, y = density, color = Species),  
            data = label_locs)
```



```
ggplot(iris, aes(Sepal.Length, fill = Species)) +  
  geom_density(alpha = 0.3,  
               color = "white") +  
  scale_fill_OkabeIto() +  
  scale_color_OkabeIto() +  
  geom_text(aes(label = Species, y = density, color = Species),  
            data = label_locs) +  
  guides(color = "none",  
         fill = "none")
```




```
label_locs <- tibble(Sepal.Length = c(5.4, 6, 6.9),  
                     density = c(1, 0.75, 0.6),  
                     Species = c("setosa", "versicolor", "virginica"))  
  
ggplot(iris, aes(Sepal.Length, fill = Species)) +  
  geom_density(alpha = 0.3,  
              color = "white") +  
  scale_fill_okabeIto() +  
  scale_color_okabeIto() +  
  geom_text(aes(label = Species, y = density),  
            color = "gray40",  
            data = label_locs) +  
  guides(fill = "none")
```

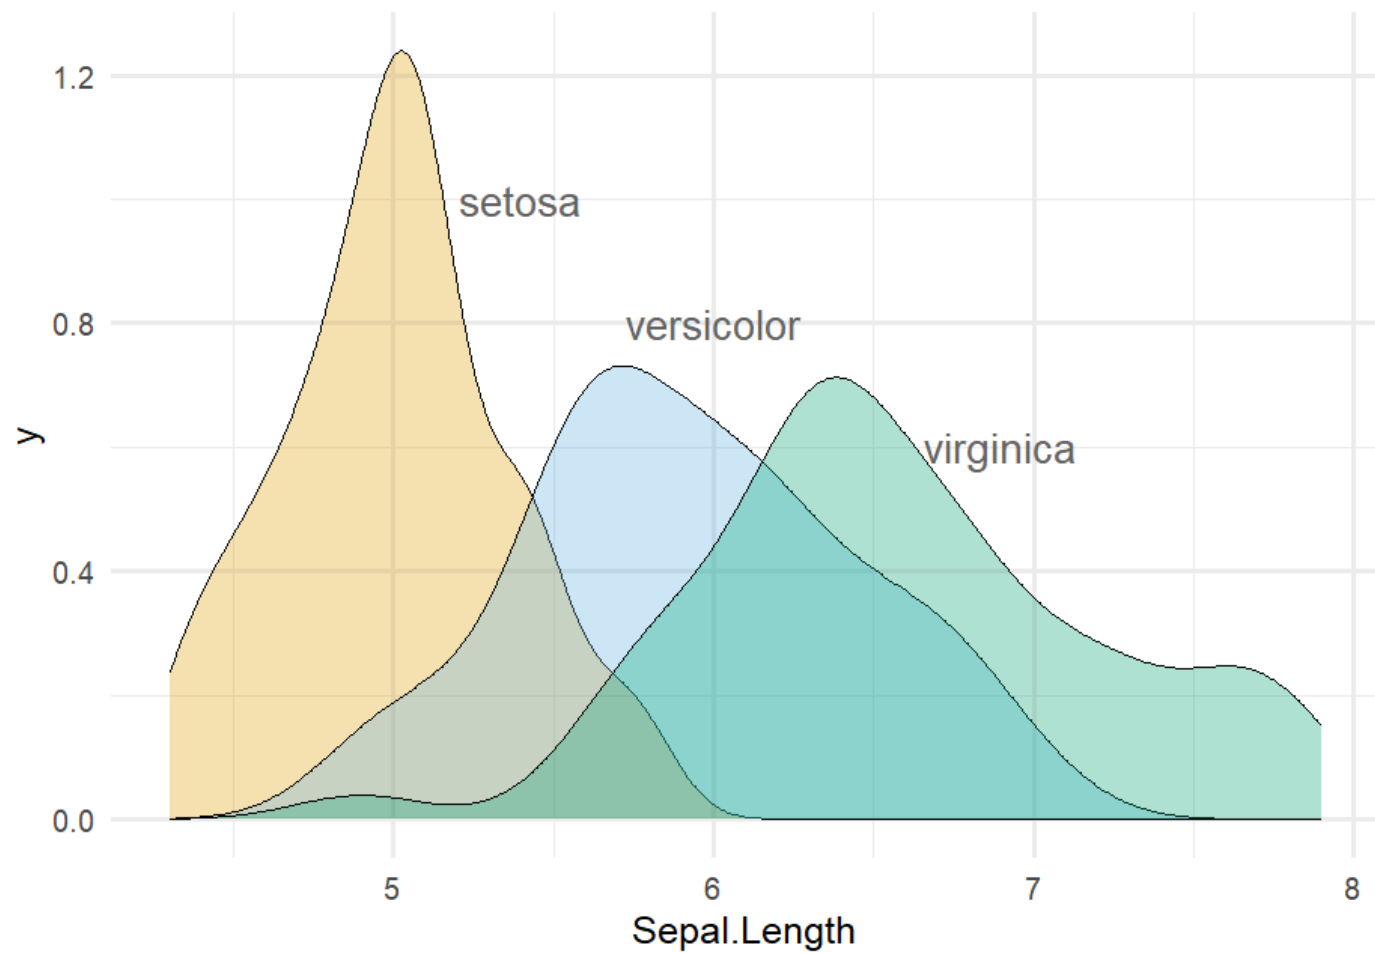


Other options

- Rather than using a new data frame, you could use multiple calls to `annotate`.
- One is not necessarily better than the other, but I prefer the data frame method
- Keep in mind you can always use multiple data sources within a single plot
 - Each layer can have its own data source
 - Common in geographic data in particular

Annotate example

```
ggplot(iris, aes(Sepal.Length, fill = Species)) +  
  geom_density(alpha = 0.3) +  
  scale_fill_OkabeIto() +  
  scale_color_OkabeIto() +  
  annotate("text", label = "setosa", x = 5.45, y = 1, color = "green") +  
  annotate("text", label = "versicolor", x = 6, y = 0.8, color = "red") +  
  annotate("text", label = "virginica", x = 7, y = 0.6, color = "blue") +  
  guides(fill = "none")
```



Practice

Use the diamonds dataset again.

- Compute the mean carat size for each color
- Create a bar chart
- Include labels for each bar rounding the actual value to two decimals
- Make any other modifications you'd like

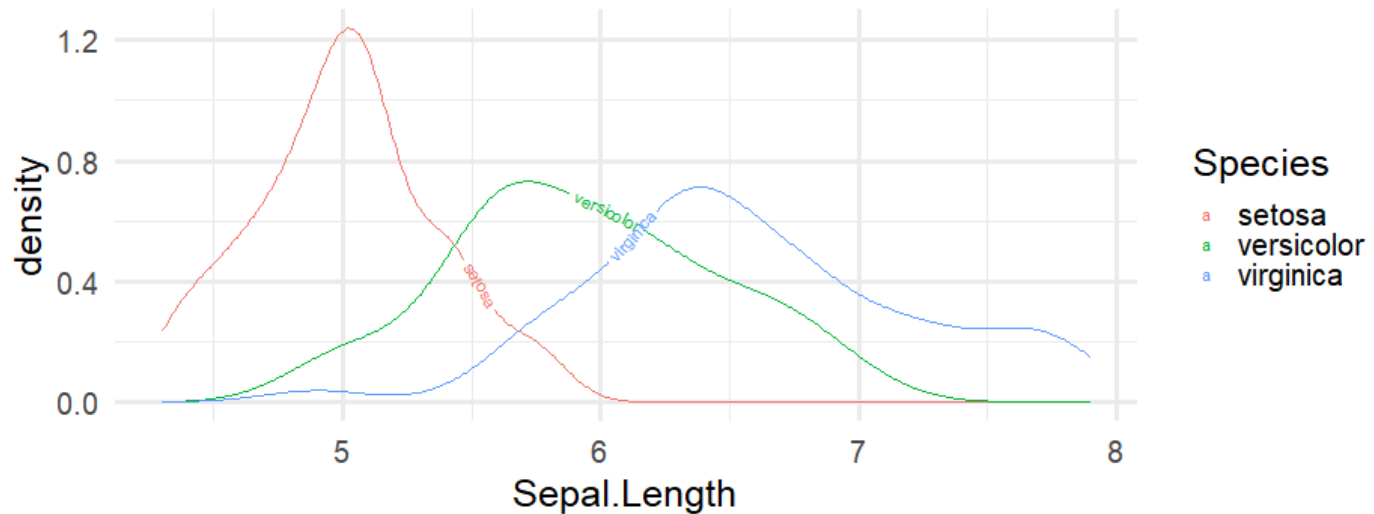
05:00

{geomtextpath}

Not super new package, but looks really cool; still I haven't used much

{geomtextpath}

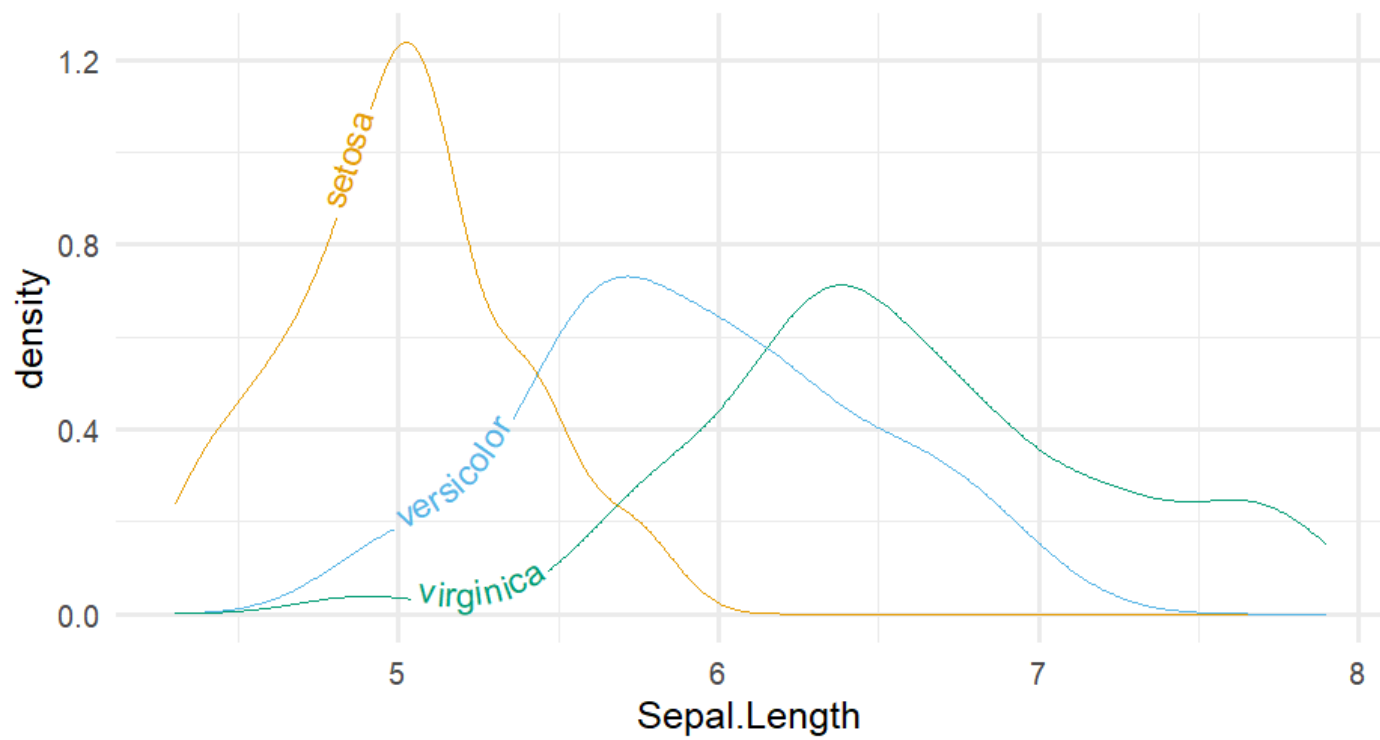
```
#install.packages("geomtextpath")  
library(geomtextpath)  
  
ggplot(iris, aes(Sepal.Length)) +  
  geom_textdensity(aes(color = Species, label = Species))
```



Slight modifications

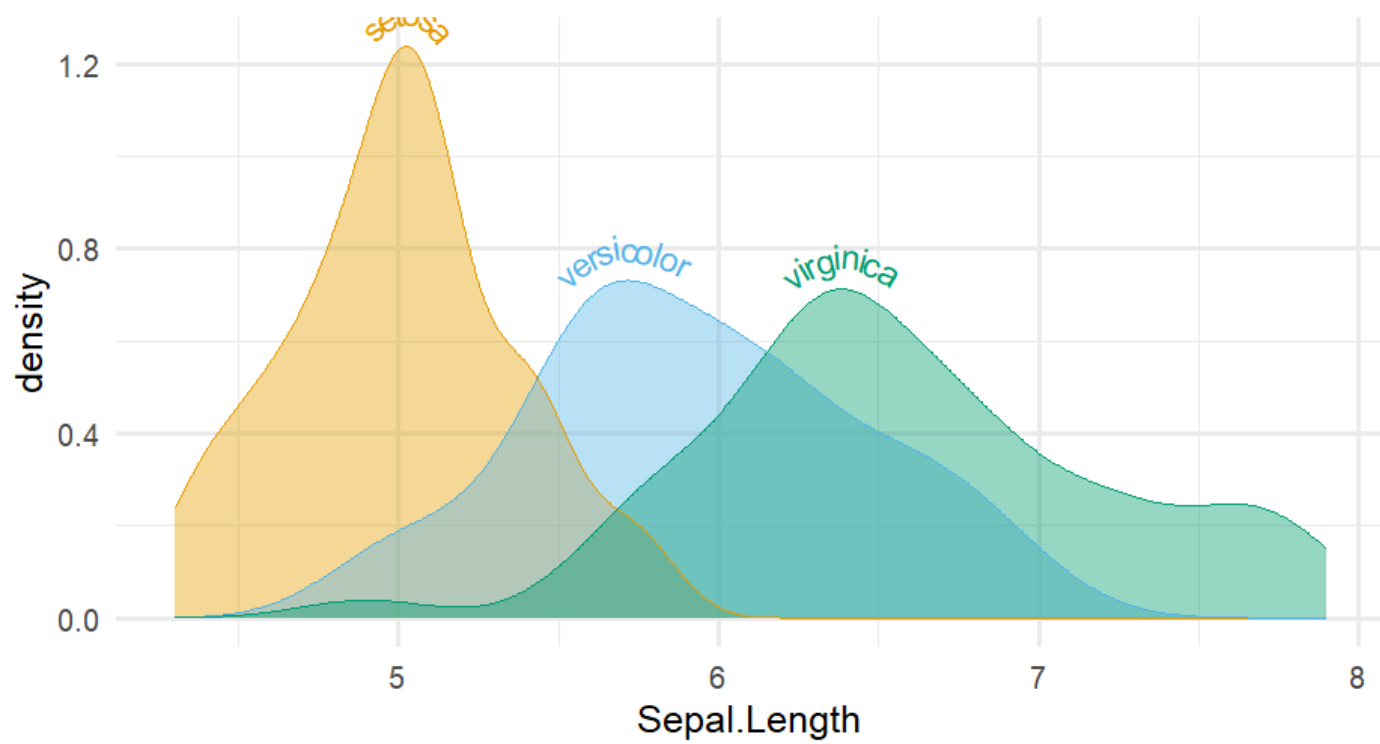
```
ggplot(iris, aes(Sepal.Length)) +  
  geom_textdensity(  
    aes(color = Species, label = Species),  
    hjust = 0.2,  
    vjust = 0.3,  
    size = 8 # bigger than you'll probs need  
  ) +  
  scale_color_OkabeIto() +  
  theme(legend.position = "none")
```

Note I couldn't get `aes(fill = Species)` to work.



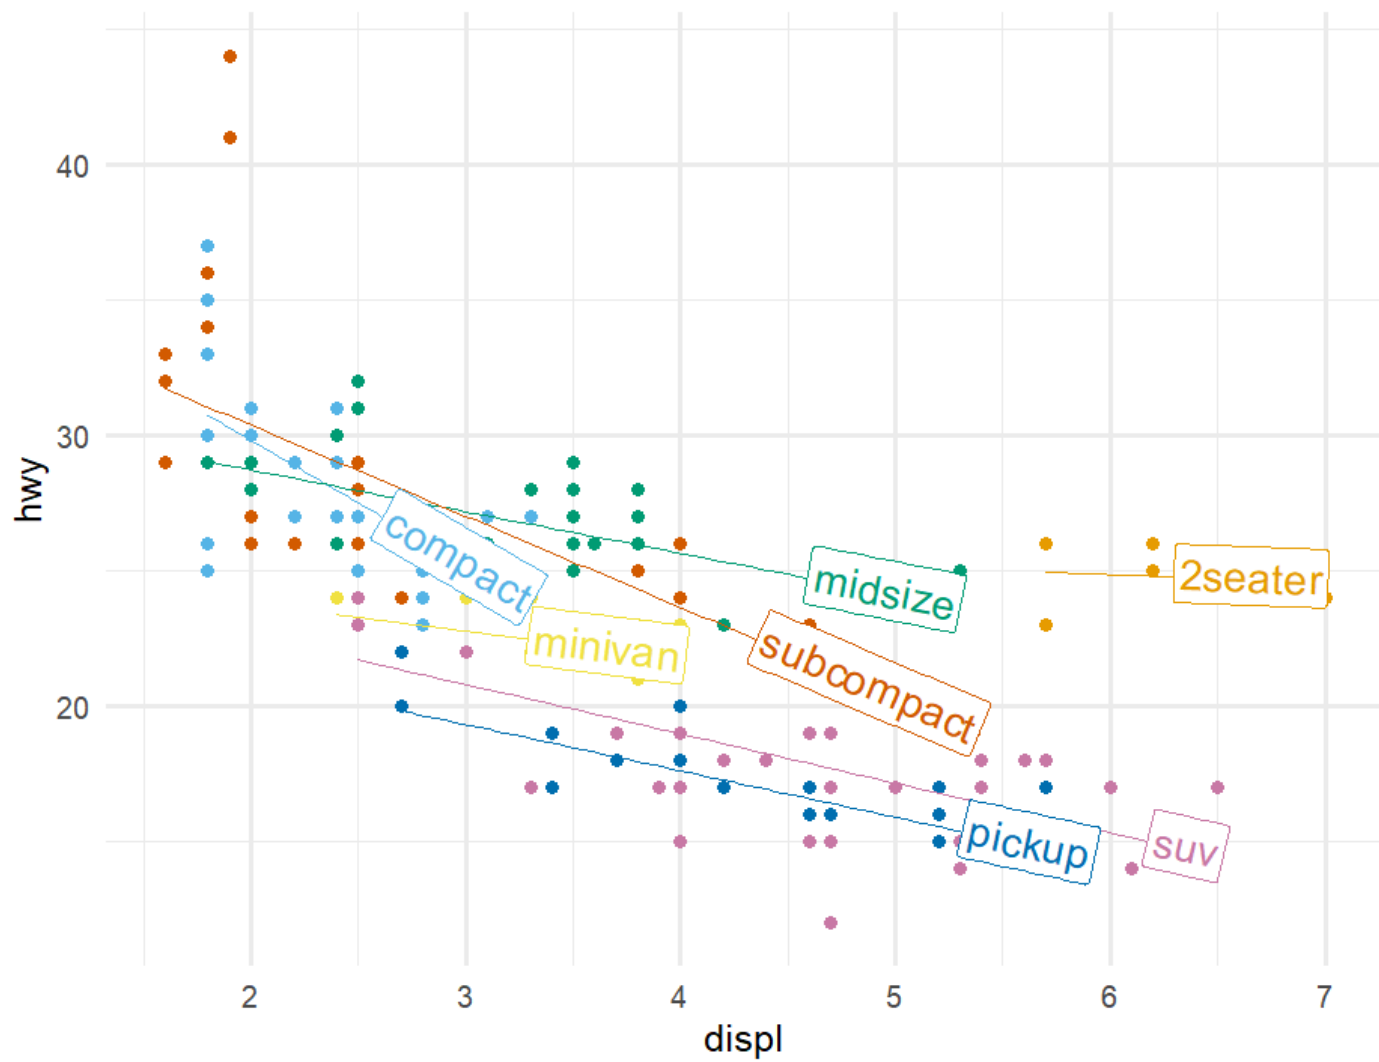
A workaround

```
ggplot(iris, aes(Sepal.Length)) +  
  geom_density(aes(fill = Species), alpha = 0.4) +  
  geom_textdensity(  
    aes(color = Species, label = Species),  
    hjust = "ymax",  
    vjust = -0.3,  
    size = 8  
  ) +  
  scale_color_okabeIto() +  
  scale_fill_okabeIto() +  
  theme(legend.position = "none")
```



Smooths

```
ggplot(mpg, aes(displ, hwy, color = class)) +  
  geom_point() +  
  geom_labelsmooth(  
    aes(label = class),  
    method = "lm",  
    hjust = "xmax"  
  ) +  
  scale_color_okabeIto() +  
  guides(color = "none")
```



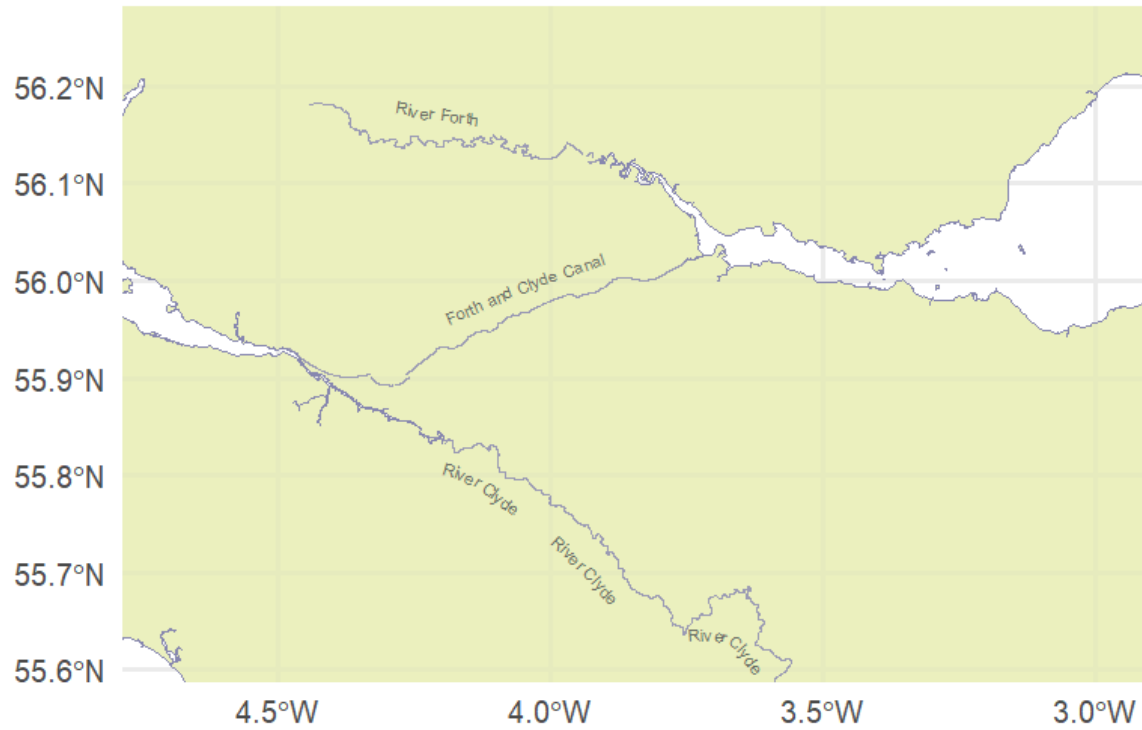
Drop-in replacements

ggplot geom	Text equivalent	Label equivalent
<code>geom_path</code>	<code>geom_textpath</code>	<code>geom_labelpath</code>
<code>geom_segment</code>	<code>geom_textsegment</code>	<code>geom_labelsegment</code>
<code>geom_line</code>	<code>geom_textline</code>	<code>geom_labelline</code>
<code>geom_abline</code>	<code>geom_textabline</code>	<code>geom_labelabline</code>
<code>geom_hline</code>	<code>geom_texthline</code>	<code>geom_labelhline</code>
<code>geom_vline</code>	<code>geom_textvline</code>	<code>geom_labelvline</code>
<code>geom_curve</code>	<code>geom_textcurve</code>	<code>geom_labelcurve</code>
<code>geom_density</code>	<code>geom_textdensity</code>	<code>geom_labeldensity</code>
<code>geom_smooth</code>	<code>geom_textsmooth</code>	<code>geom_labelsmooth</code>

Even works w/Maps!

```
library(sf)

ggplot() +
  geom_textsf(
    data = waterways,
    aes(label = name),
    text_smoothing = 65,
    linecolour = "#8888B3",
    vjust = -0.8,
    fill = "#E6F0B3",
    alpha = 0.8,
    size = 4
  ) +
  theme(panel.grid = element_line()) +
  lims(x = c(-4.7, -3), y = c(55.62, 56.25))
```

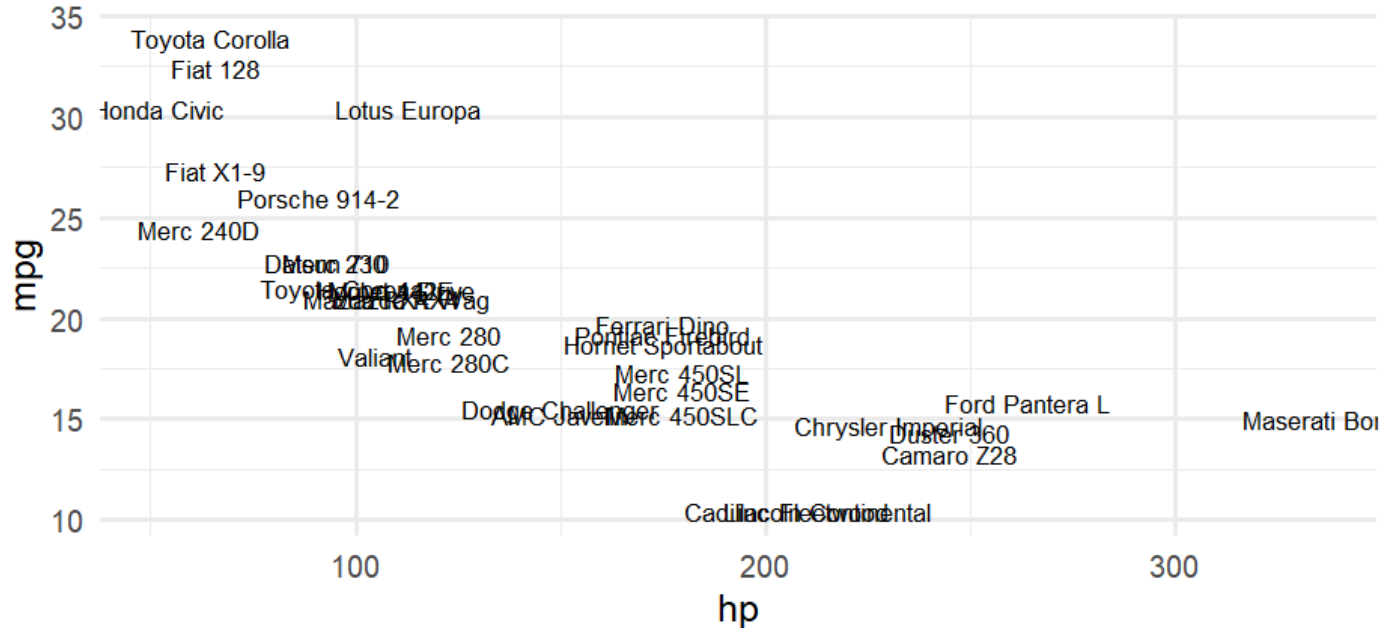



ggrepel

Plot text directly

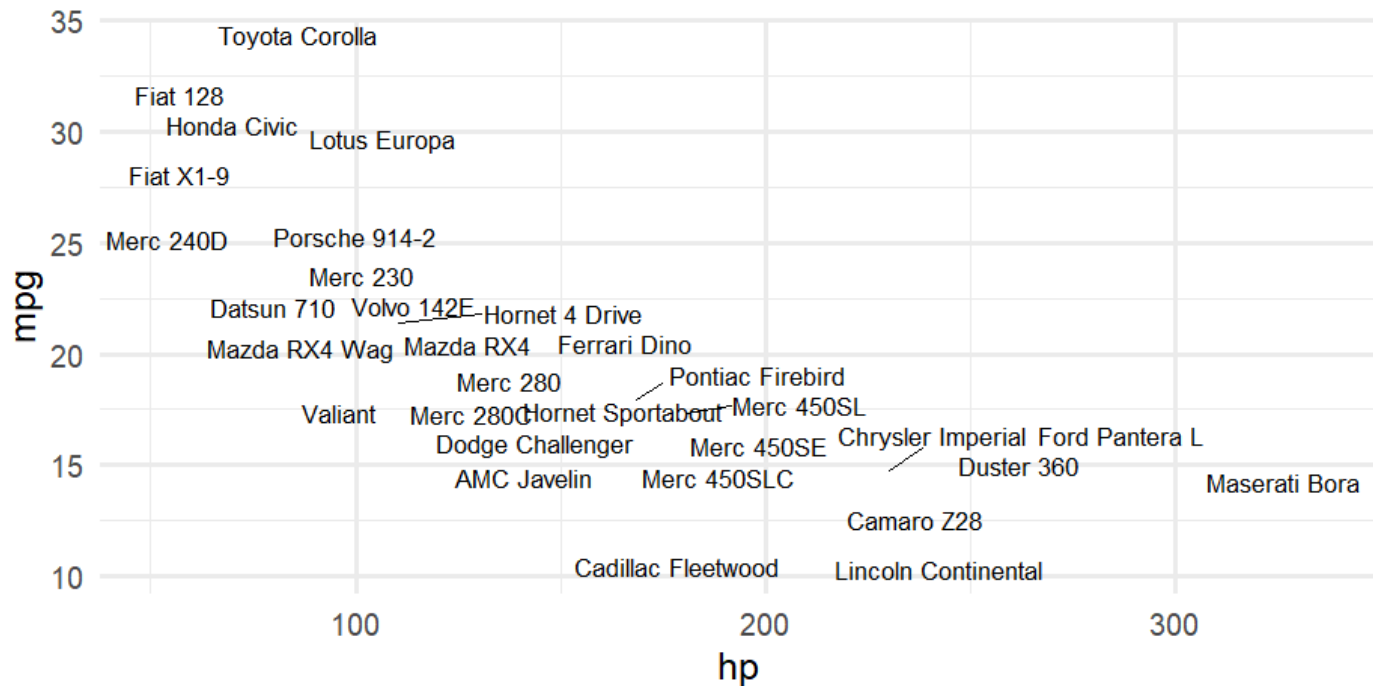
```
cars <- rownames_to_column(mtcars)

ggplot(cars, aes(hp, mpg)) +
  geom_text(aes(label = rowname))
```



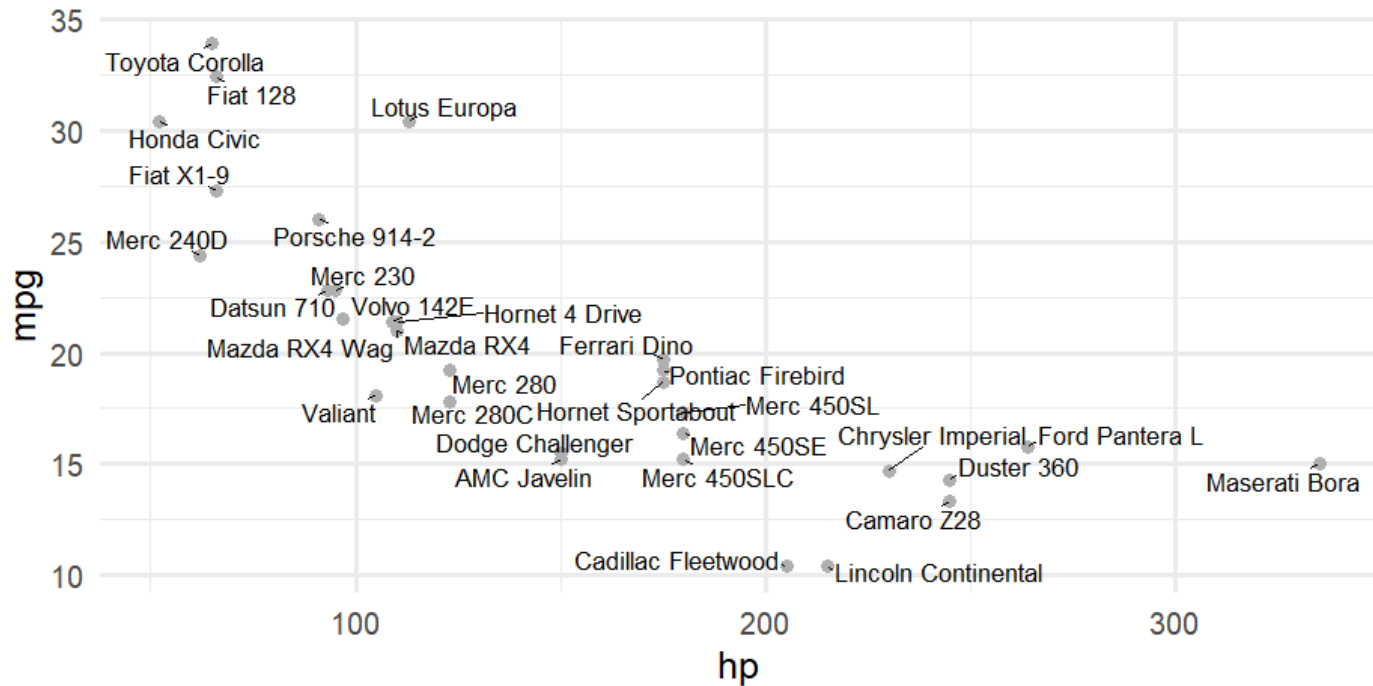
Repel text

```
library(ggrepel)
ggplot(cars, aes(hp, mpg)) +
  geom_text_repel(aes(label = rowname))
```



Slightly better

```
ggplot(cars, aes(hp, mpg)) +  
  geom_point(color = "gray70") +  
  geom_text_repel(aes(label = rowname),  
    min.segment.length = 0)
```



Common use cases

- Label some sample data that makes some theoretical sense (we've seen this before)
- Label outliers
- Label points from a specific group (e.g., similar to highlighting - can be used in conjunction)

Some new data

Please follow along

```
#remotes::install_github("kjhealy/socviz")  
library(socviz)
```

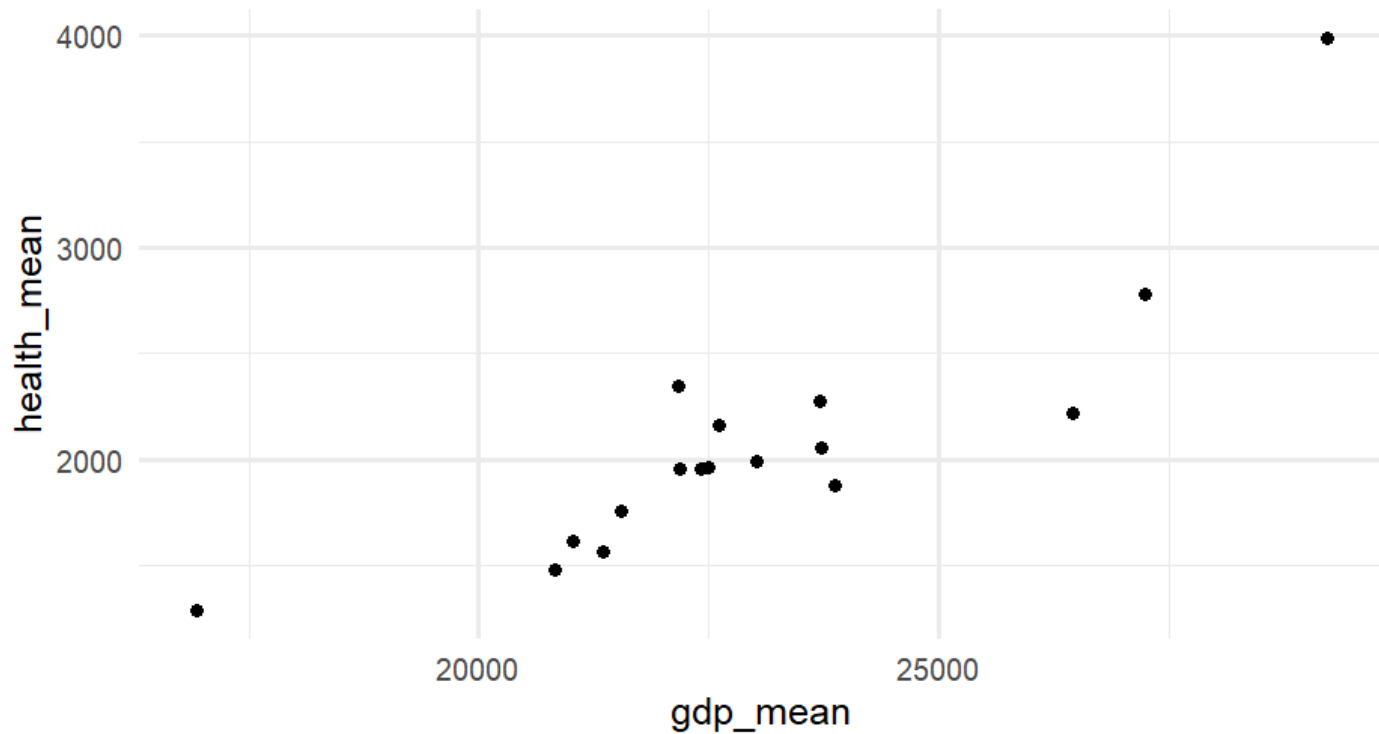
```
by_country <- organdata %>%  
  group_by(consent_law, country) %>%  
  summarize(donors_mean= mean(donors, na.rm = TRUE),  
            donors_sd = sd(donors, na.rm = TRUE),  
            gdp_mean = mean(gdp, na.rm = TRUE),  
            health_mean = mean(health, na.rm = TRUE),  
            roads_mean = mean(roads, na.rm = TRUE),  
            cerebvas_mean = mean(cerebvas, na.rm = TRUE))
```

```
by_country
```

```
## # A tibble: 17 × 8
## # Groups:   consent_law [2]
##   consent_law country      donors_mean donors_sd gdp_mean health_mean roads_mean
##   <chr>         <chr>          <dbl>      <dbl>    <dbl>      <dbl>      <dbl>
## 1 Informed      Australia        10.6       1.14    22179.     1958.     105.
## 2 Informed      Canada           14.0       0.751    23711.     2272.     109.
## 3 Informed      Denmark          13.1       1.47     23722.     2054.     102.
## 4 Informed      Germany           13.0       0.611    22163.     2349.     113.
## 5 Informed      Ireland           19.8       2.48     20824.     1480.     118.
## 6 Informed      Netherlands       13.7       1.55     23013.     1993.      76.1
## # i 11 more rows
## # i 1 more variable: cerebvas_mean <dbl>
```

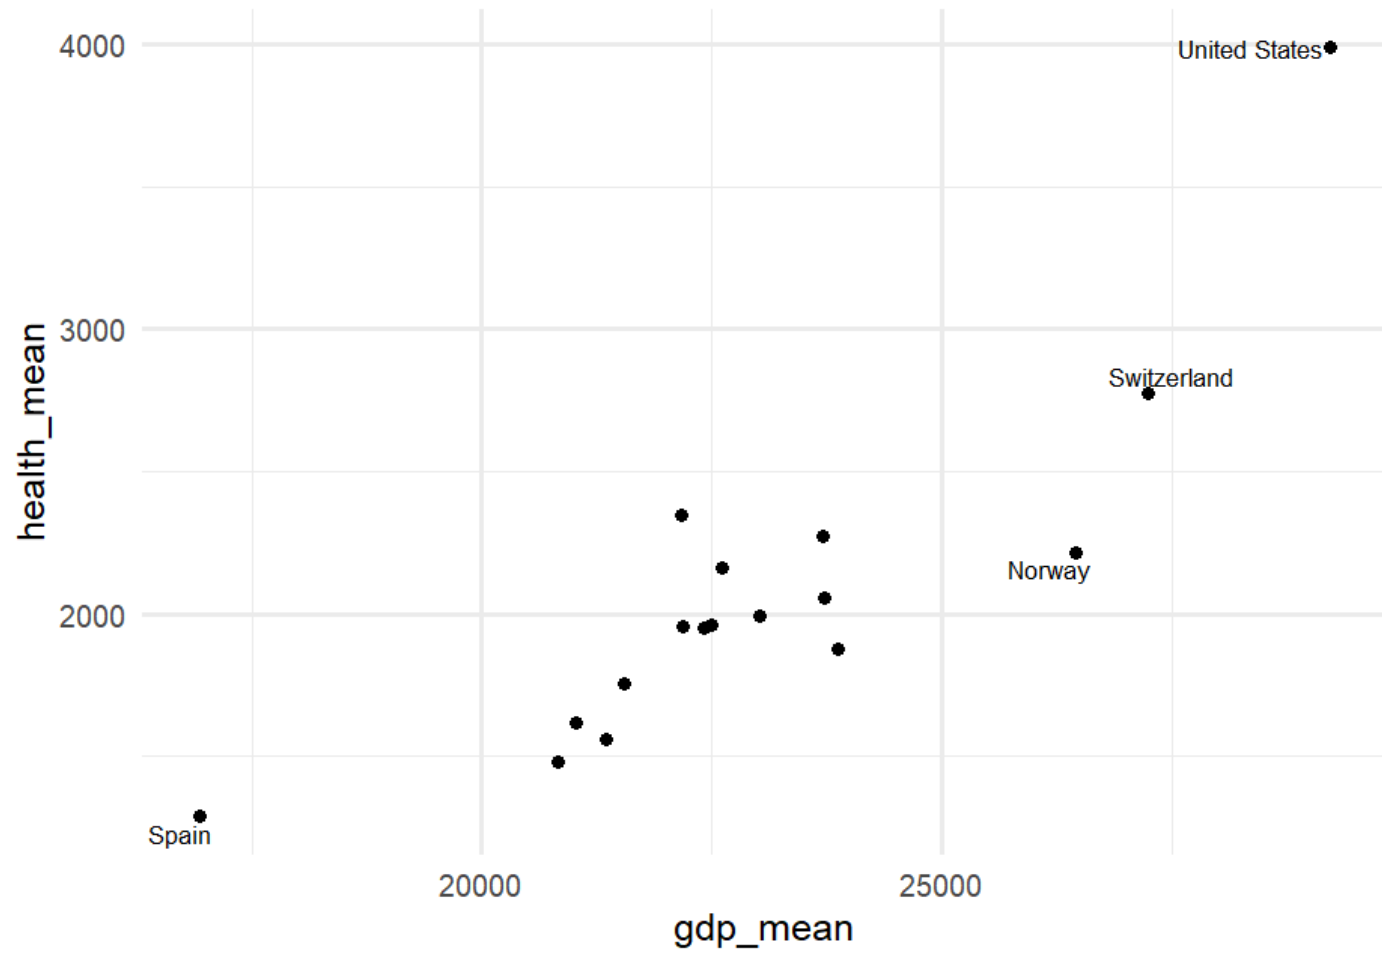

Scatterplot

```
ggplot(by_country, aes(gdp_mean, health_mean)) +  
  geom_point()
```



Outliers

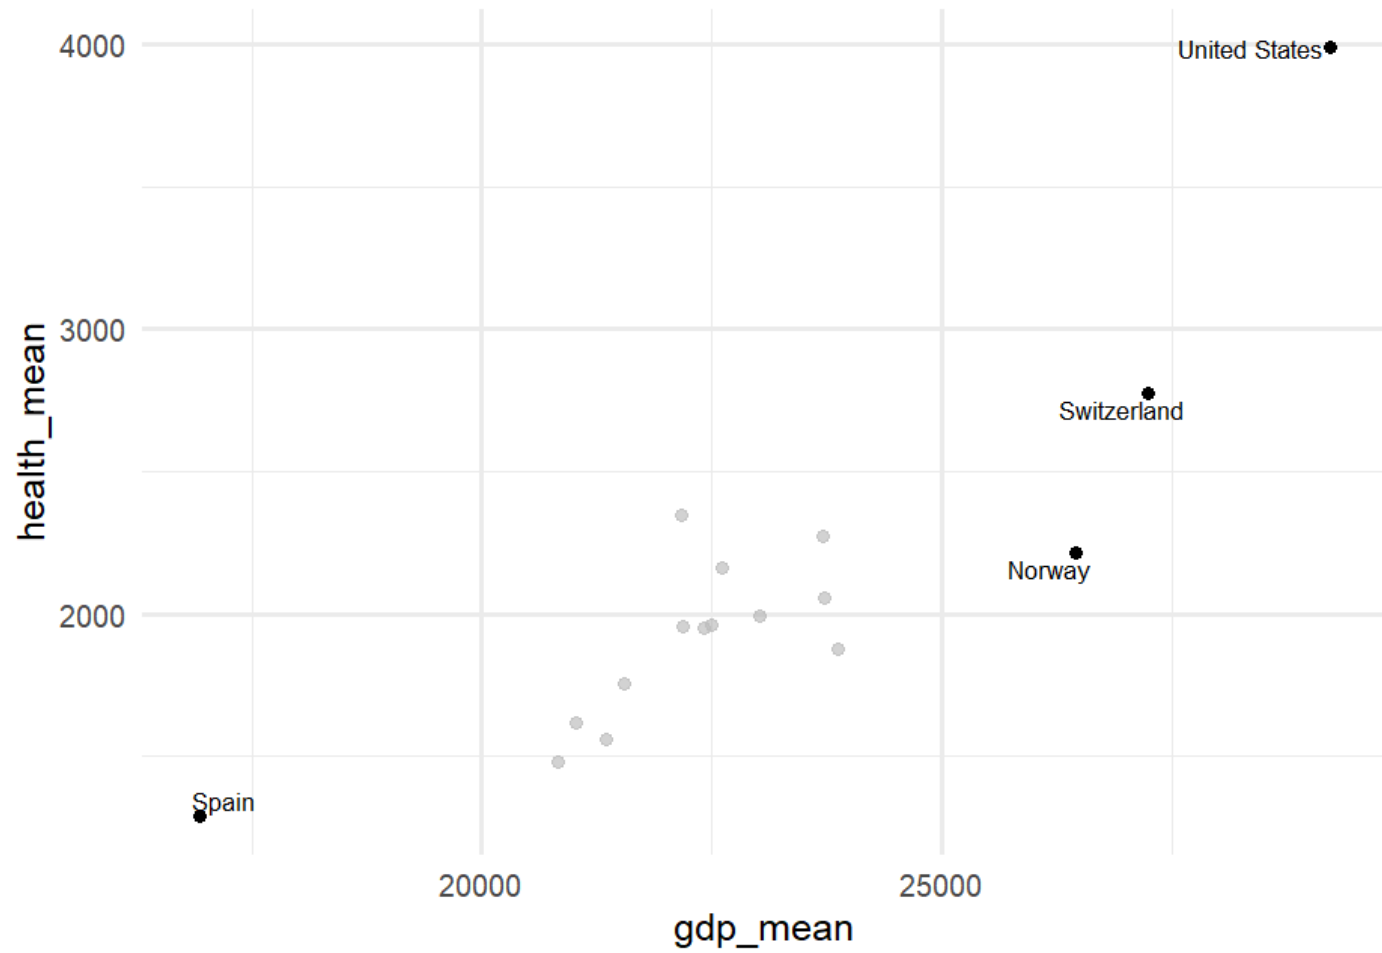
```
ggplot(by_country, aes(gdp_mean, health_mean)) +  
  geom_point() +  
  geom_text_repel(data = filter(by_country,  
                                gdp_mean > 25000 |  
                                gdp_mean < 20000),  
                  aes(label = country))
```



Combine with highlighting

```
library(gghighlight)
ggplot(by_country, aes(gdp_mean, health_mean)) +
  geom_point() +
  gghighlight(gdp_mean > 25000 | gdp_mean < 20000) +
  geom_text_repel(aes(label = country))
```

- Notice you only have to specify the points to highlight and `geom_text_repel` will then only label those points

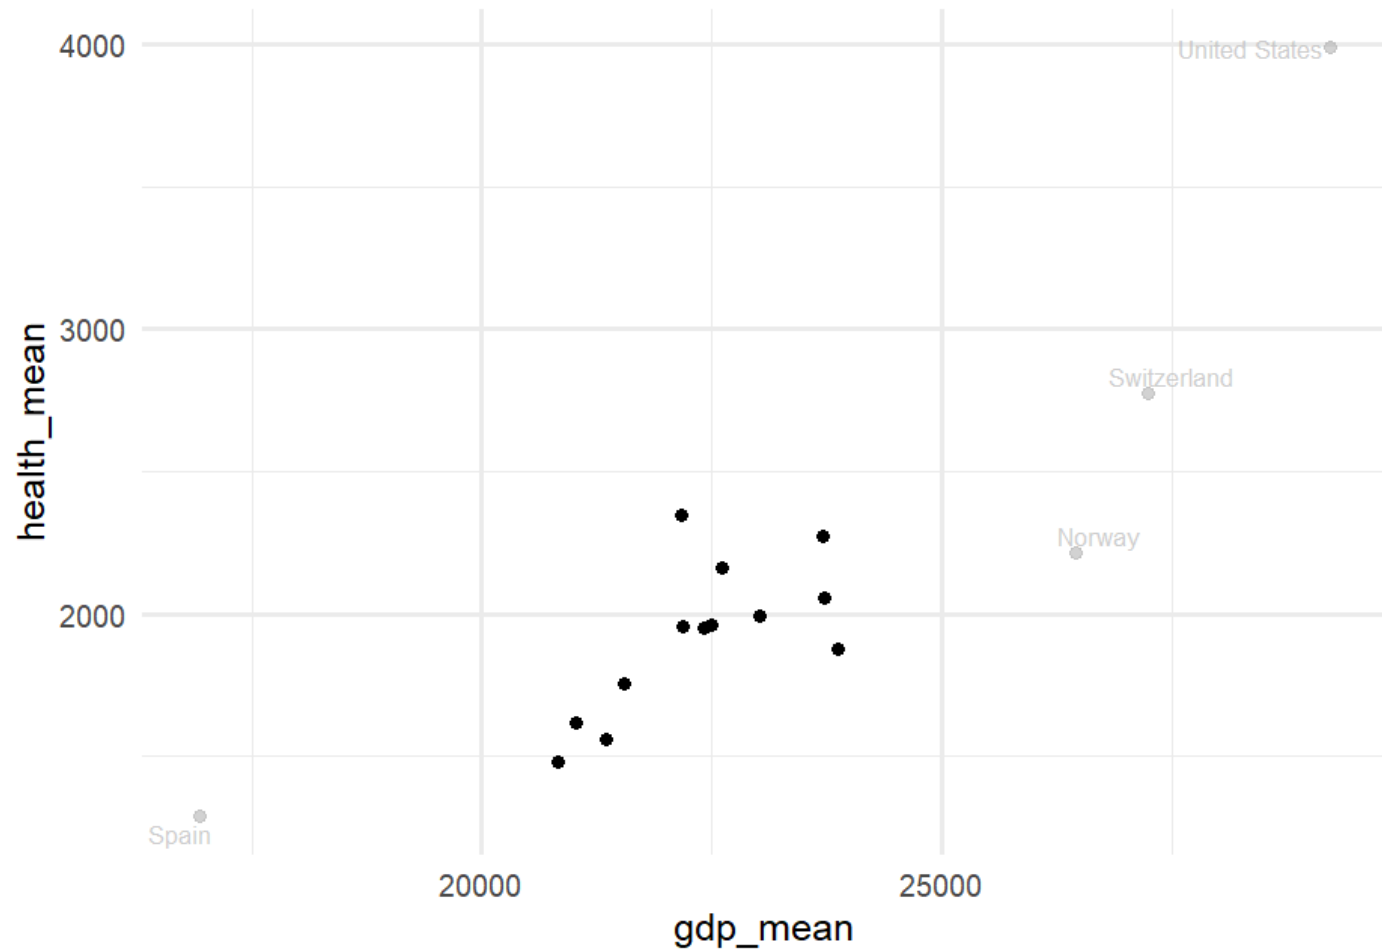


Combine with highlighting

Switch to make outliers grayed out and labeled

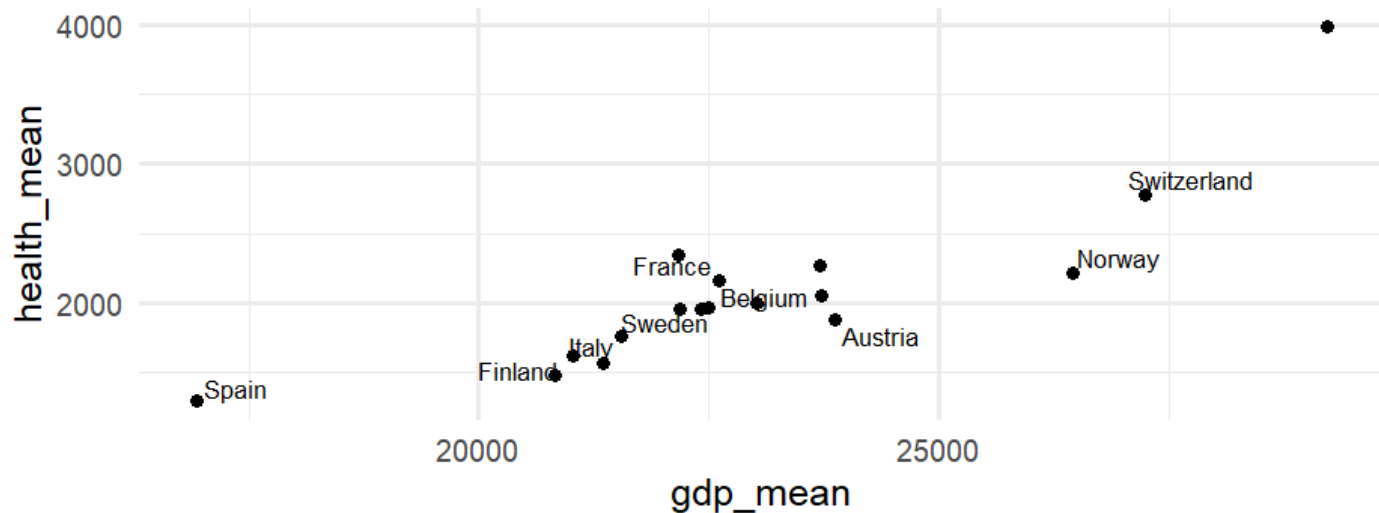
```
ggplot(by_country, aes(gdp_mean, health_mean)) +  
  geom_point() +  
  gghighlight(gdp_mean > 20000 & gdp_mean < 25000 ) +  
  geom_text_repel(data = filter(by_country,  
                                gdp_mean > 25000 |  
                                gdp_mean < 20000),  
                  aes(label = country),  
                  color = "#BEBEBEB3")
```

Note I found the exact gray color by looking at the source code. Specifically, it is the output from `ggplot2::alpha("grey", 0.7)`



By group

```
ggplot(by_country, aes(gdp_mean, health_mean)) +  
  geom_point() +  
  geom_text_repel(data = filter(by_country,  
                                consent_law == "Presumed"),  
                  aes(label = country))
```

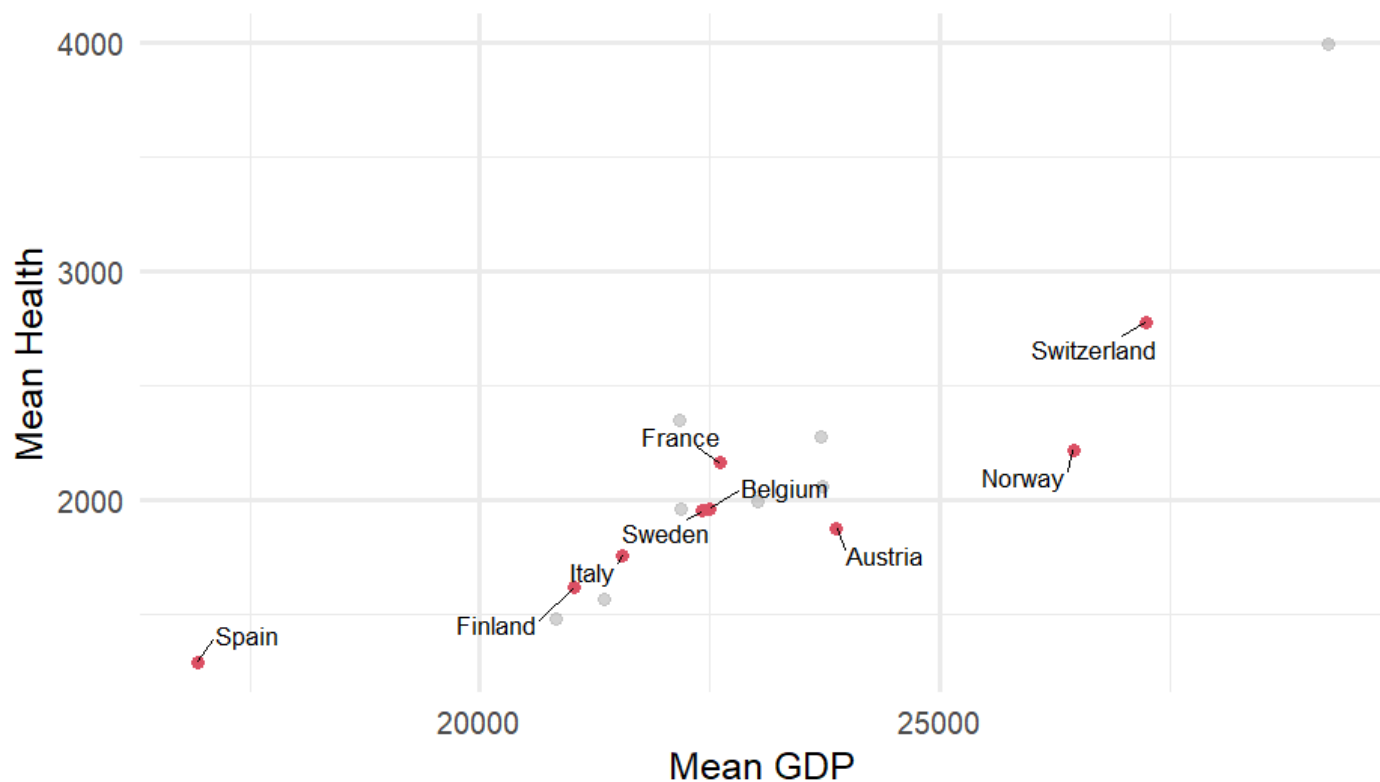


By group

```
ggplot(by_country, aes(gdp_mean, health_mean)) +  
  geom_point(color = "#DC5265") +  
  gghighlight(consent_law == "Presumed") +  
  geom_text_repel(aes(label = country),  
                  min.segment.length = 0,  
                  box.padding = 0.75) +  
  labs(title = "GDP and Health",  
        subtitle = "Countries with a presumed organ donation con",  
        caption = "Data from the General Social Science Survey",  
        x = "Mean GDP",  
        y = "Mean Health")
```

GDP and Health

Countries with a presumed organ donation consent are highlighted



Data from the General Social Science Survey, Distributed through the socviz R package

Practice

Use the mpg dataset

- Group by manufacturer
- Compute the mean highway `hwy` and mean `displ`
- Plot the relation between these means. Plot points and label the manufacturer of each point.

ggforce

Please follow along

Annotating groups of points

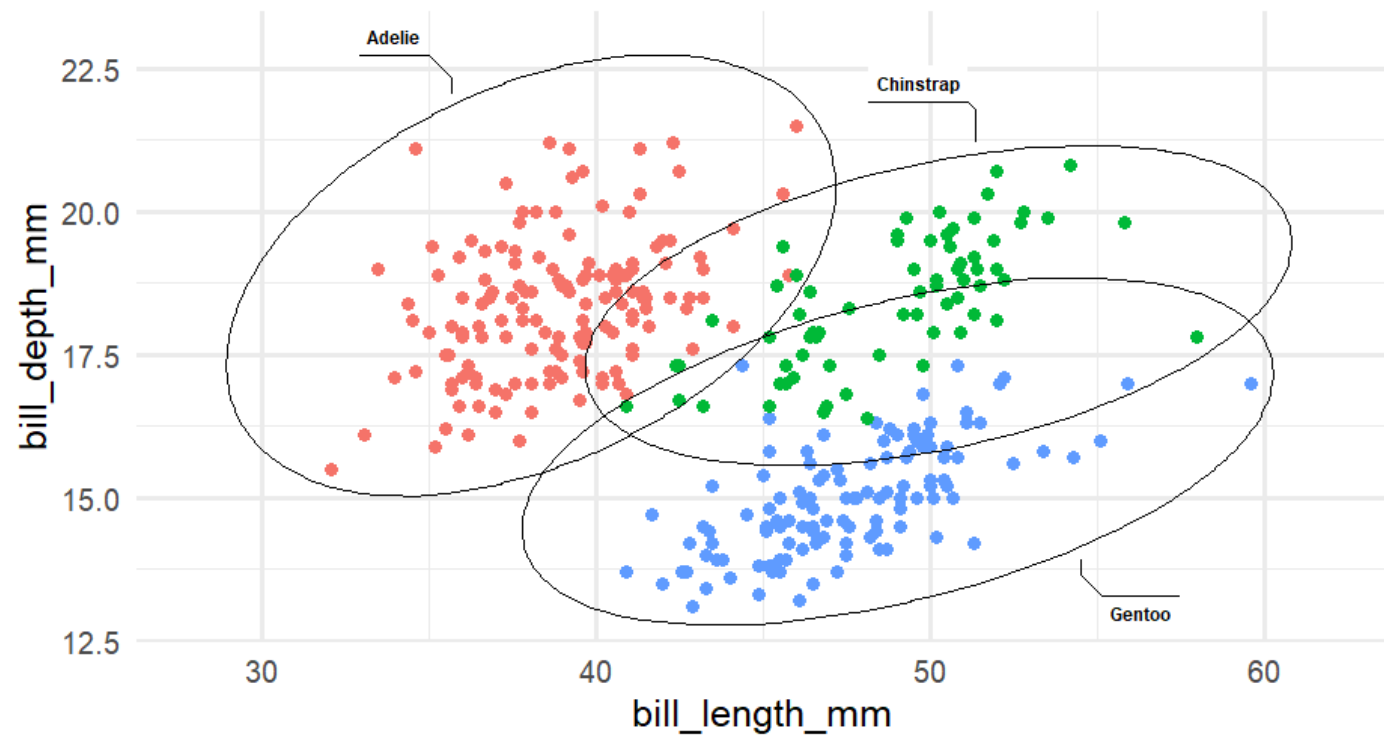
Consider using any of the following from **ggforce** to annotate specific points

- `geom_mark_rect()`
- `geom_mark_circle()`
- `geom_mark_ellipse()`
- `geom_mark_hull()`

Examples

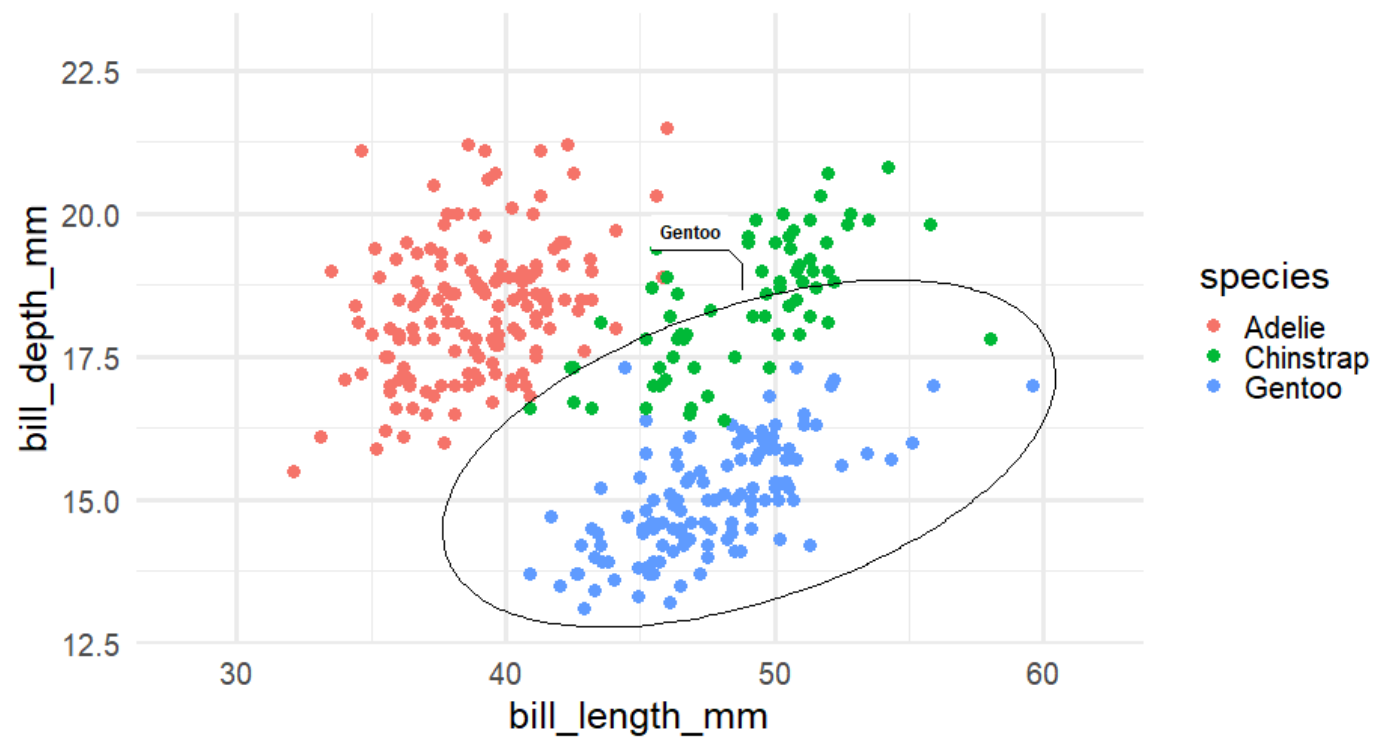
```
library(palmerpenguins)
library(ggforce)

penguins %>%
  drop_na() %>% # Can't take missing data
  ggplot(aes(bill_length_mm, bill_depth_mm)) +
  geom_mark_ellipse(aes(group = species, label = species)) +
  geom_point(aes(color = species)) +
  coord_cartesian(xlim = c(28, 62), ylim = c(13, 23)) +
  guides(color = "none")
```



Limit to a single group

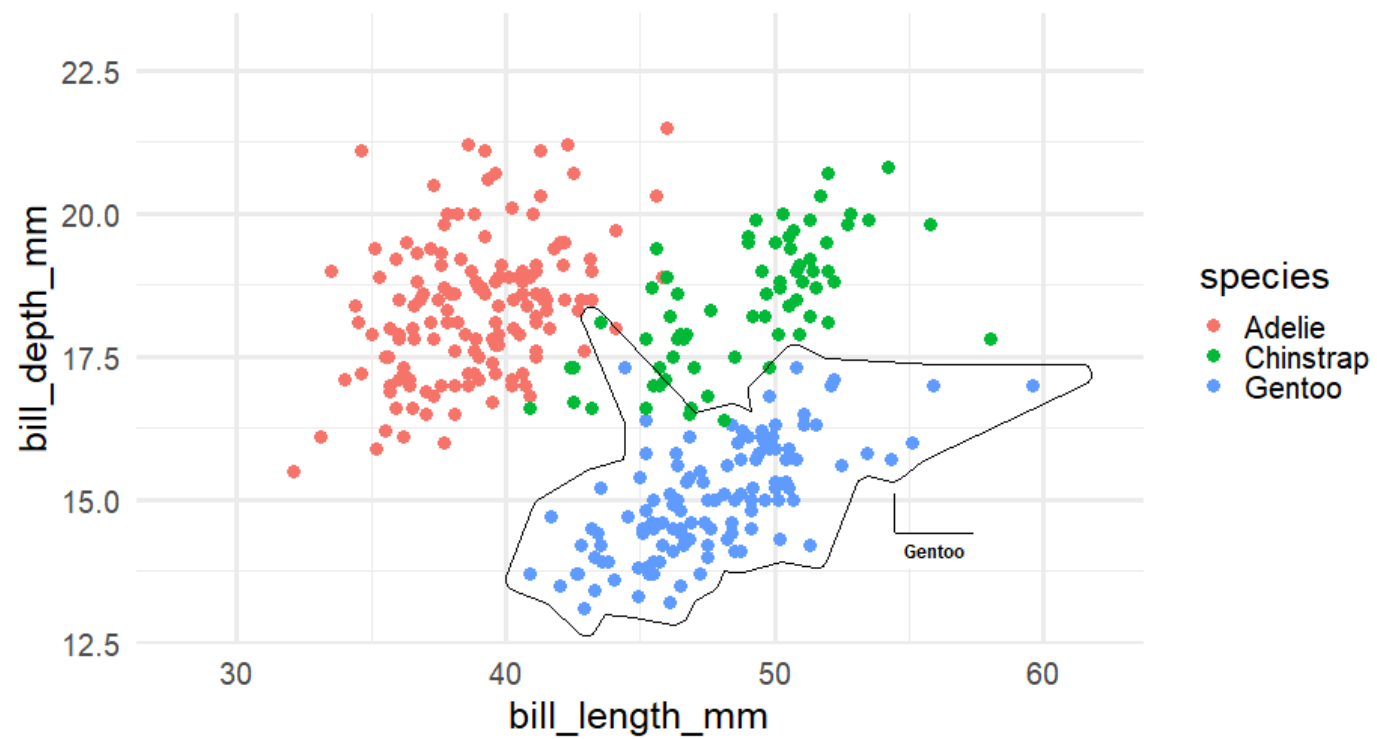
```
penguins %>%  
  drop_na() %>%  
  ggplot(aes(bill_length_mm, bill_depth_mm)) +  
    geom_point(aes(color = species)) +  
    geom_mark_ellipse(aes(group = species, label = species),  
                      data = filter(drop_na(penguins),  
                                    species == "Gentoo")) +  
    coord_cartesian(xlim = c(28, 62), ylim = c(13, 23))
```

Switch to hull

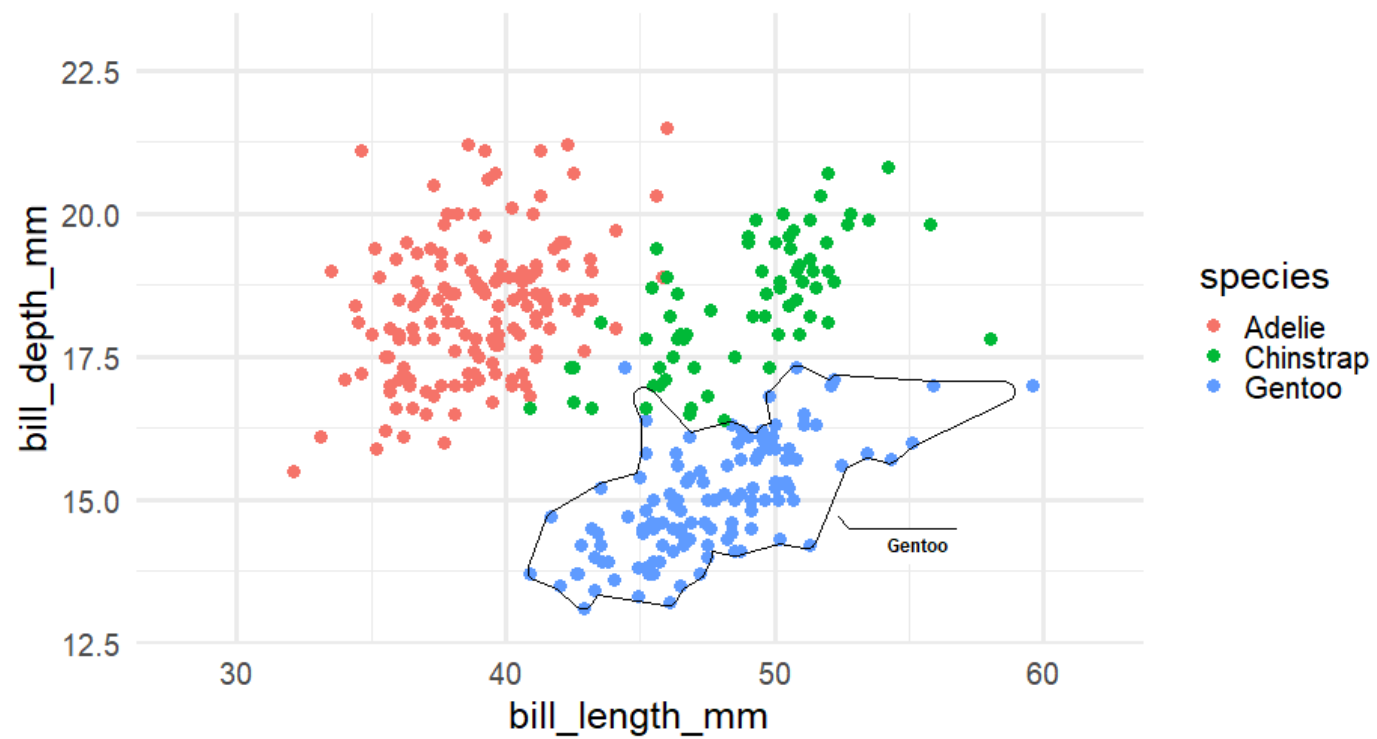
Note - requires the **concaveman** package be installed

```
penguins %>%  
  drop_na() %>%  
  ggplot(aes(bill_length_mm, bill_depth_mm)) +  
    geom_point(aes(color = species)) +  
    geom_mark_hull(aes(group = species, label = species),  
                  data = filter(drop_na(penguins),  
                                species == "Gentoo")) +  
    coord_cartesian(xlim = c(28, 62), ylim = c(13, 23))
```



Change expand

```
penguins %>%  
  drop_na() %>%  
  ggplot(aes(bill_length_mm, bill_depth_mm)) +  
    geom_point(aes(color = species)) +  
    geom_mark_hull(aes(group = species, label = species),  
                   expand = unit(1, "mm"),  
                   data = filter(drop_na(penguins),  
                                   species == "Gentoo")) +  
  coord_cartesian(xlim = c(28, 62), ylim = c(13, 23))
```



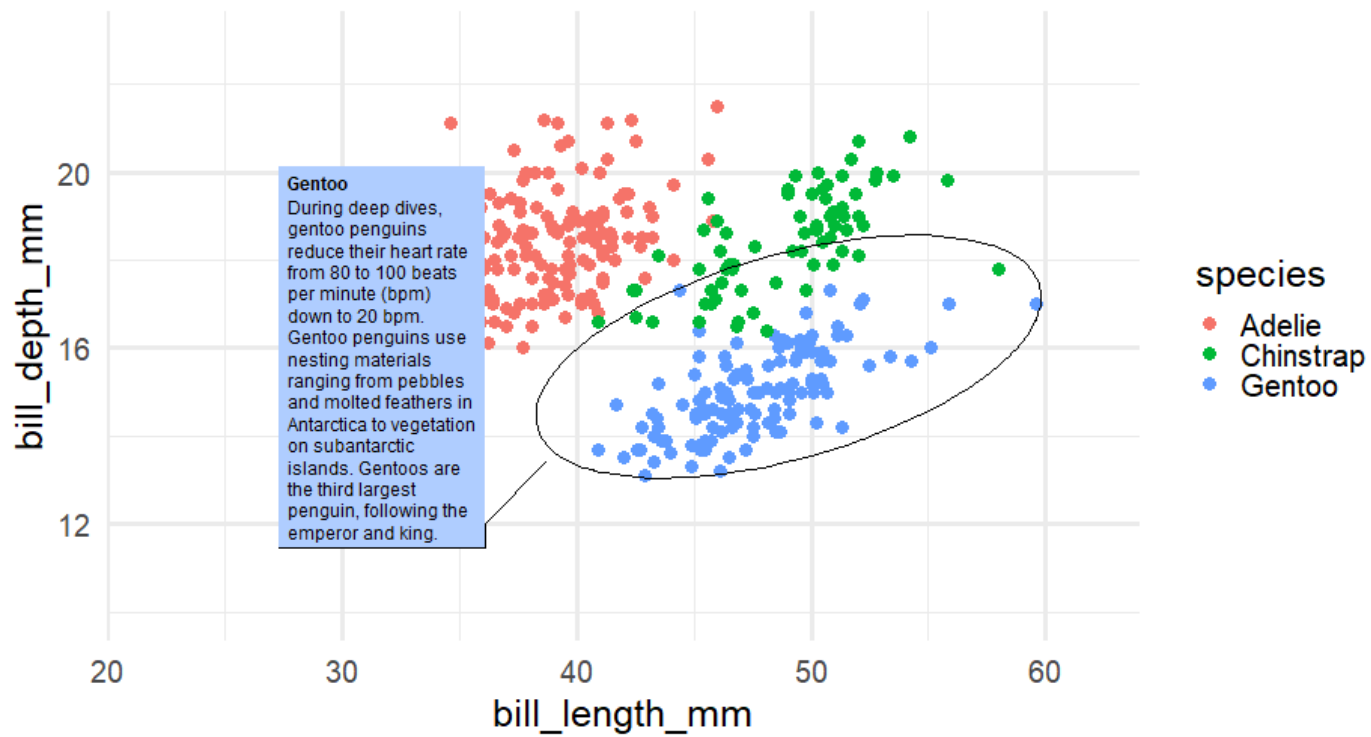
More in-depth annotations

First create a description

```
penguins <- penguins %>%  
  mutate(desc = ifelse(species != "Gentoo", "", "During deep dive
```

Now add as a description

```
penguins %>%  
  drop_na() %>%  
  ggplot(aes(bill_length_mm, bill_depth_mm)) +  
    geom_point(aes(color = species)) +  
    geom_mark_ellipse(aes(group = species,  
                          label = species,  
                          description = desc),  
                      data = filter(drop_na(penguins),  
                                     species == "Gentoo"),  
                      label.fill = "#b3cfff") +  
  coord_cartesian(xlim = c(28, 62), ylim = c(13, 23))
```

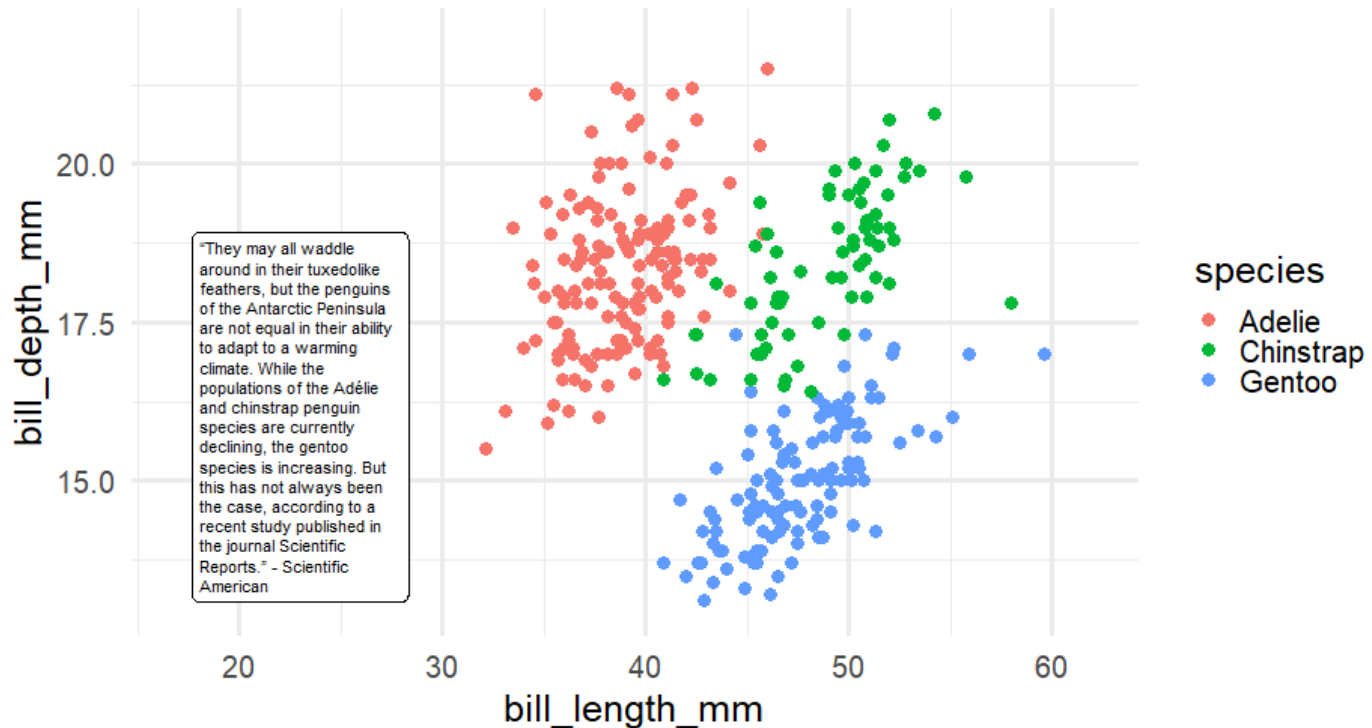


Similar

We can also just add a textbox through **{ggtext}**

```
txtbox <- tibble(  
  bill_length_mm = 23,  
  bill_depth_mm = 16,  
  lab = '"They may all waddle around in their tuxedolike feathers'  
)
```

```
penguins %>%
  drop_na() %>%
  ggplot(aes(bill_length_mm, bill_depth_mm)) +
    geom_point(aes(color = species)) +
    ggtext::geom_textbox(aes(label = lab),
                        data = txtbox) +
    coord_cartesian(xlim = c(17, 62), ylim = c(13, 22))
```



Last bit

The **ggforce** package is well worth exploring more.

See [here](#) for a nice walkthrough that has good data viz and uses some of the **ggforce** functions (as well as illustrating a few other cool packages)

Saving plots

And potentially making additional edits

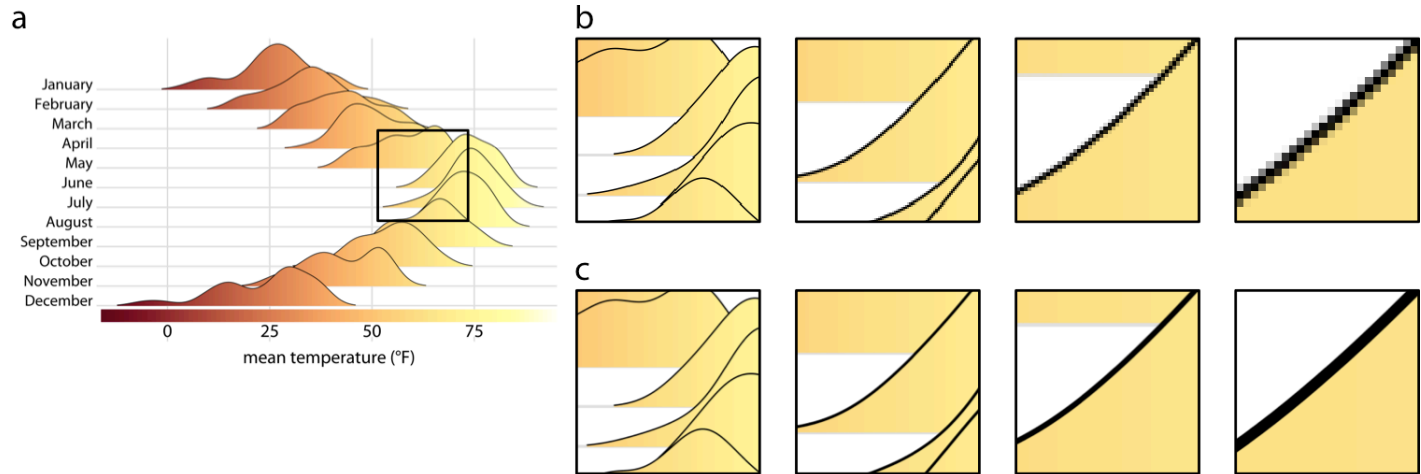
Raster/Vector

We'll talk about this again when discussing maps, but it relates to saving as well.

Raster (also called bitmap) stores images as a grid of points (pixels)

Vector store instructions for how the figure should be drawn. Image is redrawn as it is printed/displayed on screen

Differences



Vector graphics: pdf, eps, svg

Raster graphics: png, jpeg, tiff, gif

Downsides to vector graphics

- Possible differences in appearance between displays (programs, computers, etc.)
- Very large/complex figures can balloon to giant file sizes and be slow to render

Lossy/Lossless compression

- Lossless - guarantees image is, pixel for pixel, identical to original
 - png and tiff use lossless compression
- Lossy - accepts some minor image artifacts to reduce size
 - jpeg

Practical advice

Export to PDF

If that won't work (web), use png

Practice

- Create a plot
- Save it as a PDF with `ggsave()`

Note, the first argument to `ggsave()` is the `path`, so you could do something like

```
ggsave(here::here("myplot.pdf"))
```

You can also specify the width/height.

By default, it will save the last plot you produced, but you can also specify it with `plot =` argument, where you pass an object that has the plot

Modifications

I rarely do this, but if I do, I tend to use Inkscape, which is free.

```
ggplot(mpg, aes(displ, hwy)) +  
  geom_point(color = "gray80") +  
  geom_point(color = "#FD7A43", data = filter(mpg, cyl == 4))  
  
ggsave("~/Desktop/example-plot.pdf", width = 6.5, height = 6.5)
```

Compound figures

Please follow along

Options

My favorite: `{patchwork}`

- `{cowplot}`
- `{ggpubr}`
- `{gridExtra}`

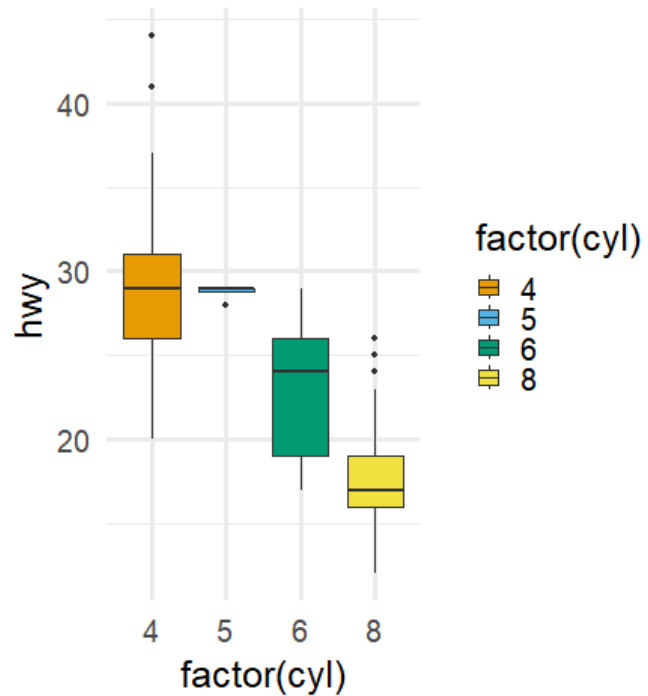
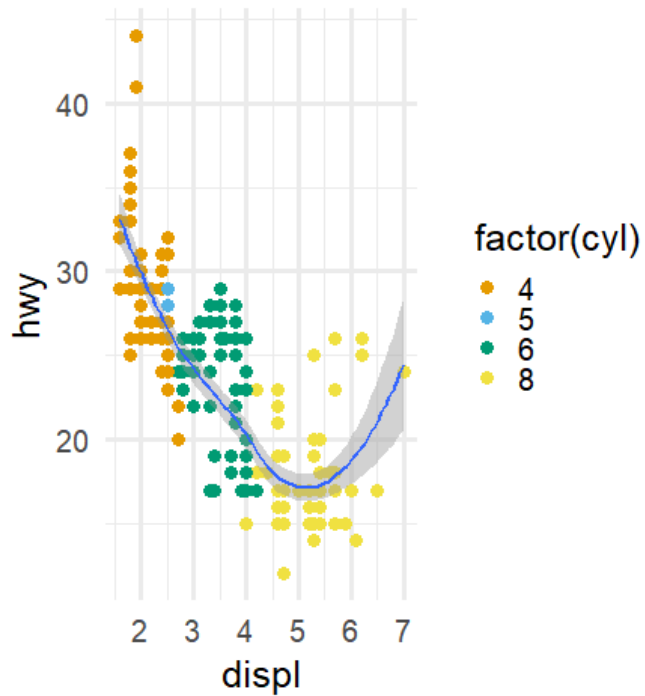
Example

- First, create two plots

```
p1 <- ggplot(mpg, aes(displ, hwy)) +  
  geom_point(aes(color = factor(cyl))) +  
  geom_smooth() +  
  scale_color_okabeIto()  
  
p2 <- ggplot(mpg, aes(factor(cyl), hwy)) +  
  geom_boxplot(aes(fill = factor(cyl))) +  
  scale_fill_okabeIto()
```

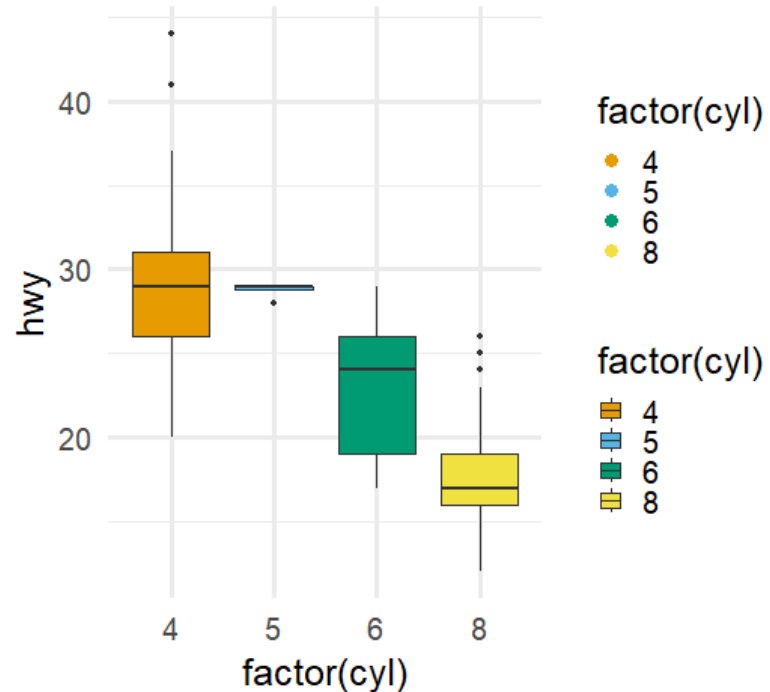
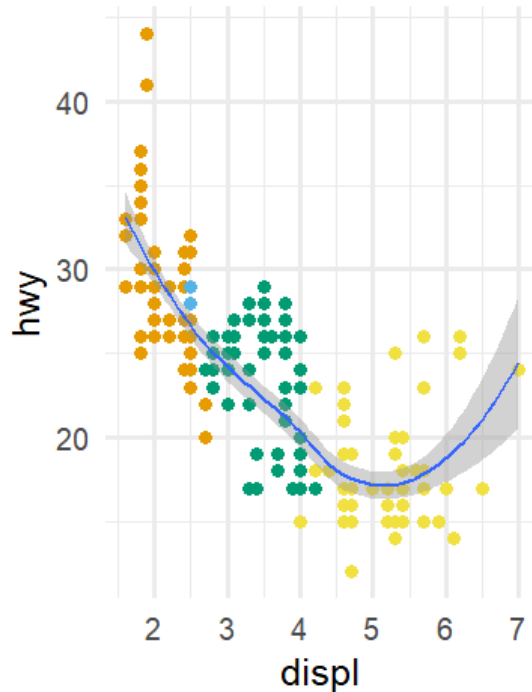
Side by side

```
library(patchwork)
p1 + p2
```



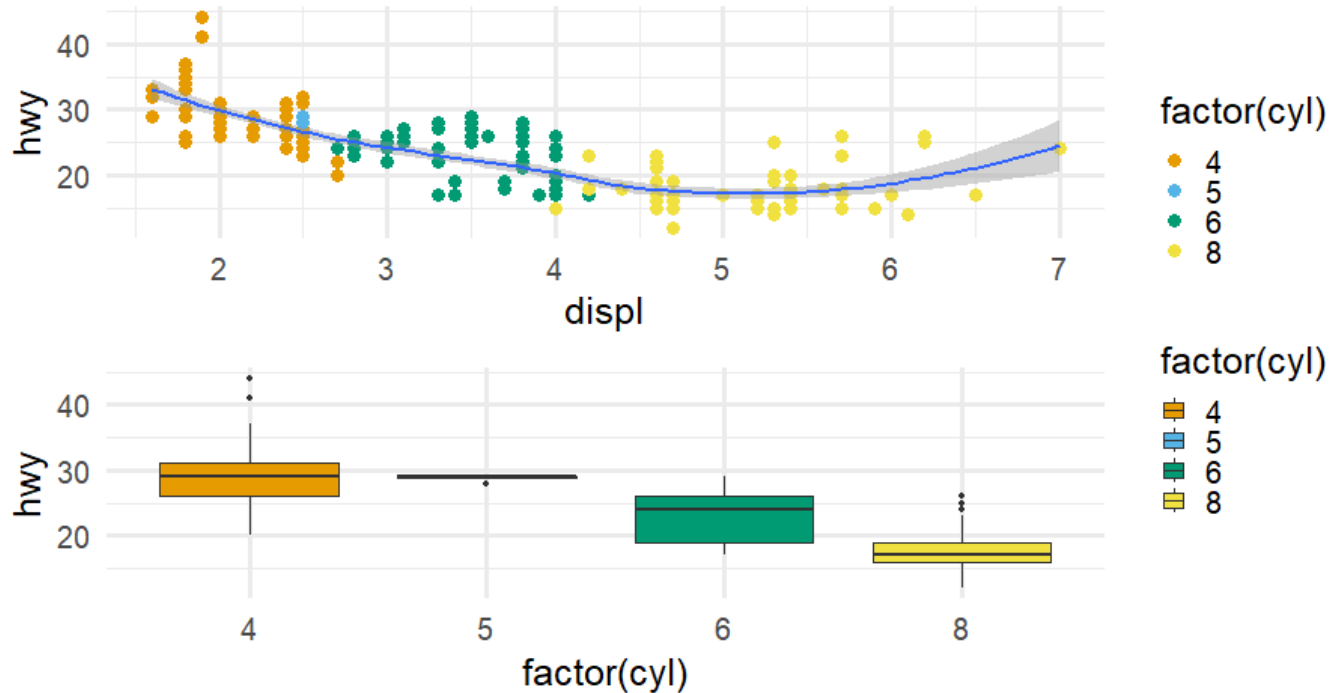
Collect legends

```
p1 + p2 + plot_layout(guides = "collect")
```



Stack vertically

```
p1 / p2 + plot_layout(guides = "collect")
```

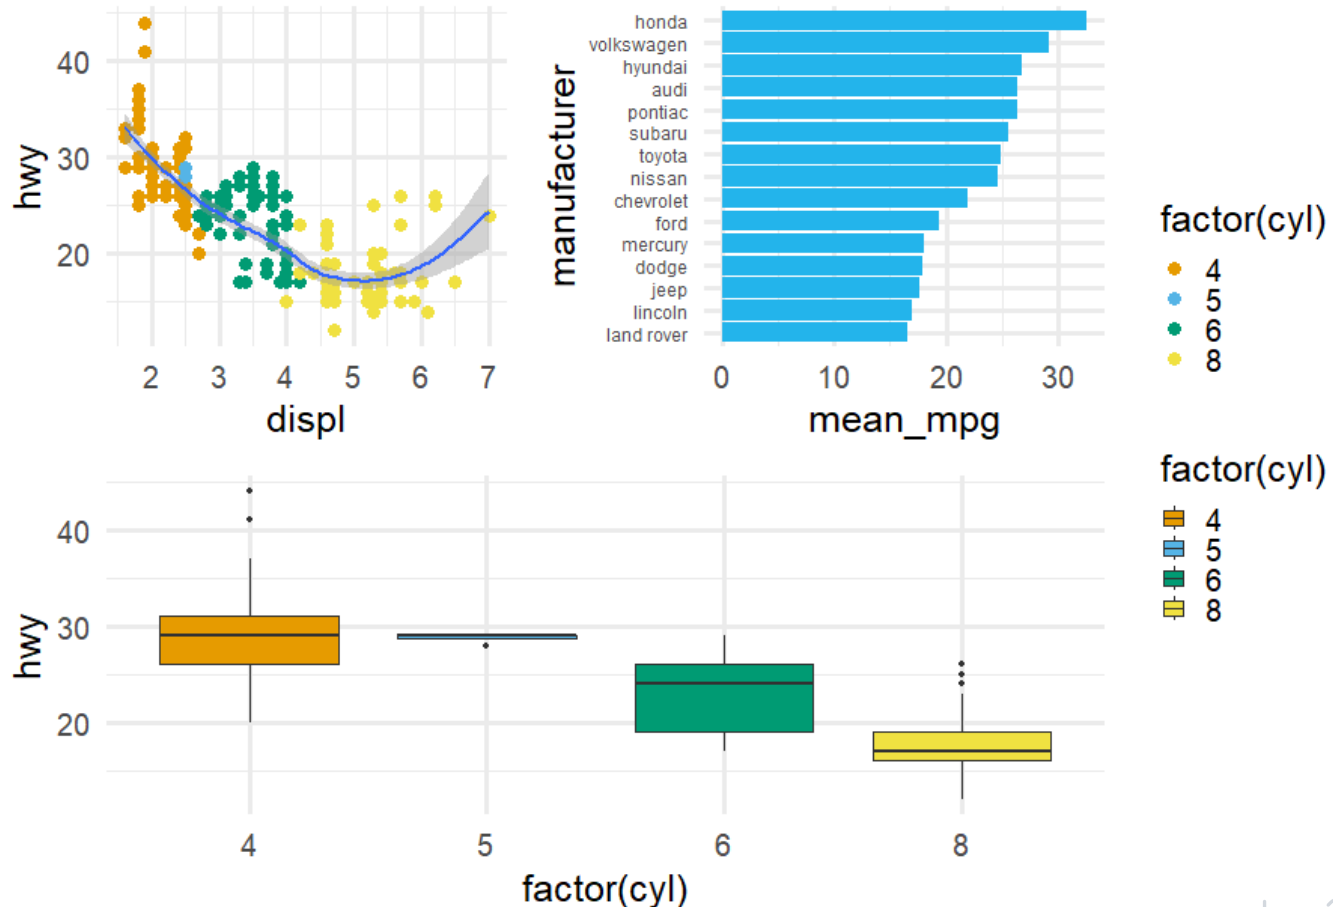


Add a third plot

```
p3 <- mpg %>%  
  group_by(manufacturer) %>%  
  summarize(mean_mpg = mean(hwy, na.rm = TRUE)) %>%  
  mutate(manufacturer = fct_reorder(manufacturer, mean_mpg)) %>%  
  ggplot(aes(mean_mpg, manufacturer)) +  
    geom_col(fill = "#25B6EE") +  
    theme(axis.text.y = element_text(size = 12))
```

Put box plot on bottom

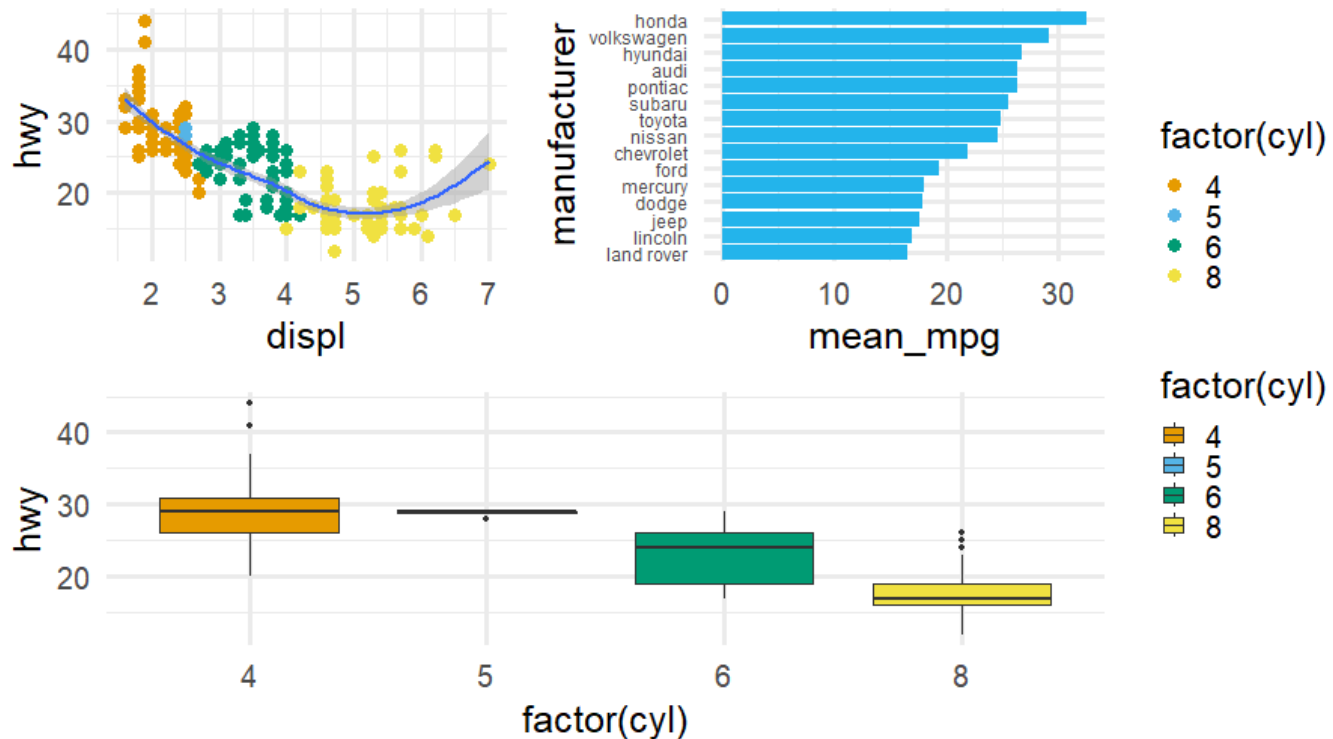
```
(p1 + p3) / p2 + plot_layout(guides = "collect")
```



Overall title

```
(p1 + p3) / p2 + plot_layout(guides = "collect") +  
  plot_annotation("Some cool plots")
```

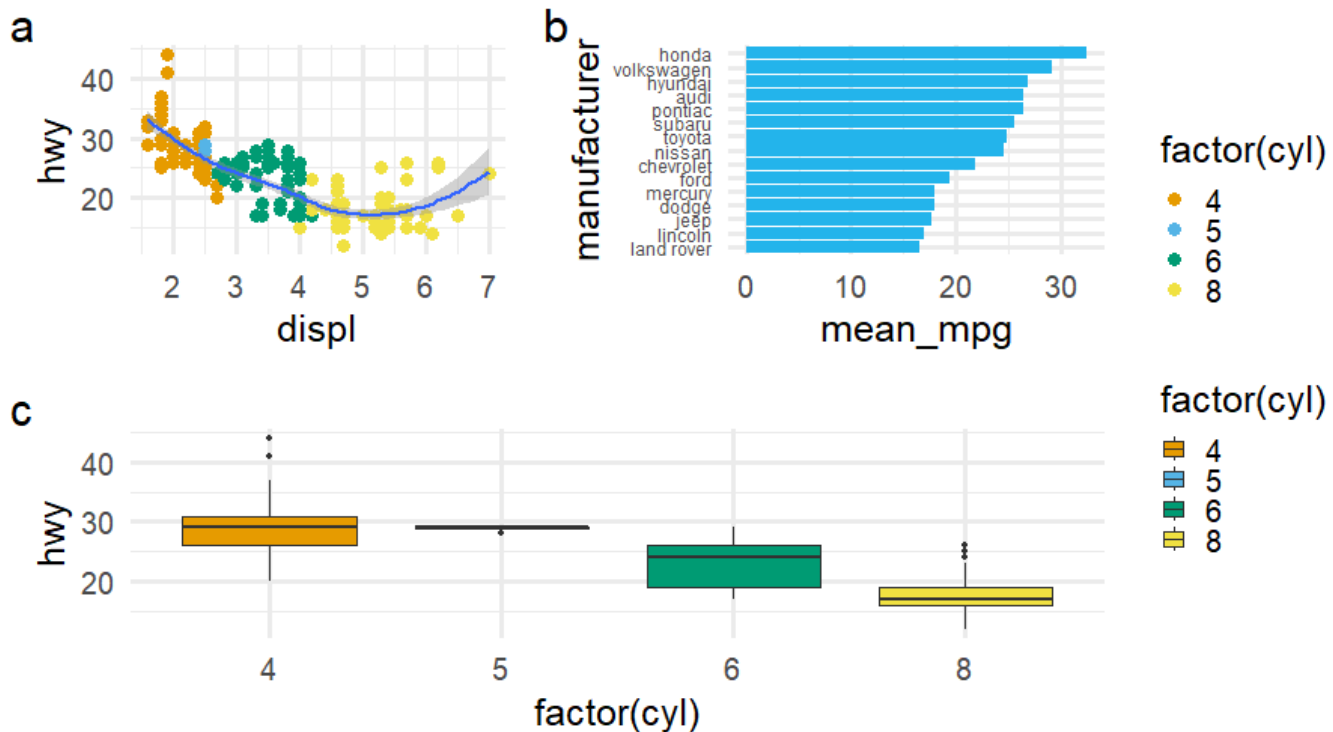
Some cool plots



Tags

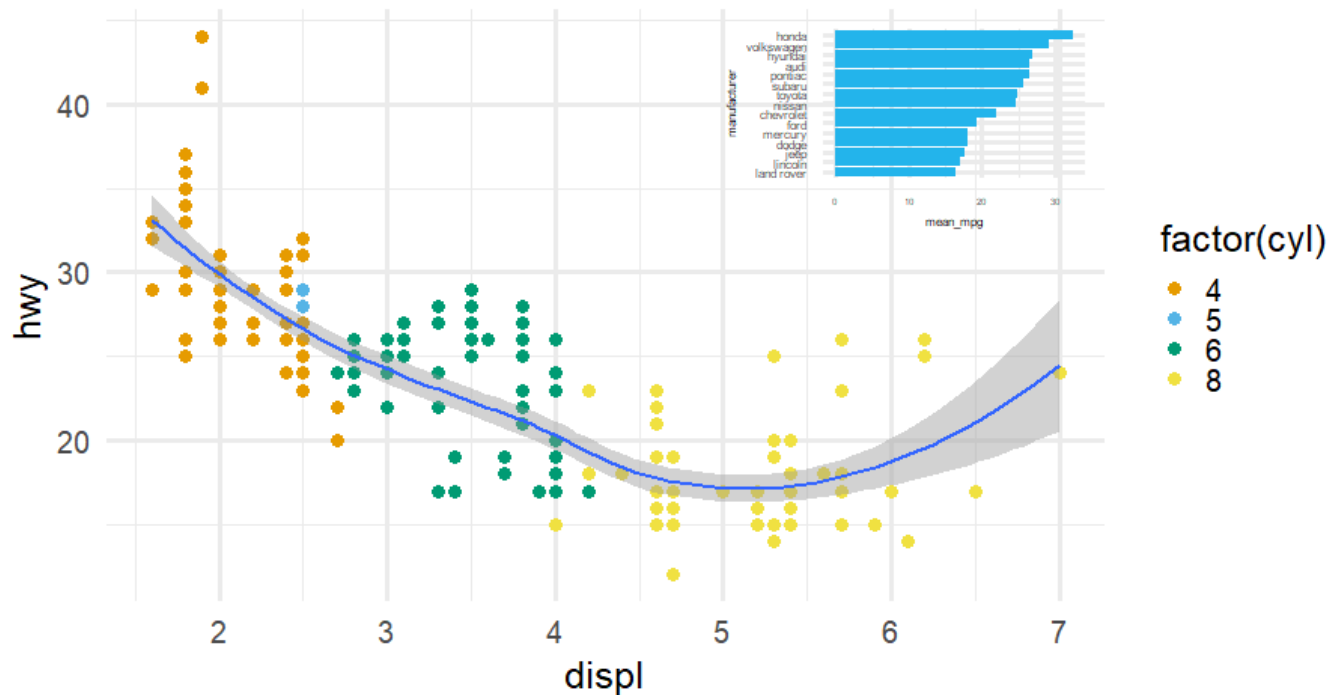
```
(p1 + p3) / p2 + plot_layout(guides = "collect") +  
  plot_annotation("Some cool plots", tag_levels = "a")
```

Some cool plots

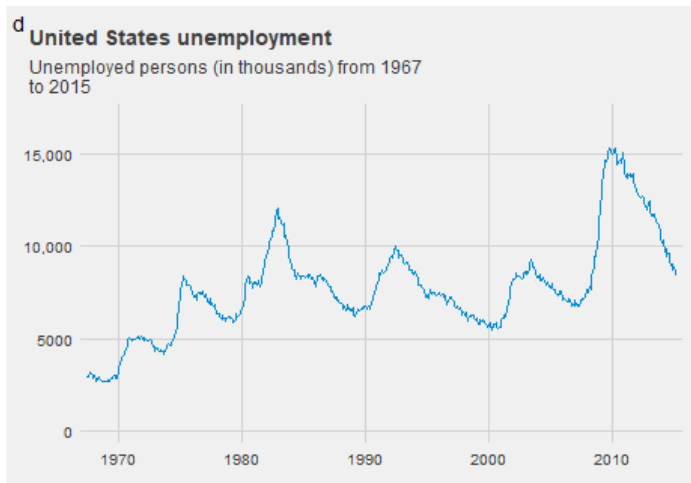
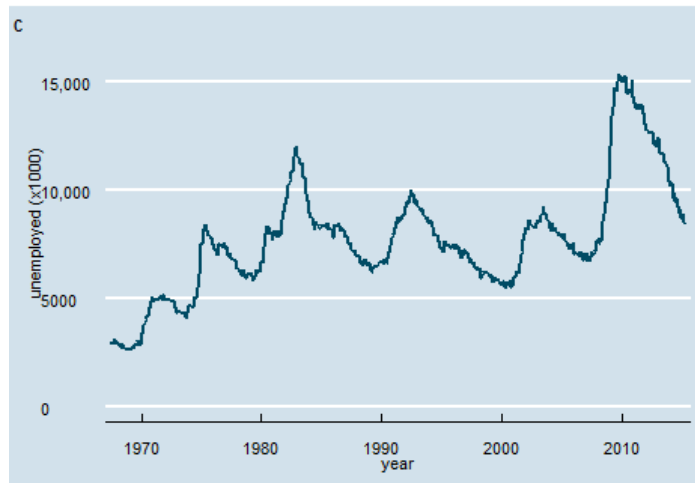
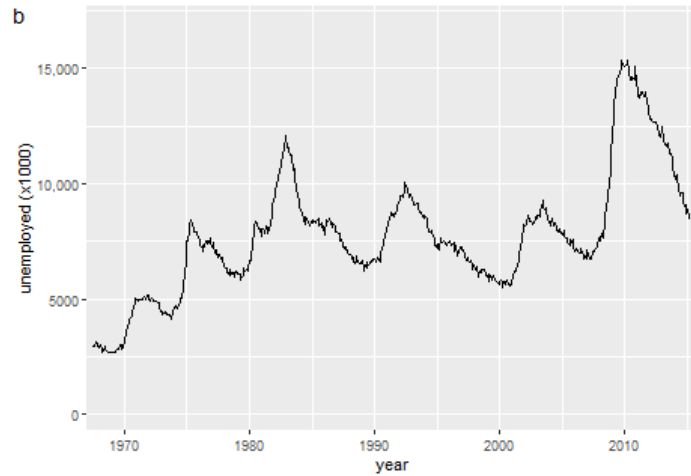
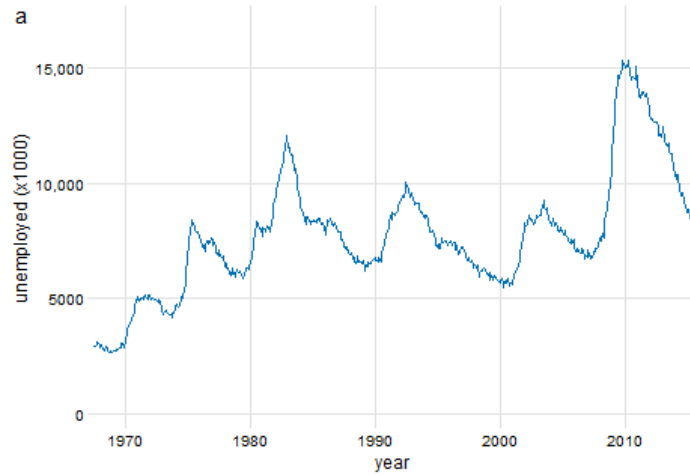


Insets

```
p3_small_txt <- p3 +  
  theme(axis.text.y = element_text(size = 8),  
        text = element_text(size = 8))  
  
p1 + inset_element(p3_small_txt, 0.6, 0.6, 1, 1)
```



Themes Refresher



ggthemes

- Good place to start. All sorts of themes.
- Includes color scales, etc., that align with themes
- You can even conform with other software
 - fit into an economics conference with `theme_stata`

See the themes [here](#)

BBC

The BBC uses ggplot for most of its graphics. They've developed a package with a theme and some functions to help make it match their style more.

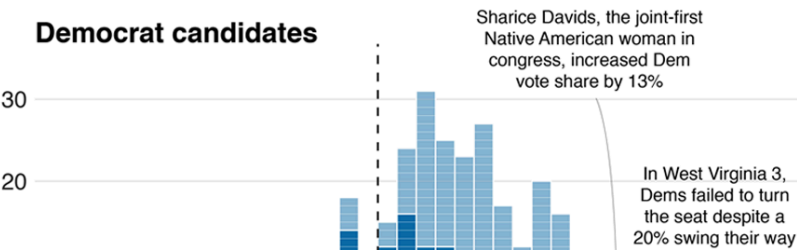
See the repo [here](#)

Their [Journalism Cookbook](#) is really nice too

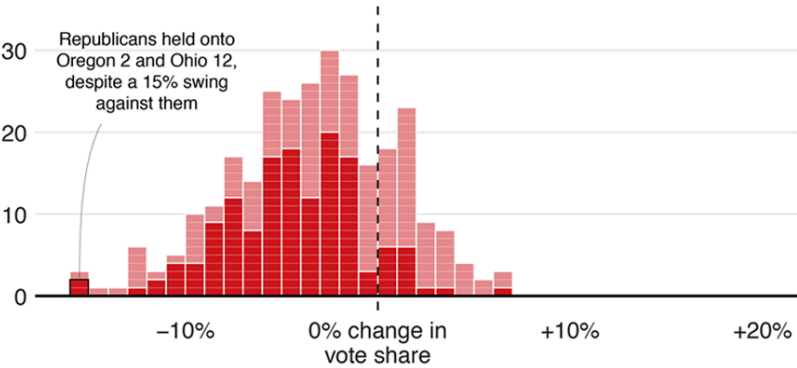
Blue wave

Won seat Didn't win

Democrat candidates



Republican candidates

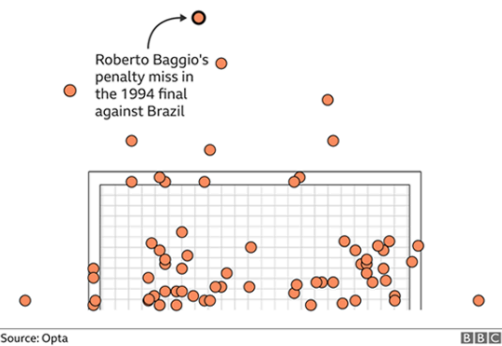


Source: AP, 19:01 ET



Where penalties are saved

World Cup shootout misses and saves, 1982-2014

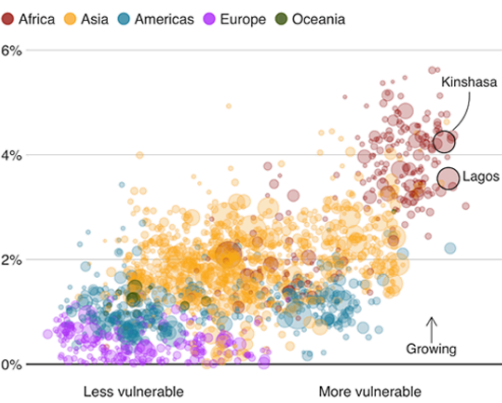


Source: Opta



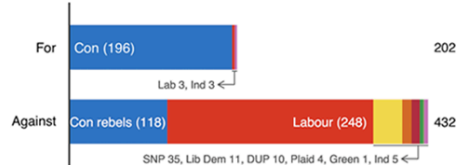
Fast-growing cities face worse climate risks

Population growth 2018-2035 over climate change vulnerability



Source: Verisk Maplecroft. Circle size represents current population.

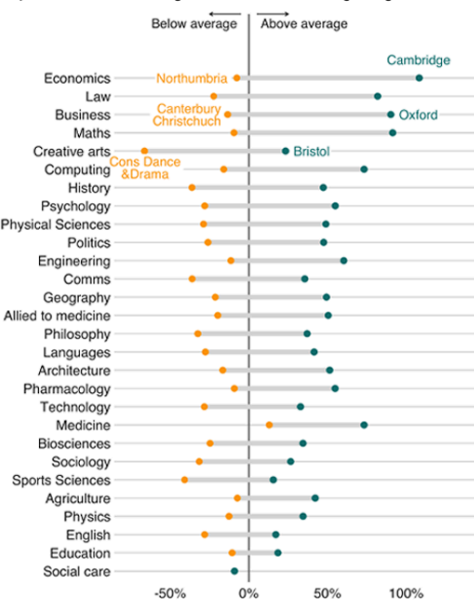
MPs rejected Theresa May's deal by 230 votes



Source: Commons Votes Services. Excludes 'tellers', the Speaker and deputies

Earnings vary across units even within subjects

Impact on men's earnings relative to the average degree



Source: Institute for Fiscal Studies



Similarly, the Urban Institute Visual Guide

See the repo [here](#)

So, I created one!

- Based on UO's visual guide [here](#)

[demo]

ggthemeassist

- Another great place to start with making major modifications/creating your own custom theme
- Can't do everything, but can do a lot
- See [here](#)

[demo]

theme() for everything else

- You can basically change your plot to look however you want through `theme`
- Generally a bit more complicated
- I've used ggplot for *years* and only really now gaining fluency with it

Last Quick example

The *google_trends* dataset comes from a [fivethirtyeight story](#) about how the media covered hurricanes and Trump.

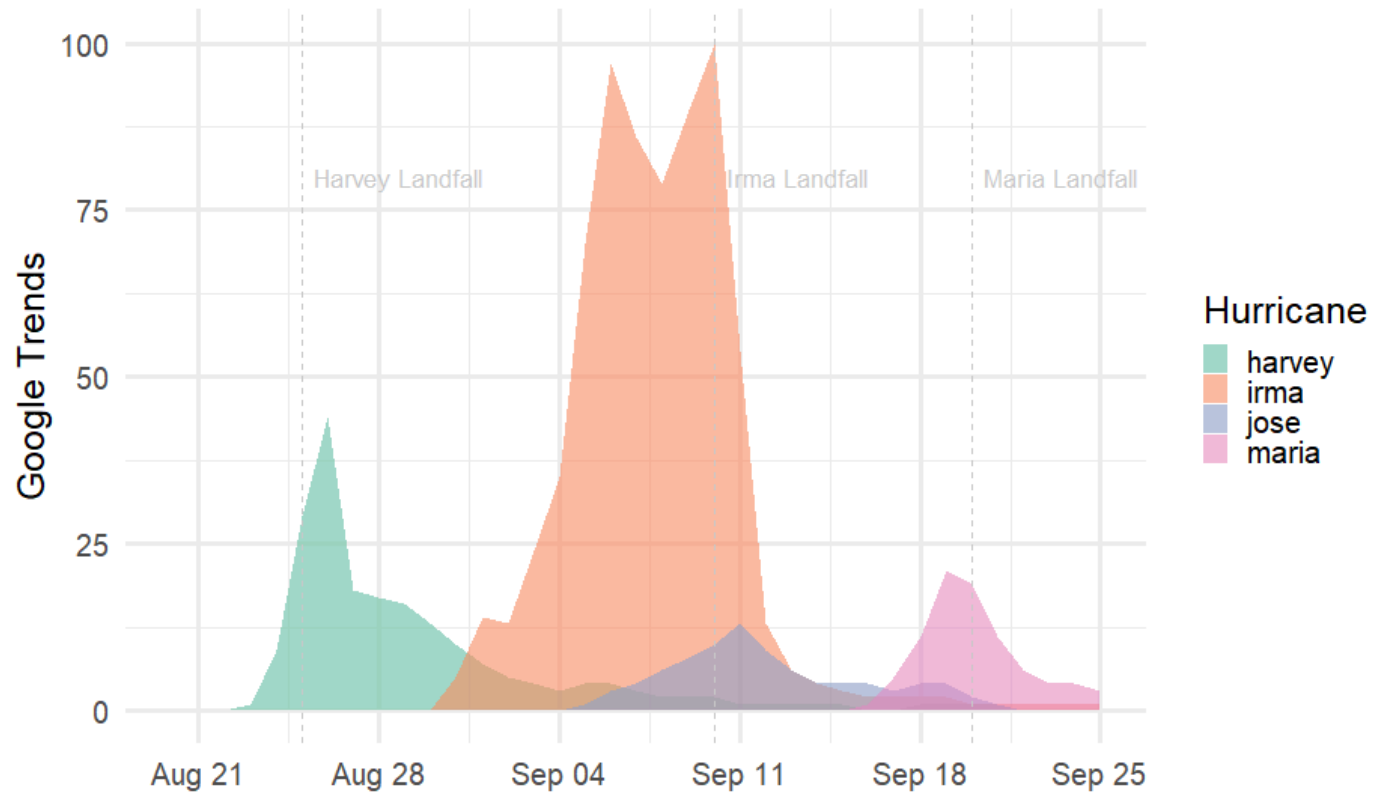
```
library(fivethirtyeight)
g <- google_trends %>%
  pivot_longer(starts_with("hurricane"),
               names_to = "hurricane",
               values_to = "interest",
               names_pattern = "_(.+)_")

landfall <- tibble(
  date = lubridate::mdy(
    c("August 25, 2017", "September 10, 2017", "September 20, 2017")
  ),
  hurricane = c("Harvey Landfall", "Irma Landfall", "Maria Landfall")
)
```

Let's start by visualizing the change in trends for each hurricane over time in one plot with three scales. We can map color to a discrete scale here. We can add vertical lines to show when each of them made landfall!

```
p <- ggplot(g, aes(date, interest)) +  
  geom_ribbon(aes(fill = hurricane, ymin = 0, ymax = interest),  
            alpha = 0.6) +  
  geom_vline(aes(xintercept = date), landfall,  
            color = "gray80",  
            lty = "dashed") +  
  geom_text(aes(x = date, y = 80, label = hurricane), landfall,  
            color = "gray80",  
            nudge_x = 0.5,  
            hjust = 0) +  
  labs(x = "",  
       y = "Google Trends",  
       title = "Hurricane Google trends over time",  
       caption = "Source: https://github.com/fivethirtyeight/data",  
       scale_fill_brewer("Hurricane", palette = "Set2")
```


Hurricane Google trends over time



Source: <https://github.com/fivethirtyeight/data/tree/master/puerto-rico-media>

We can use `ggthemeassist` to make a whole bunch of changes!

```
p + theme(  
  panel.grid.major = element_line(colour = "gray30"),  
  panel.grid.minor = element_line(colour = "gray30"),  
  axis.text = element_text(colour = "gray80"),  
  axis.text.x = element_text(colour = "gray80"),  
  axis.text.y = element_text(colour = "gray80"),  
  axis.title = element_text(colour = "gray80"),  
  legend.text = element_text(colour = "gray80"),  
  legend.title = element_text(colour = "gray80"),  
  panel.background = element_rect(fill = "gray10"),  
  plot.background = element_rect(fill = "gray10"),  
  legend.background = element_rect(fill = NA, color = NA),  
  legend.position = c(0.2, -0.1),  
  legend.direction = "horizontal",  
  plot.margin = margin(10, 10, b = 20, 10),  
  plot.caption = element_text(colour = "gray80", vjust = 1),  
  plot.title = element_text(colour = "gray80")  
)
```

Lab PS-3

Lab 5

Next time

Intro to Websites, Flex dashboards,
CSS customizations?

Note: Change in schedule and Lab-PS3- Last Lab problem
set in this class- is also posted